




# ***HiFi 1 Neural Network Library***

## **Programmer's Guide**

For HiFi DSPs



Cadence Design Systems, Inc.  
2655 Seely Ave.  
San Jose, CA 95134  
[www.cadence.com](http://www.cadence.com)

© 2021 Cadence Design Systems, Inc.  
All rights reserved worldwide

This publication is provided "AS IS." Cadence Design Systems, Inc. (hereafter "Cadence") does not make any warranty of any kind, either expressed or implied, including, but not limited to, the implied warranties of merchantability and fitness for a particular purpose. Information in this document is provided solely to enable system and software developers to use our processors. Unless specifically set forth herein, there are no express or implied patent, copyright or any other intellectual property rights or licenses granted hereunder to design or fabricate Cadence integrated circuits or integrated circuits based on the information in this document. Cadence does not warrant that the contents of this publication, whether individually or as one or more groups, meets your requirements or that the publication is error-free. This publication could include technical inaccuracies or typographical errors. Changes may be made to the information herein, and these changes may be incorporated in new editions of this publication.

© 2019 Cadence, the Cadence logo, Allegro, Assura, Broadband Spice, CDNLIVE!, Celtic, Chipestimate.com, Conformal, Connections, Denali, Diva, Dracula, Encounter, Flashpoint, FLIX, First Encounter, Incisive, Incyte, InstallScape, NanoRoute, NC-Verilog, OrCAD, OSKit, Palladium, PowerForward, PowerSI, PSpice, Purespec, Puresuite, Quickcycles, SignalStorm, Sigrity, SKILL, SoC Encounter, SourceLink, Spectre, Specman, Specman-Elite, SpeedBridge, Stars & Strikes, Tensilica, TripleCheck, TurboXim, Vectra, Virtuoso, VoltageStorm, Xcelium, Xplorer, Xtensa, and Xtreme are either trademarks or registered trademarks of Cadence Design Systems, Inc. in the United States and/or other jurisdictions.

OSCI, SystemC, Open SystemC, Open SystemC Initiative, and SystemC Initiative are registered trademarks of Open SystemC Initiative, Inc. in the United States and other countries and are used with permission. All other trademarks are the property of their respective holders.

Version 0.9.0  
July 2021

## Contents

---

1.	Introduction to the HiFi 1 NN Library .....	1
1.1	Organization of the HiFi 1 NN Library Package .....	1
1.1.1	Document Overview .....	2
1.2	HiFi 1 NN Library Specification .....	2
1.2.1	Low Level Kernels .....	2
1.2.2	Layers .....	3
1.2.3	Support for TensorFlow Lite Micro Operators .....	3
1.3	HiFi NN Library Performance .....	5
1.3.1	Memory Requirements .....	6
1.3.2	Runtime Memory – Low Level Kernels .....	6
1.3.3	Runtime Memory – Layers .....	10
1.3.4	Timings – Low-Level Kernels .....	14
1.3.5	Timings – Layers .....	28
2.	Generic HiFi NN Layer API .....	32
2.1	Shape .....	32
2.2	Memory Management .....	34
2.2.1	API Handle / Persistent Memory .....	34
2.2.2	Scratch Memory .....	34
2.2.3	Weights and Biases Memory .....	34
2.2.4	Input Buffer .....	34
2.2.5	Output Buffer .....	34
2.3	Generic API Errors .....	35
2.3.1	Common API Errors .....	35
2.4	C Language API .....	36
2.4.1	Startup Functions .....	37
2.4.2	Query Functions .....	38
2.4.3	Initialization Functions .....	38
2.4.4	Execution Functions .....	38
3.	HiFi 1 NN Library – Low-Level Kernels .....	39
3.1	Matrix X Vector Multiplication Kernels .....	39
3.1.1	Matrix X Vector Kernels .....	39
3.1.2	Fused (Activation) Matrix X Vector Kernels .....	43
3.1.3	Matrix X Vector Batch Kernels .....	47
3.1.4	Matrix Multiplication Kernels .....	49
3.1.5	Matrix X Vector Kernels with Output Stride .....	52
3.1.6	Matrix X Vector Batch Kernels with Accumulation .....	54
3.2	Convolution Kernels .....	56
3.2.1	Standard 2D Convolution Kernel .....	56

3.2.2 Standard 1D Convolution Kernel .....	60
3.2.3 Depthwise Separable 2D Convolution Kernel .....	64
3.2.3.1 Depthwise 2D Convolution Kernel.....	65
3.2.3.2 Pointwise 2D Convolution Kernel.....	70
3.3 Activation Kernels .....	74
3.3.1 Sigmoid.....	74
3.3.2 Tanh.....	76
3.3.3 Rectifier Linear Unit (ReLU).....	79
3.3.4 Softmax.....	82
3.3.5 Activation Min Max.....	83
3.3.6 Hard Swish.....	85
3.3.7 Parametric ReLU (PReLU).....	86
3.3.8 Leaky ReLU.....	88
3.4 Pooling Kernels .....	90
3.4.1 Average Pool Kernel.....	90
3.4.2 Max Pool Kernel.....	93
3.5 Fully connected Layer.....	97
3.5.1 Fully Connected Kernel.....	97
3.6 Basic Operations and Miscellaneous Kernels .....	100
3.6.1 Interpolation Kernel.....	100
3.6.2 Dot Product Kernels .....	101
3.6.3 Elementwise Requantize Kernels .....	102
3.6.4 Elementwise Dequantize Kernels .....	104
3.6.5 Basic Kernels.....	105
3.6.6 Broadcast Kernel .....	108
3.6.7 Memory Operation Kernels.....	109
3.7 Normalization Kernels.....	111
3.7.1 L2 Normalization Kernel.....	111
3.8 Reorg Kernels .....	112
3.8.1 Depth to Space Kernels .....	112
3.8.2 Space to Depth Kernels .....	114
3.8.3 Pad Kernel.....	116
3.8.4 Batch to Space Kernels .....	118
3.8.5 Space to Batch Kernels .....	122
4. HiFi 1 NN Library – Layers .....	125
4.1 GRU Layer.....	125
4.1.1 GRU Layer Specification.....	125
4.1.2 Error Codes Specific to GRU.....	126
4.1.3 API Functions Specific to GRU.....	127
4.1.3.1 Query Functions .....	127
4.1.3.2 Initialization Stage.....	129
4.1.3.3 Execution Stage.....	130
4.1.4 Structures Specific to GRU.....	133
4.1.5 Enums Specific to GRU.....	134

4.2	LSTM Layer .....	136
4.2.1	LSTM Layer Specification .....	136
4.2.2	Error Codes Specific to LSTM.....	136
4.2.3	API Functions Specific to LSTM.....	138
4.2.3.1	Query Functions .....	138
4.2.3.2	Initialization Stage.....	140
4.2.3.3	Execution Stage.....	141
4.2.4	Structures Specific to LSTM .....	145
4.2.5	Enums Specific to LSTM.....	146
4.3	CNN Layer.....	148
4.3.1	CNN Layer Specification.....	148
4.3.2	Error Codes Specific to CNN .....	148
4.3.3	API Functions Specific to CNN.....	149
4.3.3.1	Query Functions .....	149
4.3.3.2	Initialization Stage.....	152
4.3.3.3	Execution Stage.....	154
4.3.4	Structures Specific to CNN.....	157
4.3.5	Enums Specific to CNN .....	158
5.	Introduction to the Example Testbench.....	160
5.1	Making the Library.....	160
5.2	Making the Executable.....	160
5.3	Sample Testbench for Matrix X Vector Multiplication Kernels .....	161
5.3.1	Usage.....	161
5.4	Sample Testbench for Convolution Kernels.....	162
5.4.1	Usage.....	162
5.5	Sample Testbench for Activation Kernels.....	165
5.5.1	Usage.....	165
5.6	Sample Testbench for Pooling Kernels .....	166
5.6.1	Usage.....	166
5.7	Sample Testbench for Basic Operations Kernels.....	168
5.7.1	Usage.....	168
5.8	Sample Testbench for Normalization Kernels .....	170
5.8.1	Usage.....	170
5.9	Sample Testbench for Reorg Kernels.....	170
5.9.1	Usage.....	171
5.10	Sample Testbench for GRU Layer .....	172
5.10.1	Usage .....	172
5.11	Sample Testbench for LSTM Layer.....	173
5.11.1	Usage .....	173
5.12	Sample Testbench for CNN Layer .....	174
5.12.1	Usage .....	174
6.	References .....	176

## Figures

---

Figure 2-1 HiFi NN Layer Interfaces .....	32
Figure 2-2 Matrix and Cube (SHAPE_CUBE_DWH_T) Shape Representation .....	33
Figure 2-3 NN Layer Flow Overview .....	36
Figure 3-1 Broadcasting a 1x4x1 tensor to 1x4x3 and 2x4x3 .....	108
Figure 3-2 Depth to space conversion for 4x4x8 input with block size of 2 .....	112
Figure 3-3 Space to depth conversion for a 8x8x2 input with a block size of 2 .....	114
Figure 3-4 batch_to_space and space to batch conversion (for simplicity crop_sizes and pad_sizes are assumed to be 0) .....	121

## Tables

---

Table 2-1 Library Identification Functions .....	37
Table 4-1 GRU Get Persistent Size Function .....	127
Table 4-2 GRU Get Scratch Size Function .....	128
Table 4-3 GRU Init Function .....	129
Table 4-4 GRU Execution Function .....	130
Table 4-5 GRU Set Parameter Function Details .....	131
Table 4-6 GRU Get Parameter Function Details .....	132
Table 4-7 GRU Config Structure xa_nnlib_gru_init_config_t .....	133
Table 4-8 xa_nnlib_gru_weights_t Parameter Type .....	133
Table 4-9 xa_nnlib_gru_biases_t Parameter Type .....	134
Table 4-10 Enum xa_nnlib_gru_precision_t .....	134
Table 4-11 GRU Specific Parameters .....	135
Table 4-12 LSTM Get Persistent Size Function .....	138
Table 4-13 LSTM Get Scratch Size Function .....	139
Table 4-14 LSTM Init Function .....	140
Table 4-15 LSTM Execution Function .....	141
Table 4-16 LSTM Set Parameter Function Details .....	143
Table 4-17 LSTM Get Parameter Function Details .....	144
Table 4-18 LSTM Config Structure xa_nnlib_lstm_init_config_t .....	145
Table 4-19 xa_nnlib_lstm_weights_t Parameter Type .....	145
Table 4-20 xa_nnlib_lstm_biases_t Parameter Type .....	146

Table 4-21 Enum xa_nnlib_lstm_precision_t .....	146
Table 4-22 LSTM Specific Parameters .....	147
Table 4-23 CNN Get Persistent Size Function .....	149
Table 4-24 CNN Get Scratch Size Function.....	150
Table 4-25 CNN Init Function.....	152
Table 4-26 CNN Execution Function .....	154
Table 4-27 CNN Set Parameter Function Details .....	155
Table 4-28 CNN Get Parameter Function Details .....	156
Table 4-29 CNN Config Structure xa_nnlib_cnn_init_config_t.....	157
Table 4-30 Enum xa_nnlib_cnn_precision_t.....	158
Table 4-31 Enum xa_nnlib_cnn_algo_t.....	158
Table 4-32 CNN Specific Parameters.....	159

## Document Change History

---

Version	Changes
0.9.0	Initial version

DRAFT



# 1. Introduction to the HiFi 1 NN Library

The HiFi 1 Neural Network (NN) Library is a HiFi-optimized implementation of various NN layers and low level NN kernels. The library is designed with speech and audio neural network domain focus. The low level NN kernels are HiFi-optimized building blocks for NN layer implementation with a generic and simple interface. The NN layers are built using low level kernels and accept input in the form of 'shapes'<sup>1</sup> (up to four dimensions) and produce the output, also in the form of shapes. The layers use the weights or coefficients and biases stored 'externally'<sup>2</sup> for their operation. The shape of the input, output, weights and biases are as per the layer's design.

This guide refers to the HiFi 1 NN Library as HiFi NN Library, NN layers simply as layers and low level NN kernels as low-level kernels. The current version of the library implements GRU, LSTM (forward path), and CNN layers. It also implements matrix vector multiply, activation, pooling, and convolution functions as low-level kernels.

---

**Note** This version of the library supports HiFi 1 DSPs. The SP-VFPU (Single Precision Vector Floating Point Unit) is optional. The library can be compiled for HiFi 1 DSPs with or without the SP-VFPU enabled.

---

## 1.1 Organization of the HiFi 1 NN Library Package

The HiFi NN Library package includes the HiFi NN library containing all layers and low-level kernels implementations, and a set of sample test applications (for layers and low-level kernels).

The HiFi NN library implements a set of NN layers. The application can instantiate these layers and connect inputs and outputs across the layers to form a Neural Network system.

The HiFi NN library also provides a set of low level NN kernels. The application can use these kernels to implement or optimize performance of other NN layers.

The sample test applications implement a file-based application to test an instance of a layer or low level NN kernels for the given specification using pre-generated input, weight or coefficients and bias shapes stored in files in raw binary format.

---

<sup>1</sup> Refer to Section 2.1 Shape

<sup>2</sup> Refer to Section 2.2.3 Weights and Biases Memory

## 1.1.1 Document Overview

This document covers all the information required to integrate the HiFi NN Library into a Neural Network system. All the layers implement “HiFi NN layer APIs”, which is generic and explained in Section 2. The low level NN kernels are explained in Section 3. The APIs for each layer are described in Section 3.6.2. Section 5 provides details about available sample testbenches. References are listed in Section 6.

## 1.2 HiFi 1 NN Library Specification

The current version of the HiFi NN Library provides the following HiFi-optimized low-level kernels and layer implementations.

### 1.2.1 Low Level Kernels

- Matrix-vector multiplication kernels
- Convolution kernels
- Activation kernels
- Pooling kernels
- Basic operations kernels

These kernels support fixed point 8 bit, 16 bit, and single precision floating point data types for weights or coefficients, biases, input, and output. Refer to Section 3 for details.

Additionally, 8-bit and 16-bit quantized datatypes as defined in TensorFlow (TF), TensorFlow Lite for Microcontrollers (TFLM) are also supported for select kernels<sup>[3]</sup>. These datatypes use 8-bit/16-bit quantized values (asym8u – asymmetric 8-bit unsigned, asym8s – asymmetric 8-bit signed, sym8s – symmetric 8-bit signed) for weights or coefficients, input, and output. Biases are 32-bit quantized values.

8-bit quantized types are either unsigned (0, 255) or signed (-128, 127) 8-bit integer with 3 additional parameters.

Three numbers are associated with a quantized 8-bit value that can be used to convert the 8-bit integer to the real value and vice versa. These numbers are:

- Shift: an integer value indicating the amount of shift. If the value is positive, it is left shift and if negative, it is right shift
- Multiplier: a 32 bit (Q31) fixed point value greater than zero.
- Zero point: a 32 bit integer, in range [0, 255] for unsigned type, in range [-128, 127] for signed type.

The formula is:

$$\text{real\_value} = (\text{quantized\_value} - \text{zero\_point}) * 2^{(\text{shift})} * \text{multiplier}$$

The ‘sym8s’ type is symmetrical around 0, this means that quantized values are between -127 to 127 and zero point is 0, so all the calculation required due to zero point is avoided.

To match the asym8u/asym8s/sym8s APIs with Tensorflow, we define zero point as zero\_bias in the NN library APIs. The zero\_bias is an integer value having range asym8u - [0, 255], asym8s – [-128, 127] (or asym8u - [-255, 0], asym8s – [-127, 128] in case of the reverse operation depending on the corresponding Tensorflow kernel).

In addition to the quantized 8-bit datatypes, a similar 16-bit quantized datatype (asym16s) is used for few kernels. The zero\_bias for asym16s datatype is an integer value having range – [-32768, 32767].

## 1.2.2 Layers

- GRU layer (8x16, 16x16 precision)
- LSTM (forward path) layer (8x16, 16x16 precision)
- CNN layer (8x8, 8x16, 16x16, and float32xfloat32 precision)

**Note:** MxN precision above denotes (weights or coefficients) x (input, output, bias) precision. Refer to Section 3.6.2 for details.

## 1.2.3 Support for TensorFlow Lite Micro Operators

The HiFi 1 NN Library low level kernels can be used to implement the following operators of TensorFlow Lite Micro:

No.	Operator	Float32 Datatype Support	Uint8 (asymmetric quantized uint8) Datatype Support	Int8 (quantized int8) Datatype Support	Boolean (1 Byte) Datatype Support
1	FULLY_CONNECTED		Yes	Yes	
2	MAX_POOL_2D	Yes		Yes	
3	SOFTMAX	Yes	Yes	Yes	
4	LOGISTIC	Yes		Yes	
5	SVDF			Yes	
6	CONV_2D	Yes	Yes	Yes	
7	DEPTHWISE_CONV_2D	Yes	Yes	Yes	
8	AVERAGE_POOL_2D	Yes	Yes	Yes	
9	FLOOR	Yes			
10	RELU	Yes		Yes	
11	RELU6	Yes	Yes		
12	ADD	Yes	Yes		
13	MUL	Yes	Yes		
14	QUANTIZE <sup>3</sup>			Yes	
15	HARDSWISH			Yes	

<sup>3</sup> QUANTIZE operator has different input and output quantized data types, HiF1 NN Library has kernels for Int16 to Int8, Int8 to Int32, Int16 to Int32.

16	MAXIMUM			Yes	
17	MINIMUM			Yes	
18	PRELU			Yes	
19	SUB			Yes	
20	TANH			Yes	
21	L2 NORM	Yes		Yes	
22	ABS	Yes			
23	SIN	Yes			
24	COS	Yes			
25	LOG	Yes			
26	SQRT	Yes			
27	RSQRT	Yes			
28	SQUARE	Yes			
29	FILL	Yes			
30	CEIL	Yes			
31	ROUND	Yes			
32	NEG	Yes			
33	DEQUANTIZE			Yes <sup>4</sup>	
34	LEAKY_RELU			Yes	
35	PAD			Yes	
36	CIRCULAR_BUFFER			Yes	
37	DEPTH_TO_SPACE			Yes	
38	BATCH_TO_SPACE_ND			Yes	
39	SPACE_TO_BATCH_ND			Yes	

Following TFLM operators get optimized out of box on HiFi1 and don't require any HiFi1 NNLib kernels:

No.	Operator	Float32 Datatype Support	Uint8 (asymmetric quantized uint8) Datatype Support	Int8 (quantized int8) Datatype Support	Int32	Int64	Boolean (1 Byte) Datatype Support
1	PACK	Yes	Yes	Yes	Yes	Yes	
2	EXPAND_DIMS	Yes		Yes			
3	RESHAPE <sup>5</sup>						
4	SQUEEZE <sup>5</sup>						

<sup>4</sup> For TFLM DEQUANTIZE operator output is always single precision float whereas multiple input data types are supported. HiFi1 NN Library has kernel for quantized Int8 input datatype.

<sup>5</sup> For RESHAPE and SQUEEZE datatype is not specified in Tensorflow Lite Micro.

## ***1.3 HiFi NN Library Performance***

The HiFi NN library from Cadence was characterized on the 2-stage HiFi DSP. The memory usage and performance figures are provided for design reference.

DRAFT

### 1.3.1 Memory Requirements

The NN library is provided as a single library archive and the Text and Read-Only Data (ROData) sizes of the archive are shown in Table 1-1. Floating point variants of kernels/layers are only included in library compiled for HiFi DSPs with the SP-VFPU (optional Single Precision Vector Floating Point Unit).

Table 1-1 Library Text and ROData Sizes

DSP	Neural Network Library	
	Text (in KBytes)	Data (in KBytes)
HiFi 1 (without VFPU)	263.49	1.41
HiFi 1 (with VFPU)	305.42	1.81

### 1.3.2 Runtime Memory – Low Level Kernels

Some low-level kernels require temporary scratch memory for algorithm implementation. The runtime memory requirements for low-level kernels is shown in Table 1-2.

Table 1-2 Low Level Kernels Runtime Memory Requirements

LowLevel Kernel	Parameters	Scratch (Bytes)
matXvec_16x16_16_sigmoid	rows=1024; cols1=1024; cols2=1024; bias_prec=16	4096
matXvec_16x16_16_sigmoid	rows=1024; cols1=1024; cols2=1024; bias_prec=64	4096
matXvec_8x16_16_sigmoid	rows=1024; cols1=1024; cols2=1024; bias_prec=16	4096
matXvec_8x16_16_sigmoid	rows=1024; cols1=1024; cols2=1024; bias_prec=64	4096
matXvec_8x8_8_sigmoid	rows=1024; cols1=1024; cols2=1024; bias_prec=8	4096
matXvec_8x8_8_sigmoid	rows=1024; cols1=1024; cols2=1024; bias_prec=32	4096
matXvec_f32xf32_f32_sigmoid	rows=1024; cols1=1024; cols2=1024; bias_prec=-1	4096
matXvec_16x16_16_tanh	rows=1024; cols1=1024; cols2=1024; bias_prec=16	4096
matXvec_16x16_16_tanh	rows=1024; cols1=1024; cols2=1024; bias_prec=64	4096
matXvec_8x16_16_tanh	rows=1024; cols1=1024; cols2=1024; bias_prec=16	4096
matXvec_8x16_16_tanh	rows=1024; cols1=1024; cols2=1024; bias_prec=64	4096
matXvec_8x8_8_tanh	rows=1024; cols1=1024; cols2=1024; bias_prec=8	4096
matXvec_8x8_8_tanh	rows=1024; cols1=1024; cols2=1024; bias_prec=32	4096
matXvec_f32xf32_f32_tanh	rows=1024; cols1=1024; cols2=1024; bias_prec=-1	4096
matXvec_16x16_16_sigmoid	rows=256; cols1=256; cols2=256; bias_prec=16	1024
matXvec_16x16_16_sigmoid	rows=256; cols1=256; cols2=256; bias_prec=64	1024
matXvec_8x16_16_sigmoid	rows=256; cols1=256; cols2=256; bias_prec=16	1024
matXvec_8x16_16_sigmoid	rows=256; cols1=256; cols2=256; bias_prec=64	1024
matXvec_8x8_8_sigmoid	rows=256; cols1=256; cols2=256; bias_prec=8	1024
matXvec_8x8_8_sigmoid	rows=256; cols1=256; cols2=256; bias_prec=32	1024
matXvec_f32xf32_f32_sigmoid	rows=256; cols1=256; cols2=256; bias_prec=-1	1024
matXvec_16x16_16_tanh	rows=256; cols1=256; cols2=256; bias_prec=16	1024

LowLevel Kernel	Parameters	Scratch (Bytes)
matXvec_16x16_16_tanh	rows=256; cols1=256; cols2=256; bias_prec=64	1024
matXvec_8x16_16_tanh	rows=256; cols1=256; cols2=256; bias_prec=16	1024
matXvec_8x16_16_tanh	rows=256; cols1=256; cols2=256; bias_prec=64	1024
matXvec_8x8_8_tanh	rows=256; cols1=256; cols2=256; bias_prec=8	1024
matXvec_8x8_8_tanh	rows=256; cols1=256; cols2=256; bias_prec=32	1024
matXvec_f32xf32_f32_tanh	rows=256; cols1=256; cols2=256; bias_prec=-1	1024
conv2d_depth_8x8	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	992
conv2d_depth_8x8	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	3968
conv2d_depth_8x8_nhwc	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	5296
conv2d_depth_8x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	1376
conv2d_depth_8x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	5792
conv2d_depth_8x16_nhwc	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	10552
conv2d_depth_16x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	1376
conv2d_depth_16x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	5792
conv2d_depth_16x16_nhwc	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	10552
conv2d_depth_f32xf32	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	1856
conv2d_depth_f32xf32	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	8384
conv2d_depth_f32xf32_nhwc	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	20512
conv2d_depth_asym8xasym8	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	704
conv2d_depth_asym8xasym8	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	2912

LowLevel Kernel	Parameters	Scratch (Bytes)
conv2d_depth_asym8xasym8_nhwc	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	5296
conv2d_depth_sym8sxasym8s	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	704
conv2d_depth_sym8sxasym8s	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	2912
conv2d_depth_sym8sxasym8s_nhwc	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	5296
conv1d_std_8x8	input_height=32; input_width=40; input_channels=32; kernel_height=7; out_channels=24; out_height=26; ;	8991
conv1d_std_8x8	input_height=128; input_width=128; input_channels=8; kernel_height=11; out_channels=8; out_height=128; ;	11295
conv1d_std_8x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; out_channels=24; out_height=26; ;	17951
conv1d_std_8x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; out_channels=8; out_height=128; ;	22559
conv1d_std_16x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; out_channels=24; out_height=26; ;	17951
conv1d_std_16x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; out_channels=8; out_height=128; ;	22559
conv1d_std_f32xf32	input_height=32; input_width=40; input_channels=32; kernel_height=7; out_channels=24; out_height=26; ;	35871
conv1d_std_f32xf32	input_height=128; input_width=128; input_channels=8; kernel_height=11; out_channels=8; out_height=128; ;	45087
conv1d_std_asym8xasym8	input_height=32; input_width=40; input_channels=32; kernel_height=7; out_channels=24; out_height=26; ;	8991
conv1d_std_asym8xasym8	input_height=128; input_width=128; input_channels=8; kernel_height=11; out_channels=8; out_height=128; ;	11295
conv2d_std_8x8	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	5166
conv2d_std_8x8	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	12190
conv2d_std_8x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	10286
conv2d_std_8x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	24334
conv2d_std_16x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	10286
conv2d_std_16x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	24334



LowLevel Kernel	Parameters	Scratch (Bytes)
conv2d_std_f32xf32	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	20526
conv2d_std_f32xf32	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	48622
conv2d_std_asym8xasym8	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	5166
conv2d_std_asym8xasym8	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	12190
conv2d_std_sym8sxasym8s	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	5166
conv2d_std_sym8sxasym8s	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	12190
avgpool_8	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	584
avgpool_8	input_height=384; input_width=128; input_channels=8; kernel_height=12; kernel_width=4; out_height=43; out_width=43	1768
avgpool_8_nhwc	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	4960
avgpool_16	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	584
avgpool_16	input_height=384; input_width=128; input_channels=8; kernel_height=12; kernel_width=4; out_height=43; out_width=43	1768
avgpool_16_nhwc	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	9664
avgpool_f32	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	464
avgpool_f32	input_height=384; input_width=128; input_channels=8; kernel_height=12; kernel_width=4; out_height=43; out_width=43	1424
avgpool_f32_nhwc	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	13048
maxpool_8	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	232
maxpool_8	input_height=384; input_width=128; input_channels=8; kernel_height=12; kernel_width=4; out_height=43; out_width=43	712
maxpool_8_nhwc	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	3200
maxpool_16	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	232
maxpool_16	input_height=384; input_width=128; input_channels=8; kernel_height=12; kernel_width=4; out_height=43; out_width=43	712
maxpool_16_nhwc	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	3200
maxpool_f32	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	456

LowLevel Kernel	Parameters	Scratch (Bytes)
maxpool_f32	input_height=384; input_width=128; input_channels=8; kernel_height=12; kernel_width=4; out_height=43; out_width=43	1416
maxpool_f32_nhwc	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	6272
avgpool_asym8	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	584
avgpool_asym8	input_height=384; input_width=128; input_channels=8; kernel_height=12; kernel_width=4; out_height=43; out_width=43	1768
maxpool_asym8	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	456
maxpool_asym8	input_height=384; input_width=128; input_channels=8; kernel_height=12; kernel_width=4; out_height=43; out_width=43	1416
avgpool_asym8_nhwc	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	4960
maxpool_asym8_nhwc	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	3200

**Note** f32 denotes kernel/layer variants using single precision floating point data type.

### 1.3.3 Runtime Memory – Layers

The runtime memory requirements specific to each layer are shown in Table 1-3.

Table 1-3 Runtime Memory Requirements

Layer	Parameters	Persistent (Bytes)	Scratch (Bytes)	Input (Bytes)	Output (Bytes)
gru_16x16	in_feats=32; out_feats=256	904	2576	64	512
gru_16x16	in_feats=36; out_feats=256	904	2576	72	512
gru_16x16	in_feats=256; out_feats=256	904	2576	512	512
gru_16x16	in_feats=128; out_feats=1024	2440	10256	256	2048
gru_16x16	in_feats=1024; out_feats=1024	2440	10256	2048	2048
gru_8x16	in_feats=32; out_feats=256	904	2576	64	512
gru_8x16	in_feats=36; out_feats=256	904	2576	72	512
gru_8x16	in_feats=256; out_feats=256	904	2576	512	512
gru_8x16	in_feats=128; out_feats=1024	2440	10256	256	2048
gru_8x16	in_feats=1024; out_feats=1024	2440	10256	2048	2048
lstm_16x16	in_feats=32; out_feats=256	2064	2576	64	512
lstm_16x16	in_feats=36; out_feats=256	2064	2576	72	512
lstm_16x16	in_feats=256; out_feats=256	2064	2576	512	512
lstm_16x16	in_feats=128; out_feats=1024	6672	10256	256	2048

lstm_16x16	in_feats=1024 out_feats=1024	6672	10256	2048	2048
lstm_8x16	in_feats=32; out_feats=256	2064	2576	64	512
lstm_8x16	in_feats=36; out_feats=256	2064	2576	72	512
lstm_8x16	in_feats=256; out_feats=256	2064	2576	512	512
lstm_8x16	in_feats=128; out_feats=1024	6672	10256	256	2048
lstm_8x16	in_feats=1024 out_feats=1024	6672	10256	2048	2048
cnn_conv2d_depth_8x8	input_height=32 input_width=40 input_channels=32 kernel_height=7 kernel_width=5 out_channels=24 out_height=26; out_width=36	368	30944	40960	22464
cnn_conv2d_depth_8x8	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	368	135040	131072	131072
cnn_conv2d_depth_8x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	368	61280	81920	44928
cnn_conv2d_depth_8x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	368	267936	262144	262144
cnn_conv2d_depth_16x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	368	61280	81920	44928
cnn_conv2d_depth_16x16	input_height=128; input_width=128; input_channels=8;	368	267936	262144	262144

	kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128				
cnn_conv2d_depth_f32xf32	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	368	121664	163840	89856
cnn_conv2d_depth_f32xf32	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	368	532672	524288	524288
cnn_conv1d_std_8x8	input_height=32; input_width=40; input_channels=32; kernel_height=7; out_channels=24; out_height=26;;	368	8991	40960	624
cnn_conv1d_std_8x8	input_height=128; input_width=128; input_channels=8; kernel_height=11; out_channels=8; out_height=128;;	368	11295	131072	1024
cnn_conv1d_std_8x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; out_channels=24; out_height=26;;	368	17951	81920	1248
cnn_conv1d_std_8x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; out_channels=8; out_height=128;;	368	22559	262144	2048
cnn_conv1d_std_16x16	input_height=32; input_width=40; input_channels=32; kernel_height=7;	368	17951	81920	1248

	out_channels=24; out_height=26;				
cnn_conv1d_std_16x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; out_channels=8; out_height=128;;	368	22559	262144	2048
cnn_conv1d_std_f32xf32	input_height=32; input_width=40; input_channels=32; kernel_height=7; out_channels=24; out_height=26;;	368	35871	163840	2496
cnn_conv1d_std_f32xf32	input_height=128; input_width=128; input_channels=8; kernel_height=11; out_channels=8; out_height=128;;	368	45087	524288	4096
cnn_conv2d_std_8x8	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	368	5166	40960	22464
cnn_conv2d_std_8x8	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	368	12190	131072	131072
cnn_conv2d_std_8x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	368	10286	81920	44928
cnn_conv2d_std_8x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8;	368	24334	262144	262144

	out_height=128; out_width=128				
cnn_conv2d_std_16x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	368	10286	81920	44928
cnn_conv2d_std_16x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	368	24334	262144	262144
cnn_conv2d_std_f32xf32	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	368	20526	163840	89856
cnn_conv2d_std_f32xf32	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	368	48622	524288	524288

### 1.3.4 Timings – Low-Level Kernels

Table 1-4 Low-Level Kernels Timings

Low Level Kernel	Parameters	Average Cycles	Performance Metric
matXvec_16x16_16_sigmoid	rows=1024; cols1=1024; cols2=1024; bias_prec=16	618317	3.39 (MACs/cyc)
matXvec_16x16_16_sigmoid	rows=1024; cols1=1024; cols2=1024; bias_prec=64	617291	3.40 (MACs/cyc)
matXvec_8x16_16_sigmoid	rows=1024; cols1=1024; cols2=1024; bias_prec=16	619086	3.39 (MACs/cyc)
matXvec_8x16_16_sigmoid	rows=1024; cols1=1024; cols2=1024; bias_prec=64	618700	3.39 (MACs/cyc)

Low Level Kernel	Parameters	Average Cycles	Performance Metric
matXvec_8x8_8_sigmoid	rows=1024; cols1=1024; cols2=1024; bias_prec=8	619339	3.39 (MACs/cyc)
matXvec_8x8_8_sigmoid	rows=1024; cols1=1024; cols2=1024; bias_prec=32	618571	3.39 (MACs/cyc)
matXvec_f32xf32_f32_sigmoid	rows=1024; cols1=1024; cols2=1024; bias_prec=-1	1342547	1.56 (MACs/cyc)
matXvec_16x16_16_tanh	rows=1024; cols1=1024; cols2=1024; bias_prec=16	618830	3.39 (MACs/cyc)
matXvec_16x16_16_tanh	rows=1024; cols1=1024; cols2=1024; bias_prec=64	617804	3.39 (MACs/cyc)
matXvec_8x16_16_tanh	rows=1024; cols1=1024; cols2=1024; bias_prec=16	619598	3.38 (MACs/cyc)
matXvec_8x16_16_tanh	rows=1024; cols1=1024; cols2=1024; bias_prec=64	619212	3.39 (MACs/cyc)
matXvec_8x8_8_tanh	rows=1024; cols1=1024; cols2=1024; bias_prec=8	620369	3.38 (MACs/cyc)
matXvec_8x8_8_tanh	rows=1024; cols1=1024; cols2=1024; bias_prec=32	619601	3.38 (MACs/cyc)
matXvec_f32xf32_f32_tanh	rows=1024; cols1=1024; cols2=1024; bias_prec=-1	1347295	1.56 (MACs/cyc)
matXvec_16x16_16	rows=1024; cols1=1024; cols2=1024; bias_prec=16	605688	3.46 (MACs/cyc)
matXvec_16x16_32	rows=1024; cols1=1024; cols2=1024; bias_prec=16	605816	3.46 (MACs/cyc)
matXvec_16x16_64	rows=1024; cols1=1024; cols2=1024; bias_prec=16	605292	3.46 (MACs/cyc)
matXvec_8x16_16	rows=1024; cols1=1024; cols2=1024; bias_prec=16	606585	3.46 (MACs/cyc)
matXvec_8x16_32	rows=1024; cols1=1024; cols2=1024; bias_prec=16	606202	3.46 (MACs/cyc)
matXvec_8x16_64	rows=1024; cols1=1024; cols2=1024; bias_prec=16	605423	3.46 (MACs/cyc)
matXvec_8x8_8	rows=1024; cols1=1024; cols2=1024; bias_prec=8	606459	3.46 (MACs/cyc)
matXvec_8x8_16	rows=1024; cols1=1024; cols2=1024; bias_prec=8	605306	3.46 (MACs/cyc)
matXvec_8x8_32	rows=1024; cols1=1024; cols2=1024; bias_prec=8	605049	3.47 (MACs/cyc)
matXvec_f32xf32_f32	rows=1024; cols1=1024; cols2=1024; bias_prec=-1	1325180	1.58 (MACs/cyc)
matXvec_asym8xasym8_asym8	rows=1024; cols1=1024; cols2=1024; bias_prec=32	1199036	1.75 (MACs/cyc)
matXvec_sym8sxasym8s_asym8s	rows=1024; cols1=1024; cols2=1024; bias_prec=32	6318178	0.33 (MACs/cyc)
matXvec_16x16_16_sigmoid	rows=256; cols1=256; cols2=256; bias_prec=16	44141	2.97 (MACs/cyc)
matXvec_16x16_16_sigmoid	rows=256; cols1=256; cols2=256; bias_prec=64	43883	2.99 (MACs/cyc)
matXvec_8x16_16_sigmoid	rows=256; cols1=256; cols2=256; bias_prec=16	44334	2.96 (MACs/cyc)



Low Level Kernel	Parameters	Average Cycles	Performance Metric
matXvec_8x16_16_sigmoid	rows=256; cols1=256; cols2=256; bias_prec=64	44236	2.96 (MACs/cyc)
matXvec_8x8_8_sigmoid	rows=256; cols1=256; cols2=256; bias_prec=8	44395	2.95 (MACs/cyc)
matXvec_8x8_8_sigmoid	rows=256; cols1=256; cols2=256; bias_prec=32	44203	2.97 (MACs/cyc)
matXvec_f32xf32_f32_sigmoid	rows=256; cols1=256; cols2=256; bias_prec=-1	90035	1.46 (MACs/cyc)
matXvec_16x16_16_tanh	rows=256; cols1=256; cols2=256; bias_prec=16	44270	2.96 (MACs/cyc)
matXvec_16x16_16_tanh	rows=256; cols1=256; cols2=256; bias_prec=64	44012	2.98 (MACs/cyc)
matXvec_8x16_16_tanh	rows=256; cols1=256; cols2=256; bias_prec=16	44462	2.95 (MACs/cyc)
matXvec_8x16_16_tanh	rows=256; cols1=256; cols2=256; bias_prec=64	44364	2.95 (MACs/cyc)
matXvec_8x8_8_tanh	rows=256; cols1=256; cols2=256; bias_prec=8	44657	2.94 (MACs/cyc)
matXvec_8x8_8_tanh	rows=256; cols1=256; cols2=256; bias_prec=32	44465	2.95 (MACs/cyc)
matXvec_f32xf32_f32_tanh	rows=256; cols1=256; cols2=256; bias_prec=-1	91231	1.44 (MACs/cyc)
matXvec_16x16_16	rows=256; cols1=256; cols2=256; bias_prec=16	40920	3.20 (MACs/cyc)
matXvec_16x16_32	rows=256; cols1=256; cols2=256; bias_prec=16	40952	3.20 (MACs/cyc)
matXvec_16x16_64	rows=256; cols1=256; cols2=256; bias_prec=16	40812	3.21 (MACs/cyc)
matXvec_8x16_16	rows=256; cols1=256; cols2=256; bias_prec=16	41145	3.19 (MACs/cyc)
matXvec_8x16_32	rows=256; cols1=256; cols2=256; bias_prec=16	41050	3.19 (MACs/cyc)
matXvec_8x16_64	rows=256; cols1=256; cols2=256; bias_prec=16	40847	3.21 (MACs/cyc)
matXvec_8x8_8	rows=256; cols1=256; cols2=256; bias_prec=8	41115	3.19 (MACs/cyc)
matXvec_8x8_16	rows=256; cols1=256; cols2=256; bias_prec=8	40826	3.21 (MACs/cyc)
matXvec_8x8_32	rows=256; cols1=256; cols2=256; bias_prec=8	40761	3.22 (MACs/cyc)
matXvec_f32xf32_f32	rows=256; cols1=256; cols2=256; bias_prec=-1	85628	1.53 (MACs/cyc)
matXvec_asym8xasym8_asym8	rows=256; cols1=256; cols2=256; bias_prec=32	78716	1.67 (MACs/cyc)
matXvec_sym8sxasym8s_asym8s	rows=256; cols1=256; cols2=256; bias_prec=32	399970	0.33 (MACs/cyc)
matXvec_batch_16x16_64	rows=1024; cols1=1024; bias_prec=16; vec_count=4	1080930	3.88 (MACs/cyc)
matXvec_batch_8x16_64	rows=1024; cols1=1024; bias_prec=16; vec_count=4	1086055	3.86 (MACs/cyc)



Low Level Kernel	Parameters	Average Cycles	Performance Metric
matXvec_batch_8x8_32	rows=1024; cols1=1024; bias_prec=8; vec_count=4	1087585	3.86 (MACs/cyc)
matXvec_batch_f32xf32_f32	rows=1024; cols1=1024; bias_prec=-1; vec_count=4	2131023	1.97 (MACs/cyc)
matXvec_batch_asym8xasym8_asym8	rows=1024; cols1=1024; bias_prec=32; vec_count=4	1882733	2.23 (MACs/cyc)
matXvec_batch_16x16_64	rows=256; cols1=256; bias_prec=16; vec_count=4	74082	3.54 (MACs/cyc)
matXvec_batch_8x16_64	rows=256; cols1=256; bias_prec=16; vec_count=4	75367	3.48 (MACs/cyc)
matXvec_batch_8x8_32	rows=256; cols1=256; bias_prec=8; vec_count=4	75716	3.46 (MACs/cyc)
matXvec_batch_f32xf32_f32	rows=256; cols1=256; bias_prec=-1; vec_count=4	139983	1.87 (MACs/cyc)
matXvec_batch_asym8xasym8_asym8	rows=256; cols1=256; bias_prec=32; vec_count=4	127085	2.06 (MACs/cyc)
matXvec_acc_batch_sym8sx8_asym16s	rows=1024; cols1=1024; bias_prec=32; vec_count=4	2008909	2.09 (MACs/cyc)
matXvec_acc_batch_sym8sx8_asym16s	rows=256; cols1=256; bias_prec=32; vec_count=4	142029	1.85 (MACs/cyc)
fully_connected_16x16_16	rows=1024	304759	3.44 (MACs/cyc)
fully_connected_8x16_16	rows=1024	305529	3.43 (MACs/cyc)
fully_connected_8x8_8	rows=1024	306297	3.42 (MACs/cyc)
fully_connected_f32	rows=1024	662171	1.58 (MACs/cyc)
fully_connected_asym8xasym8_asym8	rows=1024	604079	1.74 (MACs/cyc)
fully_connected_16x16_16	rows=256	20983	3.12 (MACs/cyc)
fully_connected_8x16_16	rows=256	21177	3.09 (MACs/cyc)
fully_connected_8x8_8	rows=256	21369	3.07 (MACs/cyc)
fully_connected_f32	rows=256	42779	1.53 (MACs/cyc)
fully_connected_asym8xasym8_asym8	rows=256	40559	1.62 (MACs/cyc)
fully_connected_sym8sxasym8s_asym8s	rows=1024	325172	3.22 (MACs/cyc)
fully_connected_sym8sxasym8s_asym8s	rows=256	23060	2.84 (MACs/cyc)
sigmoid_32x32	N=200	2478	12.39 (cyc/point)
tanh_32x32	N=200	2576	12.88 (cyc/point)
relu_8x8	N=200	164	0.82 (cyc/point)

Low Level Kernel	Parameters	Average Cycles	Performance Metric
relu_std_8x8	N=200	160	0.80 (cyc/point)
relu_16x16	N=200	161	0.81 (cyc/point)
relu_std_16x16	N=200	157	0.79 (cyc/point)
relu_32x32	N=200	301	1.50 (cyc/point)
relu1_32x32	N=200	300	1.50 (cyc/point)
relu6_32x32	N=200	300	1.50 (cyc/point)
relu_std_32x32	N=200	301	1.50 (cyc/point)
softmax_32x32	N=200	1932	9.66 (cyc/point)
sigmoid_32x16	N=200	2567	12.84 (cyc/point)
tanh_32x16	N=200	2665	13.32 (cyc/point)
sigmoid_16x16	N=200	3028	15.14 (cyc/point)
tanh_16x16	N=200	3242	16.21 (cyc/point)
sigmoid_f32xf32	N=200	3502	17.51 (cyc/point)
tanh_f32xf32	N=200	4444	22.22 (cyc/point)
relu_f32xf32	N=200	246	1.23 (cyc/point)
relu1_f32xf32	N=200	245	1.23 (cyc/point)
relu6_f32xf32	N=200	247	1.24 (cyc/point)
relu_std_f32xf32	N=200	231	1.16 (cyc/point)
softmax_f32xf32	N=200	2806	14.03 (cyc/point)
sigmoid_asym8xasym8	N=200	9014	45.07 (cyc/point)
softmax_asym8xasym8	N=200	5925	29.63 (cyc/point)
relu_asym8xasym8	N=200	149	0.74 (cyc/point)
softmax_asym8sx16	N=200	5696	28.48 (cyc/point)
softmax_asym8sxasym8s	N=200	5977	29.89 (cyc/point)
sigmoid_asym8sxasym8s	N=200	8479	42.40 (cyc/point)

Low Level Kernel	Parameters	Average Cycles	Performance Metric
tanh_asym8sxasym8s	N=200	10001	50.01 (cyc/point)
leaky_relu_asym8sxasym8s	N=200	969	4.84 (cyc/point)
prelu_asym8sxasym8s	N=200	1493	7.46 (cyc/point)
hard_swish_asym8sxasym8s	N=200	937	4.68 (cyc/point)
elm_floor_f32	N=200	336	1.68 (cyc/point)
elm_sine_f32	N=200	3806	19.03 (cyc/point)
elm_cosine_f32	N=200	3798	18.99 (cyc/point)
elm_logn_f32	N=200	2795	13.97 (cyc/point)
elm_abs_f32	N=200	238	1.19 (cyc/point)
elm_ceil_f32	N=200	338	1.69 (cyc/point)
elm_round_f32	N=200	340	1.70 (cyc/point)
elm_neg_f32	N=200	238	1.19 (cyc/point)
elm_square_f32	N=200	239	1.20 (cyc/point)
elm_sqrt_f32	N=200	1650	8.25 (cyc/point)
elm_rsqrt_f32	N=200	849	4.25 (cyc/point)
elm_add_f32	N=200	343	1.72 (cyc/point)
memset_f32	N=2048	1063	0.52 (cyc/point)
elm_dequantize_asym8s_f32	N=200	297	1.49 (cyc/point)
elm_requantize_asym16s_asym8s	N=200	908	4.54 (cyc/point)
l2_norm_f32	num_elms=1024	4981	0.41 (OPs/cyc)
l2_norm_f32	num_elms=256	1333	0.38 (OPs/cyc)
depth_to_space_8	input_height=11	12548	0.72 (cyc/point)
depth_to_space_8	input_height=40	40867	1.00 (cyc/point)
space_to_depth_8	input_height=33	12610	0.72 (cyc/point)
space_to_depth_8	input_height=80	40703	0.99 (cyc/point)

Low Level Kernel	Parameters	Average Cycles	Performance Metric
pad_8	input_shape=1 3 3 256 pad_shape=2 4 output_shape=1 5 5 256 pad_values=0 0 1 1 1 1 0 0	1786	0.28 (cyc/point)
pad_8	input_shape=1 128 128 8 pad_shape=2 4 output_shape=1 138 138 8 pad_values=0 0 7 3 3 7 0 0	56847	0.37 (cyc/point)
pad_8	input_shape=1 32 40 32 pad_shape=2 4 output_shape=1 32 40 32 pad_values=0 0 0 0 0 0 0 0	15665	0.38 (cyc/point)
batch_to_space_nd_8	input_shape=18 6 8 32 block_sizes=3 2 crop_sizes=0 0 0 0 output_shape=3 18 16 32	30212	1.09 (cyc/point)
batch_to_space_nd_8	input_shape=128 6 128 block_sizes=2 3 crop_sizes=0 0 0 0 output_shape=2 16 18 128	48487	0.66 (cyc/point)
space_to_batch_nd_8	input_shape=3 18 16 32 block_sizes=3 2 pad_sizes=0 0 0 0 output_shape=18 6 8 32	33735	1.22 (cyc/point)
space_to_batch_nd_8	input_shape=2 16 18 128 block_sizes=2 3 pad_sizes=0 0 0 0 output_shape=12 8 6 128	57770	0.78 (cyc/point)
conv2d_depth_8x8	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	1644689	0.64 (MACs/cyc)
conv2d_point_8x8	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	441032	1.63 (MACs/cyc)
conv2d_depth_8x8	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	11609867	1.37 (MACs/cyc)
conv2d_point_8x8	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	1966261	0.53 (MACs/cyc)
conv2d_depth_8x8_nhwc	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	1188880	0.88 (MACs/cyc)

Low Level Kernel	Parameters	Average Cycles	Performance Metric
conv2d_point_8x8_nhwc	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	441033	1.63 (MACs/cyc)
conv2d_depth_8x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	1634243	0.64 (MACs/cyc)
conv2d_point_8x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	422327	1.70 (MACs/cyc)
conv2d_depth_8x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	11660605	1.36 (MACs/cyc)
conv2d_point_8x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	1867972	0.56 (MACs/cyc)
conv2d_depth_8x16_nhwc	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	1510885	0.69 (MACs/cyc)
conv2d_point_8x16_nhwc	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	422328	1.70 (MACs/cyc)
conv2d_depth_16x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	1543749	0.68 (MACs/cyc)
conv2d_point_16x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	413905	1.74 (MACs/cyc)
conv2d_depth_16x16	input_height=128; input_width=128; input_channels=8; kernel_height=11;	11160225	1.42 (MACs/cyc)

Low Level Kernel	Parameters	Average Cycles	Performance Metric
	kernel_width=11; out_channels=8; out_height=128; out_width=128		
conv2d_point_16x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	1818822	0.58 (MACs/cyc)
conv2d_depth_16x16_nhwc	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	1398309	0.75 (MACs/cyc)
conv2d_point_16x16_nhwc	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	413908	1.74 (MACs/cyc)
conv2d_depth_f32xf32	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	2006054	0.52 (MACs/cyc)
conv2d_point_f32xf32	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	595933	1.21 (MACs/cyc)
conv2d_depth_f32xf32	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	15922406	1.00 (MACs/cyc)
conv2d_point_f32xf32	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	2105513	0.50 (MACs/cyc)
conv2d_depth_f32xf32_nhwc	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	1266572	0.83 (MACs/cyc)
conv2d_point_f32xf32_nhwc	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	595934	1.21 (MACs/cyc)

Low Level Kernel	Parameters	Average Cycles	Performance Metric
conv2d_depth_asym8xasym8	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	1798766	0.58 (MACs/cyc)
conv2d_point_asym8xasym8	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	603930	1.19 (MACs/cyc)
conv2d_depth_asym8xasym8	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	12852024	1.23 (MACs/cyc)
conv2d_point_asym8xasym8	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	2359511	0.44 (MACs/cyc)
conv2d_depth_asym8xasym8_nhwc	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	1419466	0.74 (MACs/cyc)
conv2d_point_asym8xasym8_nhwc	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	603932	1.19 (MACs/cyc)
conv2d_depth_sym8sxasym8s	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	1860910	0.56 (MACs/cyc)
conv2d_point_sym8sxasym8s	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	607285	1.18 (MACs/cyc)
conv2d_depth_sym8sxasym8s	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	13257736	1.20 (MACs/cyc)
conv2d_point_sym8sxasym8s	input_height=128; input_width=128; input_channels=8;	2523517	0.42 (MACs/cyc)

Low Level Kernel	Parameters	Average Cycles	Performance Metric
	kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128		
conv2d_depth_sym8sxasym8s_nhwc	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	1369944	0.77 (MACs/cyc)
conv2d_point_sym8sxasym8s_nhwc	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	607287	1.18 (MACs/cyc)
conv1d_std_8x8	input_height=32; input_width=40; input_channels=32; kernel_height=7; out_channels=24; out_height=26; ;	1770230	3.16 (MACs/cyc)
conv1d_std_8x8	input_height=128; input_width=128; input_channels=8; kernel_height=11; out_channels=8; out_height=128; ;	8727888	1.32 (MACs/cyc)
conv1d_std_8x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; out_channels=24; out_height=26; ;	1608740	3.48 (MACs/cyc)
conv1d_std_8x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; out_channels=8; out_height=128; ;	5882064	1.96 (MACs/cyc)
conv1d_std_16x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; out_channels=24; out_height=26; ;	1607570	3.48 (MACs/cyc)
conv1d_std_16x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; out_channels=8; out_height=128; ;	5878864	1.96 (MACs/cyc)
conv1d_std_f32xf32	input_height=32; input_width=40; input_channels=32; kernel_height=7; out_channels=24; out_height=26; ;	3200622	1.75 (MACs/cyc)
conv1d_std_f32xf32	input_height=128; input_width=128; input_channels=8; kernel_height=11; out_channels=8; out_height=128; ;	11723076	0.98 (MACs/cyc)
conv1d_std_asym8xasym8	input_height=32; input_width=40; input_channels=32;	3172333	1.76 (MACs/cyc)



Low Level Kernel	Parameters	Average Cycles	Performance Metric
	kernel_height=7; out_channels=24; out_height=26; ;		
conv1d_std_asym8xasym8	input_height=128; input_width=128; input_channels=8; kernel_height=11; out_channels=8; out_height=128; ;	6569832	1.76 (MACs/cyc)
conv2d_std_8x8	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	8627218	2.92 (MACs/cyc)
conv2d_std_8x8	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	42834936	2.96 (MACs/cyc)
conv2d_std_8x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	6913418	3.64 (MACs/cyc)
conv2d_std_8x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	34838279	3.64 (MACs/cyc)
conv2d_std_16x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	6915938	3.64 (MACs/cyc)
conv2d_std_16x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	34853895	3.64 (MACs/cyc)
conv2d_std_f32xf32	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	13408097	1.88 (MACs/cyc)
conv2d_std_f32xf32	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	66141661	1.92 (MACs/cyc)

Low Level Kernel	Parameters	Average Cycles	Performance Metric
conv2d_std_asym8xasym8	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	11842368	2.12 (MACs/cyc)
conv2d_std_asym8xasym8	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	59067354	2.15 (MACs/cyc)
conv2d_std_sym8sxasym8s	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	8221645	3.06 (MACs/cyc)
conv2d_std_sym8sxasym8s	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	42698223	2.97 (MACs/cyc)
avgpool_8	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	88720	0.23 (OPs/cyc)
avgpool_8	input_height=384; input_width=128; input_channels=8; kernel_height=12; kernel_width=4; out_height=43; out_width=43	1021888	0.71 (OPs/cyc)
avgpool_8_nhwc	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	34679	0.58 (OPs/cyc)
avgpool_16	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	73109	0.27 (OPs/cyc)
avgpool_16	input_height=384; input_width=128; input_channels=8; kernel_height=12; kernel_width=4; out_height=43; out_width=43	945461	0.77 (OPs/cyc)
avgpool_16_nhwc	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	38263	0.52 (OPs/cyc)
avgpool_f32	input_height=10; input_width=49; input_channels=32;	85887	0.23 (OPs/cyc)

Low Level Kernel	Parameters	Average Cycles	Performance Metric
	kernel_height=2; kernel_width=2; out_height=5; out_width=25		
avgpool_f32	input_height=384; input_width=128; input_channels=8; kernel_height=12; kernel_width=4; out_height=43; out_width=43	1094677	0.66 (OPs/cyc)
avgpool_f32_nhwc	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	45034	0.44 (OPs/cyc)
maxpool_8	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	52216	0.31 (OPs/cyc)
maxpool_8	input_height=384; input_width=128; input_channels=8; kernel_height=12; kernel_width=4; out_height=43; out_width=43	612192	1.16 (OPs/cyc)
maxpool_8_nhwc	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	26661	0.60 (OPs/cyc)
maxpool_16	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	38668	0.41 (OPs/cyc)
maxpool_16	input_height=384; input_width=128; input_channels=8; kernel_height=12; kernel_width=4; out_height=43; out_width=43	398084	1.78 (OPs/cyc)
maxpool_16_nhwc	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	17671	0.91 (OPs/cyc)
maxpool_f32	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	53034	0.30 (OPs/cyc)
maxpool_f32	input_height=384; input_width=128; input_channels=8; kernel_height=12; kernel_width=4; out_height=43; out_width=43	721674	0.98 (OPs/cyc)
maxpool_f32_nhwc	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	29385	0.54 (OPs/cyc)

Low Level Kernel	Parameters	Average Cycles	Performance Metric
avgpool_asym8	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	80501	0.25 (OPs/cyc)
avgpool_asym8	input_height=384; input_width=128; input_channels=8; kernel_height=12; kernel_width=4; out_height=43; out_width=43	997067	0.73 (OPs/cyc)
maxpool_asym8	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	48363	0.33 (OPs/cyc)
maxpool_asym8	input_height=384; input_width=128; input_channels=8; kernel_height=12; kernel_width=4; out_height=43; out_width=43	577779	1.23 (OPs/cyc)
avgpool_asym8_nhwc	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	30599	0.65 (OPs/cyc)
maxpool_asym8_nhwc	input_height=10; input_width=49; input_channels=32; kernel_height=2; kernel_width=2; out_height=5; out_width=25	25784	0.62 (OPs/cyc)

### 1.3.5 Timings – Layers

Table 1-5 provides the average cycles taken on HiFi DSP to generate one output shape by processing one input shape.

Table 1-5 HiFi Average Cycles to Generate One Output Shape

Layer	Parameters	Average Cycles	Performance Metric
gru_16x16	in_feats=32; out_feats=256	84978	NA
gru_16x16	in_feats=36; out_feats=256	85842	NA
gru_16x16	in_feats=256; out_feats=256	133362	NA
gru_16x16	in_feats=128; out_feats=1024	1084146	NA
gru_16x16	in_feats=1024; out_feats=1024	1858290	NA
gru_8x16	in_feats=32; out_feats=256	85555	NA
gru_8x16	in_feats=36; out_feats=256	86419	NA

Layer	Parameters	Average Cycles	Performance Metric
gru_8x16	in_feats=256; out_feats=256	133939	NA
gru_8x16	in_feats=128; out_feats=1024	1086451	NA
gru_8x16	in_feats=1024; out_feats=1024	1860595	NA
lstm_16x16	in_feats=32; out_feats=256	116481	NA
lstm_16x16	in_feats=36; out_feats=256	117633	NA
lstm_16x16	in_feats=256; out_feats=256	180993	NA
lstm_16x16	in_feats=128; out_feats=1024	1458273	NA
lstm_16x16	in_feats=1024; out_feats=1024	2490465	NA
lstm_8x16	in_feats=32; out_feats=256	117251	NA
lstm_8x16	in_feats=36; out_feats=256	118403	NA
lstm_8x16	in_feats=256; out_feats=256	181763	NA
lstm_8x16	in_feats=128; out_feats=1024	1461347	NA
lstm_8x16	in_feats=1024; out_feats=1024	2493539	NA
cnn_conv2d_depth_8x8	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	2086019	0.85 (MACs/cyc)
cnn_conv2d_depth_8x8	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	13576421	1.25 (MACs/cyc)
cnn_conv2d_depth_8x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	2056872	0.86 (MACs/cyc)
cnn_conv2d_depth_8x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	13528874	1.25 (MACs/cyc)
cnn_conv2d_depth_16x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	1957955	0.90 (MACs/cyc)
cnn_conv2d_depth_16x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	12979343	1.30 (MACs/cyc)
cnn_conv2d_depth_f32xf32	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	2602284	0.68 (MACs/cyc)
cnn_conv2d_depth_f32xf32	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	18028216	0.94 (MACs/cyc)

Layer	Parameters	Average Cycles	Performance Metric
cnn_conv1d_std_8x8	input_height=32; input_width=40; input_channels=32; kernel_height=7; out_channels=24; out_height=26; ;	1770385	3.16 (MACs/cyc)
cnn_conv1d_std_8x8	input_height=128; input_width=128; input_channels=8; kernel_height=11; out_channels=8; out_height=128; ;	8728043	1.32 (MACs/cyc)
cnn_conv1d_std_8x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; out_channels=24; out_height=26; ;	1608898	3.48 (MACs/cyc)
cnn_conv1d_std_8x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; out_channels=8; out_height=128; ;	5882222	1.96 (MACs/cyc)
cnn_conv1d_std_16x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; out_channels=24; out_height=26; ;	1607728	3.48 (MACs/cyc)
cnn_conv1d_std_16x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; out_channels=8; out_height=128; ;	5879022	1.96 (MACs/cyc)
cnn_conv1d_std_f32xf32	input_height=32; input_width=40; input_channels=32; kernel_height=7; out_channels=24; out_height=26; ;	3200775	1.75 (MACs/cyc)
cnn_conv1d_std_f32xf32	input_height=128; input_width=128; input_channels=8; kernel_height=11; out_channels=8; out_height=128; ;	11723229	0.98 (MACs/cyc)
cnn_conv2d_std_8x8	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	8627374	2.92 (MACs/cyc)
cnn_conv2d_std_8x8	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	42835092	2.96 (MACs/cyc)
cnn_conv2d_std_8x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	6913574	3.64 (MACs/cyc)
cnn_conv2d_std_8x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	34838435	3.64 (MACs/cyc)
cnn_conv2d_std_16x16	input_height=32; input_width=40; input_channels=32; kernel_height=7; kernel_width=5; out_channels=24; out_height=26; out_width=36	6916095	3.64 (MACs/cyc)
cnn_conv2d_std_16x16	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	34854052	3.64 (MACs/cyc)
cnn_conv2d_std_f32xf32	input_height=32; input_width=40; input_channels=32; kernel_height=7;	13408245	1.88 (MACs/cyc)

Layer	Parameters	Average Cycles	Performance Metric
	kernel_width=5; out_channels=24; out_height=26; out_width=36		
cnn_conv2d_std_f32xf32	input_height=128; input_width=128; input_channels=8; kernel_height=11; kernel_width=11; out_channels=8; out_height=128; out_width=128	66141809	1.92 (MACs/cyc)

**Note** Performance specification measurements are carried out on a cycle-accurate simulator assuming an ideal memory system, i.e., one with zero memory wait states. This is equivalent to running with all code and data in local memories or using an infinite-size, pre-filled cache model.

**Note** The measurements are carried out with Xtensa RI.7 tool chain and AE\_HiFi1\_LE5\_FP\_XC DSP.

## 2. Generic HiFi NN Layer API

**Note** This section explains an API standard that is evolving. The APIs may undergo some changes in future versions.

This section describes the API that is common to all the HiFi NN layers. The API facilitates any layer instance that works in the overall method shown in Figure 2-1.

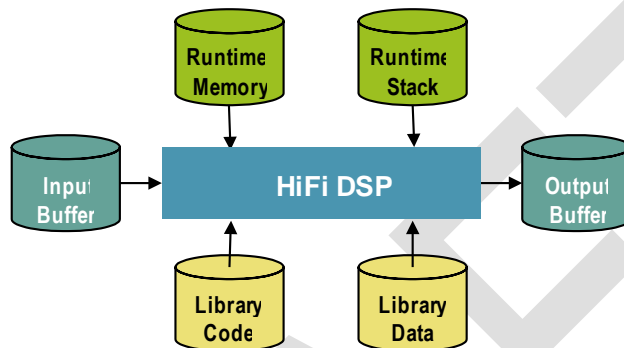


Figure 2-1 HiFi NN Layer Interfaces

All the buffers, input, output, weights and biases are described as shapes. Section 2.1 explains the shape structure.

Section 2.2 discusses all the types of runtime memory required by the layer instances. There is no state information held in static memory, therefore a single thread can perform time division processing of multiple layer instances. Additionally, multiple threads can perform concurrent layer instance processing.

The output from one instance can be fed as input to the next instance if the precision and the dimension matches.

The data types, structures, and error codes explained in this section are declared/defined in `xa_nnlib_standard.h`. By default, the API header file of each layer will include this header file. The application need not include this file.

### 2.1 Shape

The shapes are used to describe any buffer used in the NN library. The structure `xa_nnlib_shape_t` is defined in `xa_nnlib_standard.h`. The shape can be vector, matrix, or cube.

- Vector is a one-dimensional shape specified by length.
- Matrix is a two-dimensional shape specified by rows, columns, and row\_offset. This assumes that the elements in a row are stored at consecutive addresses in memory.



- Cube is a three-dimensional shape specified by height, width, depth, height\_offset, width\_offset, and depth offset. Cube supports the following shape types:
  - `SHAPE_CUBE_DWH_T`  
This assumes that elements are stored in depth, width, and height order; that is, elements with the same height and width indices are stored consecutively.
  - `SHAPE_CUBE_WHD_T`  
This assumes that elements are stored in width, height, and depth order; that is, elements with the same height and depth are stored consecutively.

Figure 2-2 shows the dimension variables of matrix and cube shapes.




Figure 2-2 Matrix and Cube (`SHAPE_CUBE_DWH_T`) Shape Representation

## 2.2 Memory Management

The HiFi NN layer API supports a flexible memory scheme and a simple interface that eases the integration into the final application. The API allows the layers to request the required memory for their operations during runtime.

The runtime memory requirement consists primarily of the scratch and persistent memory. The components also require an input buffer and output buffer for the passing of data into and out of the layer.

### 2.2.1 API Handle / Persistent Memory

The layer API stores persistent state information in a structure that is referenced via an opaque handle. The handle is passed by the application for each API call. This object contains all state and history information that is maintained from one-layer frame invocation to the next within the same thread or instance. The layers expect that the contents of the persistent memory be unchanged by the system apart from the layer itself for the complete lifetime of the layer.

### 2.2.2 Scratch Memory

This is the temporary buffer used by the layer during a single frame processing call. The contents of this memory region should not be changed if the actual layer execution process is active; that is, if the thread running the layer is inside any API call. This region can be used freely by the system between successive calls to the layer.

### 2.2.3 Weights and Biases Memory

The weights or coefficients and biases should be managed by the application, and memory should not be requested by the API. If the design requires DMA access from or to the internal memory for better performance, a ping-pong or circular buffer is allocated as part of the scratch into which the weights, biases, input, and output are copied using DMA. If required, these memories can also be persistent.

### 2.2.4 Input Buffer

This is the buffer from which the layer reads the input. This buffer must be made available for the layer before its execution call. The input buffer should have an associated shape information to describe the input data format. The input buffer pointer can be changed by the application between calls to the layer, but shape information cannot be changed. This allows the layer to read directly from the output of another layer.

### 2.2.5 Output Buffer

This is the buffer to which the layer writes the output. This buffer must be made available for the layer before its execution call. The output buffer should have an associated shape information to which the layer can describe the output data format. The output buffer pointer can be changed by the application between calls to the layer. This allows the layer to write directly to the input of another layer.

## 2.3 Generic API Errors

Layer API functions return an error code of type `Int32`, which is of type `signed int`. The format of the error codes is defined in the following table.

31	30 – 27	26–12	11 – 7	6 – 0
Fatal	Class	Reserved	Component	Sub code

The errors that can be returned from the API are subdivided into those that are fatal, which require resetting the layer; and those that are nonfatal and are provided for information to the application.

The class of an error can be API, Config, or Execution. The API category errors are concerned with the incorrect use of the API. The Config errors are produced when the layer parameters are incorrect or outside the supported usage. The Execution errors are returned after a call to the main process and indicate situations that have arisen due to the input data.

### 2.3.1 Common API Errors

The following errors are fatal and should not be encountered during normal application operation. They signal that a serious error has occurred in the application that is calling the layer.

- `XA_NNLIB_FATAL_MEM_ALLOC`  
At least one of the pointers passed into the API function is NULL.
- `XA_NNLIB_FATAL_MEM_ALIGN`  
At least one of the pointers passed into the API function is not properly aligned.
- `XA_NNLIB_FATAL_INVALID_SHAPE`  
At least one of the shapes passed to the API function is invalid.

## 2.4 C Language API

An overview of the NN layer flow is shown in Figure 2-3. The NN layer API consists of query, initialization, and execution functions.

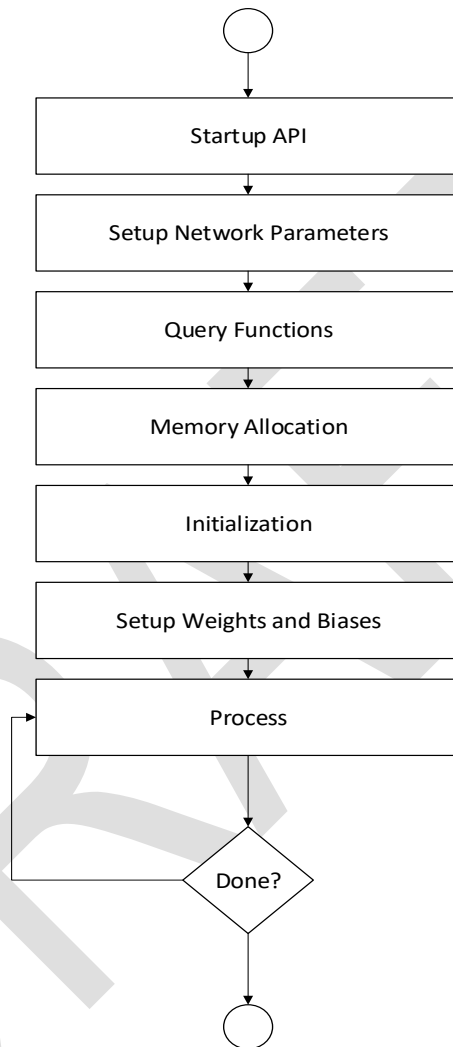


Figure 2-3 NN Layer Flow Overview

## 2.4.1 Startup Functions

The API startup functions shown in Table 2-1 get the various identification strings from the component library. They are for information only and their usage is optional. These functions do not take any input arguments and return `const char *`.

Table 2-1 Library Identification Functions

Function	Description
<code>xa_nnlib_get_lib_name_string</code>	Get the name of the library.
<code>xa_nnlib_get_lib_version_string</code>	Get the version of the library.
<code>xa_nnlib_get_lib_api_version_string</code>	Get the version of the API.

### Example

```
const char *name = xa_nnlib_get_lib_name_string();  
const char *ver = xa_nnlib_get_lib_version_string();  
const char *aver = xa_nnlib_get_lib_api_version_string();
```

### Errors

- None

## 2.4.2 Query Functions

The query functions are used in the startup and the memory allocation stages to obtain information about the memory requirements of the library.

Following is the naming convention for query functions:

```
xa_nnl-lib_<layer>_get_{persistent | scratch}_<placement>
```

Where:

<layer> indicates the module name (such as gru).

<placement> specifies fast or slow.

## 2.4.3 Initialization Functions

The initialization functions are used to reset the layer to its initial state. Because the layers are fully reentrant, the application can initialize the layer multiple times.

Following is the naming convention for initialization functions:

```
xa_nnl-lib_<layer>_init
```

## 2.4.4 Execution Functions

The execution functions are used to generate the output shape by processing one input shape.

Following is the naming convention for execution functions:

```
xa_nnl-lib_<layer>_process
```

## 3. HiFi 1 NN Library – Low-Level Kernels

This section explains the low-level kernels provided in the NN library. All the low-level kernels have a generic, simple interface.

The NN library is a single archive containing all low-level kernels and layers implementations. The following sections explain each low-level kernel in detail.

### 3.1 Matrix X Vector Multiplication Kernels

#### 3.1.1 Matrix X Vector Kernels

##### Description

These kernels perform the dual matXvec operation with bias addition; that is,  $z = \text{mat1} * \text{vec1} + \text{mat2} * \text{vec2} + \text{bias}$ . The column dimension of `mat1` must match the row dimension of `vec1` and similarly for `mat2`, `vec2`. Bias and resulting output vector `z` have as many rows as `mat1` and `mat2`.

`bias_shift` and `acc_shift` arguments are provided in the kernel API to adjust Q format of bias and output, respectively. Both `bias_shift` and `acc_shift` can be either positive or negative, where positive value denotes a left shift and negative value denotes a right shift.

`bias_shift` is the shift in number of bits applied to the bias to make it in the same Q format as matXvec multiplication – accumulation result. `acc_shift` is the shift in number of bits applied to the accumulator to obtain the output in desired Q format.

Note, `acc_shift` and `bias_shift` arguments are not relevant in case of floating-point kernels and quantized 8-bit kernels.

`row_stride1` and `row_stride2` arguments are provided in kernel API for row offsets of `mat1` and `mat2`, respectively. Note, input matrices are expected to be appropriately padded in case of `row_stride > cols`.

For conversion from higher precision accumulator to lower precision output, symmetric rounding is used.

The arguments, `mat1_zero_bias`, `mat2_zero_bias`, `vec1_zero_bias`, `vec2_zero_bias`, are provided to convert the quantized 8-bit inputs into their real values and perform matXvec operation. The `out_zero_bias`, `out_multiplier` and `out_shift` values are used to quantize real values of output back to 8-bit.

Function variants available are `xa_nn_matXvec_[p]x[q]_[r]`, where:

- `[p]`: Matrix precision in bits
- `[q]`: Vector precision in bits

- [r]: Output precision in bits

## Precision

There are twelve variants available:

Type	Description
16x16_16	16-bit matrix inputs, 16-bit vector inputs, 16-bit output
16x16_32	16-bit matrix inputs, 16-bit vector inputs, 32-bit output
16x16_64	16-bit matrix inputs, 16-bit vector inputs, 64-bit output
8x16_16	8-bit matrix inputs, 16-bit vector inputs, 16-bit output
8x16_32	8-bit matrix inputs, 16-bit vector inputs, 32-bit output
8x16_64	8-bit matrix inputs, 16-bit vector inputs, 64-bit output
8x8_8	8-bit matrix inputs, 8-bit vector inputs, 8-bit output
8x8_16	8-bit matrix inputs, 8-bit vector inputs, 16-bit output
8x8_32	8-bit matrix inputs, 8-bit vector inputs, 32-bit output
f32xf32_f32	float32 matrix inputs, float32 vector inputs, float32 output
asym8uxasym8u_asym8u	asym8u matrix inputs, asym8u vector inputs, asym8u output
sym8sxsym8s_asym8s	sym8s matrix inputs, asym8s vector inputs, asym8s output

## Algorithm

$$z_n = 2^{acc-shift} \left( \sum_{m=0}^{cols1-1} mat1_{n,m} \cdot vec1_m + \sum_{m=0}^{cols2-1} mat2_{n,m} \cdot vec2_m + 2^{bias-shift} bias_n \right)$$

For floating-point and quantized 8-bit routines, `acc_shift=0` and `bias_shift=0`.

Thus,  $2^{acc-shift} = 2^{bias-shift} = 1$

## Prototype

```
WORD32 xa_nn_matXvec_16x16_16
(WORD16 * p_out,          WORD16 * p_mat1,          WORD16 * p_mat2,
 WORD16 * p_vec1,         WORD16 * p_vec2,          WORD16 * p_bias,
 WORD32 rows,             WORD32 cols1,             WORD32 cols2,
 WORD32 row_stride1,      WORD32 row_stride2,
 WORD32 acc_shift,        WORD32 bias_shift);
WORD32 xa_nn_matXvec_16x16_32
(WORD32 * p_out,          WORD16 * p_mat1,          WORD16 * p_mat2,
 WORD16 * p_vec1,         WORD16 * p_vec2,          WORD16 * p_bias,
 WORD32 rows,             WORD32 cols1,             WORD32 cols2,
 WORD32 row_stride1,      WORD32 row_stride2,
 WORD32 acc_shift,        WORD32 bias_shift);
WORD32 xa_nn_matXvec_16x16_64
(WORD64 * p_out,          WORD16 * p_mat1,          WORD16 * p_mat2,
 WORD16 * p_vec1,         WORD16 * p_vec2,          WORD16 * p_bias,
 WORD32 rows,             WORD32 cols1,             WORD32 cols2,
 WORD32 row_stride1,      WORD32 row_stride2,
 WORD32 acc_shift,        WORD32 bias_shift);
WORD32 xa_nn_matXvec_8x16_16
(WORD16 * p_out,          WORD8 * p_mat1,          WORD8 * p_mat2,
```



```

WORD16 * p_vec1,      WORD16 * p_vec2,      WORD16 * p_bias,
WORD32 rows,          WORD32 cols1,        WORD32 cols2,
WORD32 row_stride1,   WORD32 row_stride2,
WORD32 acc_shift,     WORD32 bias_shift);
WORD32 xa_nn_matXvec_8x16_32
(WORD32 * p_out,      WORD8 * p_mat1,      WORD8 * p_mat2,
WORD16 * p_vec1,     WORD16 * p_vec2,     WORD16 * p_bias,
WORD32 rows,         WORD32 cols1,        WORD32 cols2,
WORD32 row_stride1,   WORD32 row_stride2,
WORD32 acc_shift,     WORD32 bias_shift);
WORD32 xa_nn_matXvec_8x16_64
(WORD64 * p_out,      WORD8 * p_mat1,      WORD8 * p_mat2,
WORD16 * p_vec1,     WORD16 * p_vec2,     WORD16 * p_bias,
WORD32 rows,         WORD32 cols1,        WORD32 cols2,
WORD32 row_stride1,   WORD32 row_stride2,
WORD32 acc_shift,     WORD32 bias_shift);
WORD32 xa_nn_matXvec_8x8_8
(WORD8 * p_out,       WORD8 * p_mat1,      WORD8 * p_mat2,
WORD8 * p_vec1,       WORD8 * p_vec2,      WORD8 * p_bias,
WORD32 rows,          WORD32 cols1,        WORD32 cols2,
WORD32 row_stride1,   WORD32 row_stride2,
WORD32 acc_shift,     WORD32 bias_shift);
WORD32 xa_nn_matXvec_8x8_16
(WORD16 * p_out,      WORD8 * p_mat1,      WORD8 * p_mat2,
WORD8 * p_vec1,       WORD8 * p_vec2,      WORD8 * p_bias,
WORD32 rows,          WORD32 cols1,        WORD32 cols2,
WORD32 row_stride1,   WORD32 row_stride2,
WORD32 acc_shift,     WORD32 bias_shift);
WORD32 xa_nn_matXvec_8x8_32
(WORD32 * p_out,      WORD8 * p_mat1,      WORD8 * p_mat2,
WORD8 * p_vec1,       WORD8 * p_vec2,      WORD8 * p_bias,
WORD32 rows,          WORD32 cols1,        WORD32 cols2,
WORD32 row_stride1,   WORD32 row_stride2,
WORD32 acc_shift,     WORD32 bias_shift);
WORD32 xa_nn_matXvec_f32xf32_f32
(FLOAT32 * p_out,     const FLOAT32 * p_mat1, const FLOAT32 * p_mat2,
const FLOAT32 * p_vec1, const FLOAT32 * p_vec2, const FLOAT32 * p_bias,
WORD32 rows,          WORD32 cols1,        WORD32 cols2,
WORD32 row_stride1,   WORD32 row_stride2);
WORD32 xa_nn_matXvec_asym8uxasym8u_asym8u
(UWORD8 * p_out,      const UWORD8 * p_mat1, const UWORD8 * p_mat2,
const UWORD8 * p_vec1, const UWORD8 * p_vec2, const WORD32 * p_bias,
WORD32 rows,          WORD32 cols1,        WORD32 cols2,
WORD32 row_stride1,   WORD32 row_stride2, WORD32 mat1_zero_bias,
WORD32 mat2_zero_bias, WORD32 vec1_zero_bias, WORD32 vec2_zero_bias,
WORD32 out_multiplier, WORD32 out_shift,   WORD32 out_zero_bias);

```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
WORD16 *, WORD8 *, const FLOAT32 *, const UWORD8 *	p_mat1	rows*cols1	Input matrix 1, fixed, floating point, asym8u or sym8s

Type	Name	Size	Description
const WORD8 *			
WORD16 *, WORD8 *, const FLOAT32 *, const UWORD8 *, const WORD8 *	p_mat2	rows*cols2	Input matrix 2, fixed, floating point, asym8u or sym8s
WORD16 *, WORD8 *, const FLOAT32 *, const UWORD8 *, const WORD8 *	p_vec1	cols1*1	Input vector 1, fixed, floating point, asym8u or asym8s
WORD16 *, WORD8 *, const FLOAT32 *, const UWORD8 *, const WORD8 *	p_vec2	cols2*1	Input vector 2, fixed, floating point, asym8u or asym8s
WORD16 *, WORD8 *, const WORD32 *, const FLOAT32 *	p_bias	rows*1	Bias vector, fixed or floating point
WORD32	rows		Number of rows in matrix 1, 2 and bias
WORD32	cols1		Number of columns in matrix 1 and rows in vector 1
WORD32	cols2		Number of columns in matrix 2 and rows in vector 2
WORD32	row_stride1		Row offset of matrix 1
WORD32	row_stride2		Row offset of matrix 2
WORD32	acc_shift		Shift applied to accumulator
WORD32	bias_shift		Shift applied to bias
WORD32	mat1_zero_bias		Zero offset of matrix 1
WORD32	mat2_zero_bias		Zero offset of matrix 2
WORD32	vec1_zero_bias		Zero offset of vector 1
WORD32	vec2_zero_bias		Zero offset of vector 2
WORD32	out_multiplier		Multiplier value of output
WORD32	out_shift		Shift value of output
WORD32	out_zero_bias		Zero offset of output
<b>Output</b>			
WORD8 *, UWORD8 *, WORD16 *, WORD32 *, WORD64 *, FLOAT32 *	p_out	rows*1	Output, fixed, floating point, asym8u or asym8s.

## Returns

- 0: no error

- -1: error, invalid parameters

## Restrictions

Arguments	Restrictions
row_stridel, row_stride2, cols1, cols2	row_stride1 >= cols1 row_stride2 >= cols2
p_mat1, p_mat2, p_vec1, p_vec2, p_bias, p_out	Aligned on <size of one element> boundary Should not overlap
p_mat1, p_vec1, p_out	Cannot be NULL
acc_shift, bias_shift, out_shift	{-31, ..., 31}
mat1_zero_bias, mat2_zero_bias, vec1_zero_bias, vec2_zero_bias	{-255, ..., 0} for asym8u, {-127, ..., 127} for asym8s
out_multiplier	Greater than 0
out_zero_bias	{0, ..., 255} if out type is asym8u, {-128, ..., 127} if out type is asym8s

## 3.1.2 Fused (Activation) Matrix X Vector Kernels

### Description

These kernels perform the fused dual matXvec operation with an activation function i.e.  $z = \text{activation}(\text{mat1} * \text{vec1} + \text{mat2} * \text{vec2} + \text{bias})$ . The column dimension of `mat1` must match the row dimension of `vec1` and similarly for `mat2`, `vec2`. Bias and resulting output vector `z` have as many rows as `mat1` and `mat2`.

Intermediate output of  $(\text{mat1} * \text{vec1} + \text{mat2} * \text{vec2} + \text{bias})$  is stored in temporary memory provided by the `p_scratch` argument to kernel API. Activation function is applied on this intermediate output to get final output. Note, for fixed point kernels, the activation function always takes input in Q6.25 format.

`bias_shift` and `acc_shift` arguments are provided in kernel API to adjust Q format of bias and intermediate output respectively. Both `bias_shift` and `acc_shift` can be either positive or negative, where positive value denotes a left shift and negative value denotes a right shift.

`bias_shift` is the shift in number of bits applied to the bias to make it in the same Q format as matXvec multiplication – accumulation result. `acc_shift` is the shift in number of bits applied to the accumulator to obtain the intermediate output in Q6.25 format.

**Note:** `acc_shift` and `bias_shift` are not relevant in case of floating point kernels.

`row_stridel` and `row_stride2` arguments are provided in kernel API for row offsets of `mat1` and `mat2` respectively. Note, input matrices are expected to be appropriately padded in case of `row_stride > cols`.

For conversion from higher precision accumulator to lower precision output, symmetric rounding is used.

Function variants available are `xa_nn_matXvec_[p]x[q]_[r]_<activation>`, where:

- [p]: Matrix precision in bits
- [q]: Vector precision in bits
- [r]: Output precision in bits
- <activation>: activation tag 'tanh' or 'sigmoid'

## Precision

There are eight variants available:

Type	Description
16x16_16_tanh	16-bit matrix inputs, 16-bit vector inputs, 16-bit output with tanh activation function
16x16_16_sigmoid	16-bit matrix inputs, 16-bit vector inputs, 16-bit output with sigmoid activation function
8x16_16_tanh	8-bit matrix inputs, 16-bit vector inputs, 16-bit output with tanh activation function
8x16_16_sigmoid	8-bit matrix inputs, 16-bit vector inputs, 16-bit output with sigmoid activation function
8x8_8_tanh	8-bit matrix inputs, 8-bit vector inputs, 8-bit output with tanh activation
8x8_8_sigmoid	8-bit matrix inputs, 8-bit vector inputs, 8-bit output with sigmoid activation
f32xf32_f32_tanh	float32 matrix inputs, float32 vector inputs, float32 output with tanh activation
f32xf32_f32_sigmoid	float32 matrix inputs, float32 vector inputs, float32 output with sigmoid activation

## Algorithm

$$z_n = \text{activation} \left( 2^{\text{acc\_shift}} \left( \sum_{m=0}^{\text{cols1}-1} \text{mat1}_{n,m} \cdot \text{vec1}_m + \sum_{m=0}^{\text{cols2}-1} \text{mat2}_{n,m} \cdot \text{vec2}_m \right) + 2^{\text{bias\_shift}} \text{bias}_n \right), \quad n = 0, \dots, \text{rows} - 1$$

In case of floating point routine, acc\_shift=0 and bias\_shift=0.

Thus,  $2^{\text{acc\_shift}} = 2^{\text{bias\_shift}} = 1$

activation is tanh or sigmoid

## Prototype

```
WORD32 xa_nn_matXvec_16x16_16_tanh
(WORD16 * p_out,          WORD16 * p_mat1,          WORD16 * p_mat2,
 WORD16 * p_vec1,         WORD16 * p_vec2,          VOID * p_bias,
 WORD32 rows,             WORD32 cols1,              WORD32 cols2,
 WORD32 row_stride1,      WORD32 row_stride2,      WORD32 acc_shift,
 WORD32 bias_shift,       WORD32 bias_precision,    VOID * p_scratch);
```

```

WORD32 xa_nn_matXvec_16x16_16_sigmoid
(WORD16 * p_out,          WORD16 * p_mat1,          WORD16 * p_mat2,
 WORD16 * p_vec1,         WORD16 * p_vec2,          VOID * p_bias,
 WORD32 rows,             WORD32 cols1,             WORD32 cols2,
 WORD32 row_stridel,      WORD32 row_stride2,      WORD32 acc_shift,
 WORD32 bias_shift,       WORD32 bias_precision,    VOID * p_scratch);
WORD32 xa_nn_matXvec_8x16_16_tanh
(WORD16 * p_out,          WORD8 * p_mat1,          WORD8 * p_mat2,
 WORD16 * p_vec1,         WORD16 * p_vec2,          VOID * p_bias,
 WORD32 rows,             WORD32 cols1,             WORD32 cols2,
 WORD32 row_stridel,      WORD32 row_stride2,      WORD32 acc_shift,
 WORD32 bias_shift,       WORD32 bias_precision,    VOID * p_scratch);
WORD32 xa_nn_matXvec_8x16_16_sigmoid
(WORD16 * p_out,          WORD8 * p_mat1,          WORD8 * p_mat2,
 WORD16 * p_vec1,         WORD16 * p_vec2,          VOID * p_bias,
 WORD32 rows,             WORD32 cols1,             WORD32 cols2,
 WORD32 row_stridel,      WORD32 row_stride2,      WORD32 acc_shift,
 WORD32 bias_shift,       WORD32 bias_precision,    VOID * p_scratch);
WORD32 xa_nn_matXvec_8x8_8_tanh
(WORD8 * p_out,           WORD8 * p_mat1,          WORD8 * p_mat2,
 WORD8 * p_vec1,          WORD8 * p_vec2,          VOID * p_bias,
 WORD32 rows,             WORD32 cols1,             WORD32 cols2,
 WORD32 row_stridel,      WORD32 row_stride2,      WORD32 acc_shift,
 WORD32 bias_shift,       WORD32 bias_precision,    VOID * p_scratch);
WORD32 xa_nn_matXvec_8x8_8_sigmoid
(WORD8 * p_out,           WORD8 * p_mat1,          WORD8 * p_mat2,
 WORD8 * p_vec1,          WORD8 * p_vec2,          VOID * p_bias,
 WORD32 rows,             WORD32 cols1,             WORD32 cols2,
 WORD32 row_stridel,      WORD32 row_stride2,      WORD32 acc_shift,
 WORD32 bias_shift,       WORD32 bias_precision,    VOID * p_scratch);
WORD32 xa_nn_matXvec_f32xf32_f32_tanh
(FLOAT32 * p_out,         FLOAT32 * p_mat1,      FLOAT32 * p_mat2,
 FLOAT32 * p_vec1,        FLOAT32 * p_vec2,      FLOAT32 * p_bias,
 WORD32 rows,             WORD32 cols1,          WORD32 cols2,
 WORD32 row_stridel,      WORD32 row_stride2,    FLOAT32 * p_scratch);
WORD32 xa_nn_matXvec_f32xf32_f32_sigmoid
(FLOAT32 * p_out,         FLOAT32 * p_mat1,      FLOAT32 * p_mat2,
 FLOAT32 * p_vec1,        FLOAT32 * p_vec2,      FLOAT32 * p_bias,
 WORD32 rows,             WORD32 cols1,          WORD32 cols2,
 WORD32 row_stridel,      WORD32 row_stride2,    FLOAT32 * p_scratch);

```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
WORD16 *, WORD8 *, FLOAT32 *	p_mat1	rows*cols1	Input matrix 1, fixed or floating point
WORD16 *, WORD8 *, FLOAT32 *	p_mat2	rows*cols2	Input matrix 2, fixed or floating point
WORD16 *, WORD8 *, FLOAT32 *	p_vec1	cols1*1	Input vector 1, fixed or floating point
WORD16 *, WORD8 *, FLOAT32 *	p_vec2	cols2*1	Input vector 2, fixed or floating point
VOID *, FLOAT32 *	p_bias	rows*1	Bias vector, fixed or floating point
WORD32	rows		Number of rows in matrix 1,2, bias and output
WORD32	cols1		Number of columns in matrix 1 and rows in vector 1
WORD32	cols2		Number of columns in matrix 2 and rows in vector 2
WORD32	row_stride1		Row offset of matrix 1
WORD32	row_stride2		Row offset of matrix 2
WORD32	acc_shift		Shift applied to accumulator
WORD32	bias_shift		Shift applied to bias
WORD32	bias_precision		Precision of bias in bytes
<b>Output</b>			
WORD8 *, WORD16 *, FLOAT32 *	p_out	rows*1	Output, fixed (Q7, Q15) or floating point
<b>Temporary</b>			
VOID *, FLOAT32 *	p_scratch	rows*4	Scratch (temporary) memory pointer

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions

Arguments	Restrictions
row_stride1, row_stride2, cols1, cols2	Multiples of 4 (2 in case of floating point)
p_mat1, p_mat2, p_vec1, p_vec2, p_bias, p_out	Aligned on 16-byte boundary Should not overlap
p_mat1, p_vec1, p_bias, p_out	Cannot be NULL
acc_shift, bias_shift	{-31, ..., 31}
bias_precision	{-1, 8, 16, 32, 64} (-1 in case of floating point)

### 3.1.3 Matrix X Vector Batch Kernels

#### Description

These kernels perform the operation of multiplication of a single matrix with a series of vectors along with bias addition; that is,  $z_i = \text{mat1} * \text{vec1}_i + \text{bias}$ . These kernels can also be viewed as matrix X matrix-transpose multiplication kernels. The column dimension of `mat1` must match the row dimension of vectors in `vec1`. Bias and resulting output vector sequence `z` have as many numbers of rows as `mat1`. `vec1` is a sequence of `vec_count` number of input vectors and bias is added to each resulting vector after multiplication with `mat1`. Thus, output `z` has dimensions `rows*vec_count`. `vec_count` number of input vectors and output vectors are provided as array of pointers arguments to kernel API.

`bias_shift` and `acc_shift` arguments are provided in kernel API to adjust Q format of bias and output respectively. Both `bias_shift` and `acc_shift` can be either positive or negative where positive value denotes a left shift and negative value denotes a right shift.

`bias_shift` is the shift in number of bits applied to the bias to make it in the same Q format as `matXvec` multiplication – accumulation result. `acc_shift` is the shift in number of bits applied to the accumulator to obtain the output in desired Q format.

**Note:** `acc_shift` and `bias_shift` are not relevant in case of floating point kernels.

The `row_stridel` argument is provided in kernel API for row offset of `mat1`. Note, input matrix is expected to be appropriately padded in case of `row_stridel > cols1`.

For conversion from higher precision accumulator to lower precision output, symmetric rounding is used.

Function variants available are `xa_nn_matXvec_batch_[p]x[q]_[r]`, where:

- `[p]`: Matrix precision in bits
- `[q]`: Vector precision in bits
- `[r]`: Output precision in bits

#### Precision

There are four variants available:

Type	Description
16x16_64	16-bit matrix inputs, 16-bit vector inputs, 64-bit output vectors
8x16_64	8-bit matrix inputs, 16-bit vector inputs, 64-bit output vectors
8x8_32	8-bit matrix inputs, 8-bit vector inputs, 32-bit output vectors
f32xf32_f32	float32 matrix inputs, float32 vector inputs, float32 output

## Algorithm

$$z_{n,i} = 2^{acc-shift} \left( \sum_{m=0}^{cols1-1} mat1_{n,m} \cdot vec1_{m,i} + 2^{bias-shift} bias_n \right),$$

$$n = 0, \dots, \overline{rows} - 1 ; \quad i = 0, \dots, \overline{vec-count} - 1$$

In case of floating point routine, acc\_shift=0 and bias\_shift=0.

Thus,  $2^{acc-shift} = 2^{bias-shift} = 1$

## Prototype

```
WORD32 xa_nn_matXvec_batch_16x16_64
(WORD64 ** p_out,          WORD16 * p_mat1,          WORD16 ** p_vec1,
 WORD16 * p_bias,          WORD32 rows,              WORD32 cols1,
 WORD32 row_stridel,       WORD32 acc_shift,          WORD32 bias_shift,
 WORD32 vec_count);

WORD32 xa_nn_matXvec_batch_8x16_64
(WORD64 ** p_out,          WORD8 * p_mat1,            WORD16 ** p_vec1,
 WORD16 * p_bias,          WORD32 rows,              WORD32 cols1,
 WORD32 row_stridel,       WORD32 acc_shift,          WORD32 bias_shift,
 WORD32 vec_count);

WORD32 xa_nn_matXvec_batch_8x8_32
(WORD32 ** p_out,          WORD8 * p_mat1,            WORD8 ** p_vec1,
 WORD8 * p_bias,           WORD32 rows,              WORD32 cols1,
 WORD32 row_stridel,       WORD32 acc_shift,          WORD32 bias_shift,
 WORD32 vec_count);

WORD32 xa_nn_matXvec_batch_f32xf32_f32
(FLOAT32 ** p_out,        FLOAT32 * p_mat1,          FLOAT32 ** p_vec1,
 FLOAT32 * p_bias,        WORD32 rows,              WORD32 cols1,
 WORD32 row_stridel,      WORD32 vec_count);
```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
WORD16 *, WORD8 *, FLOAT32 *	p_mat1	rows*cols1	Input matrix, fixed or floating point
WORD16 **, WORD8 **, FLOAT32 **	p_vec1	cols1*vec_count	Input vector pointers, fixed or floating point
WORD16 *, WORD8 *, FLOAT32 *	p_bias	rows*1	Bias vector, fixed or floating point
WORD32	rows		Number of rows in input matrix, bias and output
WORD32	cols1		Number of columns in input matrix and rows in input vector
WORD32	row_stridel		Row offset of input matrix
WORD32	acc_shift		Shift applied to accumulator
WORD32	bias_shift		Shift applied to bias
WORD32	vec_count		Number of input vectors



Type	Name	Size	Description
<b>Output</b>			
WORD32 **, WORD64 **, FLOAT32 **	p_out	rows*vec_count	Output vector pointers, fixed or floating point

## Returns

- 0: no error
- 1: error, invalid parameters

## Restrictions

Arguments	Restrictions
row_stride1, cols1	Multiples of 4 (2 in case of floating point)
p_mat1, p_vec1, p_bias, p_out	Aligned on 16-byte boundary Should not overlap Cannot be NULL
acc_shift, bias_shift	{-31, ..., 31}

## 3.1.4 Matrix Multiplication Kernels

### Description

These kernels perform the operation of multiplication of a matrix `mat1` with another matrix `mat2` along with bias addition; that is,  $z = \text{mat1} * \text{mat2} + \text{bias}$ . The first matrix should be stored in row major order and the second matrix should be stored in column major order. The first matrix is of dimensions `rows` x `cols`. The second matrix `mat2` is of dimensions `cols` x `vec_count`. These kernels can also be viewed as a modification of the Matrix X Vector Batch kernels. The column dimension of `mat1` matches the row dimension of `mat2` i.e. the length of each vector in `p_mat2`. Bias and resulting output vector sequence `z` have as many numbers of rows as `mat1`. `mat2` is a sequence of `vec_count` number of input vectors and bias is added to each resulting vector after multiplication with `mat1`. Thus, output `z` has dimensions `rows` \* `vec_count`. The arguments `vec_offset` and `out_offset` are offsets to the next vector and output addresses. The argument `out_stride` defines the row offset for the output matrix. For standard matrix multiplication, `vec_offset` should be equal to `cols`, `out_offset` equal to 1 and `out_stride` should be equal to `vec_count` i.e. columns of `mat2`.

The `bias_shift` and `acc_shift` arguments are provided in kernel API to adjust Q format of bias and output respectively. Both `bias_shift` and `acc_shift` can be either positive or negative where positive value denotes a left shift and negative value denotes a right shift.

The `bias_shift` is the shift in number of bits applied to the bias to make it in the same Q format as multiplication – accumulation result. `acc_shift` is the shift in number of bits applied to the accumulator to obtain the output in desired Q format.

Note, the `acc_shift` and `bias_shift` arguments are not relevant in case of quantized 8-bit kernels.

The `row_stride` argument indicates the offset to next row of `mat1`.

The `vec_offset` argument refers to the column offset of `mat2`.

Similarly, the `out_offset` and `out_stride` arguments refer to the column offset and row offset of the output matrix `rows * vec_count` respectively.

For conversion from higher precision accumulator to lower precision output, symmetric rounding is used.

The arguments, `mat1_zero_bias`, `mat2_zero_bias`, are provided to convert the quantized 8-bit inputs into their real values and perform `matXvec` batch operation. The `out_zero_bias`, `out_multiplier` and `out_shift` values are used to quantize real values of output back to quantized 8-bit values.

For the quantized int8 variant, we have per-row quantized input `mat1`.

Function variants available are `xa_nn_matmul_[p]x[q]_[r]`, where:

- `[p]`: Matrix 1 precision in bits
- `[q]`: Matrix 2 precision in bits
- `[r]`: Output precision in bits

## Precision

There are three variants available:

Type	Description
<code>8x8_8</code>	8-bit matrix inputs, 8-bit vector inputs, 8-bit output vectors
<code>asym8uxasym8u_asym8u</code>	asym8u matrix inputs, asym8u vector inputs, asym8u output vectors
<code>per_chan_sym8sxasym8s_asym8s</code>	per channel quantized sym8s matrix inputs, asym8s vector inputs, asym8s output vectors

## Algorithm

$$z_{n,i} = 2^{acc-shift} \left( \sum_{m=0}^{cols1-1} mat1_{n,m} \cdot mat2_{m,i} + 2^{bias-shift} bias_n \right),$$

$$n = 0, \dots, \overline{rows} - 1 ; \quad i = 0, \dots, \overline{vec-count} - 1$$

In case of quantized 8-bit routines, `acc_shift=0` and `bias_shift=0`.

Thus,  $2^{acc-shift} = 2^{bias-shift} = 1$

## Prototype

```
WORD32 xa_nn_matmul_8x8_8
(WORD8 * p_out,          WORD8 * p_mat1,          WORD8 * p_mat2,
 WORD8 * p_bias,         WORD32 rows,              WORD32 cols,
```

```

WORD32 row_stride,      WORD32 acc_shift,      WORD32 bias_shift,
WORD32 vec_count,      WORD32 vec_offset,    WORD32 out_offset,
WORD32 out_stride);
WORD32 xa_nn_matmul_asym8uxasym8u_asym8u
(UWORD8 * p_out,        const UWORD8 * p_mat1,  const UWORD8 * p_mat2,
 const WORD32 * p_bias, WORD32 rows,          WORD32 cols,
 WORD32 row_stride,    WORD32 vec_count,      WORD32 vec_offset,
 WORD32 out_offset,    WORD32 out_stride,      WORD32 mat1_zero_bias,
 WORD32 vec1_zero_bias, WORD32 out_multiplier, WORD32 out_shift,
 WORD32 out_zero_bias);
WORD32 xa_nn_matmul_per_chan_sym8sxasym8s_asym8s
(WORD8 * p_out,         const WORD8 * p_mat1,   const WORD8 * p_mat2,
 const WORD32 * p_bias, WORD32 rows,          WORD32 cols,
 WORD32 row_stride,    WORD32 vec_count,      WORD32 vec_offset,
 WORD32 out_offset,    WORD32 out_stride,      WORD32 vec1_zero_bias
 const WORD32 *p_out_multiplier, const WORD32 p_out_shift,
 WORD32 out_zero_bias);

```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
WORD8 *, UWORD8 *, const WORD8 *	p_mat1	rows*cols	Input matrix 1, fixed point, asym8u or sym8s
WORD8 *, UWORD8 *, const WORD8 *	p_mat2	cols * vec_count	Input matrix 2, fixed , asym8u or sym8s
WORD8 *, const WORD32 *	p_bias	rows*1	Bias vector, fixed point
WORD32	rows		Number of rows in input matrix, bias and output
WORD32	cols		Number of columns in input matrix and rows in input vector
WORD32	row_stride		Row offset of input matrix
WORD32	acc_shift		Shift applied to accumulator
WORD32	bias_shift		Shift applied to bias
WORD32	vec_count		Number of vectors (columns) in matrix 2
WORD32	vec_offset		Offset to the next vector address
WORD32	out_offset		Offset to the next output address
WORD32	out_stride		Row offset of output matrix
WORD32	mat1_zero_bias		Zero offset of matrix 1
WORD32	vec1_zero_bias		Zero offset of matrix 2
WORD32	out_multiplier		Multiplier value of output
WORD32	out_shift		Shift value of output
WORD32	out_zero_bias		Zero offset of output
<b>Output</b>			
WORD8 *, UWORD8 *	p_out	rows*vec_ count	Output matrix, fixed-point, floating point or asym8u

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions

Arguments	Restrictions
p_mat1, p_mat2, p_out	Aligned on (size of one element)-byte boundary Cannot be NULL Should not overlap
p_bias	Aligned on (size of one element)-byte boundary
acc_shift, bias_shift, out_shift	{-31, ..., 31}
vec_count	Greater than 0
vec_offset, out_offset, out_stride	Should not be 0
mat1_zero_bias,	{-255, ..., 0}
vec1_zero_bias	{-255, ..., 0} for asym8u, {-127, ..., 127} for asym8s
out_multiplier	Greater than 0
out_zero_bias	{0, ..., 255} if out type is asym8u, {-128, ..., 127} if out type is asym8s

## 3.1.5 Matrix X Vector Kernels with Output Stride

### Description

These kernels perform a single matXvec operation with bias addition; that is,  $z = \text{mat1} * \text{vec1} + \text{bias}$ . The column dimension of `mat1` must match the row dimension of `vec1`. Bias and resulting output vector `z` have as many rows as `mat1`.

`row_stride1` is provided in kernel API for row offsets of `mat1`. Note, input matrices are expected to be appropriately padded in case of `row_stride > cols`.

For conversion from higher precision accumulator to lower precision output, symmetric rounding is used.

The argument `out_stride` is helpful in storing the output at a given offset.

The argument `vec1_zero_bias` is provided to convert the quantized 8-bit inputs into their real values and perform matXvec operation. The `out_multiplier` and `out_shift` values are used to convert real values of output to 16-bit.

Function variants available are `xa_nn_matXvec_[p]x[q]_[r]`, where:

- `[p]`: Matrix precision in bits
- `[q]`: Vector precision in bits

- [r]: Output precision in bits

## Precision

There is one variant available:

Type	Description
sym8sxasym8s_16	sym8s matrix inputs, asym8s vector inputs, asym8s output

## Algorithm

$$z_n = \left( \sum_{m=0}^{cols1-1} mat1_{n,m} \cdot vec1_m + bias_n \right)$$

## Prototype

```
WORD32 xa_nn_matXvec_out_stride_sym8sxasym8s_16
(WORD16 * p_out,      const WORD8 * p_mat1,  const WORD8 * p_vec1,
 const WORD32 * p_bias, WORD32 rows,        WORD32 cols1,
 WORD32 row_stride1,  WORD32 out_stride,     WORD32 vec1_zero_bias,
 WORD32 out_multiplier, WORD32 out_shift);
```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
const WORD8 *	p_mat1	rows*cols1	Input matrix, sym8s
const WORD8 *	p_vec1	cols1*1	Input vector, asym8s
const WORD32 *	p_bias	rows*1	Bias vector
WORD32	rows		Number of rows in matrix and number of elements in bias
WORD32	cols1		Number of columns in matrix and elements in vector
WORD32	row_stride1		Row offset of matrix
WORD32	out_stride		Row offset of output
WORD32	vec1_zero_bias		Zero offset of vector
WORD32	out_multiplier		Multiplier value of output
WORD32	out_shift		Shift value of output
<b>Output</b>			
WORD16 *	p_out	rows*1	Output, 16-bit

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions

Arguments	Restrictions
row_stridel, cols1	row_stride1 >= cols1
p_mat1, p_vec1, p_bias, p_out	Aligned on <size of one element> boundary Should not overlap
p_mat1, p_vec1, p_out	Cannot be NULL
out_shift	{-31, ..., 31}
vec1_zero_bias	{-127....., 128} for asym8s
out_multiplier	Greater than 0

### 3.1.6 Matrix X Vector Batch Kernels with Accumulation

These kernels perform the operation of multiplication of a single matrix with a series of vectors along with bias addition; that is,  $z_i = z_i + \text{mat1} * \text{vec1}_i + \text{bias}$ . These kernels can also be viewed as matrix X matrix-transpose multiplication kernels. The column dimension of `mat1` must match the row dimension of vectors in `vec1`. Bias and resulting output vector sequence `z` have as many numbers of rows as `mat1`. `vec1` is a sequence of `vec_count` number of input vectors and bias is added to each resulting vector after multiplication with `mat1`. Thus, output `z` has dimensions `rows*vec_count`. `vec_count` number of input and output vectors are provided as pointers to the start of first vector, subsequent vectors are supposed to be stored contiguously in memory. The result of matrix X vector batch operation is accumulated to the values present at the output.

The `row_stridel` argument is provided in kernel API for row offset of `mat1`. Note, input matrix is expected to be appropriately padded in case of `row_stridel > cols1`.

The `out_zero_bias`, `out_multiplier` and `out_shift` values are used to quantize the output to 16-bits.

Function variants available are `xa_nn_matXvec_acc_batch_[p]x[q]_[r]`, where:

- `[p]`: Matrix precision in bits
- `[q]`: Vector precision in bits
- `[r]`: Output precision in bits

## Precision

There is one variant available:

Type	Description
<code>sym8sx8_asym16s</code>	sym8s matrix inputs, 8-bit vector inputs, asym16s output vectors

## Algorithm

$$z_{n,i} = z_{n,i} + \left( \sum_{m=0}^{cols1-1} mat1_{n,m} \cdot vec1_{m,i} + bias_n \right),$$

$$n = 0, \dots, rows - 1 ; \quad i = 0, \dots, vec\_count - 1$$

## Prototype

```
WORD32 xa_nn_matXvec_acc_batch_sym8sx8_asym16s
(WORD16 * p_out,      const WORD8 * p_mat1,      const WORD8 * p_vec1,
 const WORD32 * p_bias, WORD32 rows,              WORD32 cols1,
 WORD32 row_stride1,  WORD32 out_multiplier,      WORD32 out_shift,
 WORD32 out_zero_bias, WORD32 vec_count);
```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
const WORD8 *	p_mat1	rows*cols1	Input matrix, sym8s
const WORD8 *	p_vec1	cols1*vec_count	Input vectors, 8-bit
const WORD32 *	p_bias	rows*1	Bias vector, 32-bit
WORD32	rows		Number of rows in input matrix, bias and output
WORD32	cols1		Number of columns in input matrix and rows in input vector
WORD32	row_stride1		Row offset of input matrix
WORD32	out_multiplier		Multiplier value of output
WORD32	out_shift		Shift value of output
WORD32	out_zero_bias		Zero offset of output
WORD32	vec_count		Number of input vectors
<b>Output</b>			
WORD16	p_out	rows*vec_count	Output vectors, asym16s

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions

Arguments	Restrictions
p_mat1, p_vec1, p_bias, p_out	Aligned on <size of one element> boundary
	Cannot be NULL
	Should not overlap
rows, cols1, vec_count	Should be greater than 0.
row_stride1	Cannot be less than cols1

out_shift	{-31, ..., 31}
out_zero_bias	{-32768, ..., 32767}

## 3.2 Convolution Kernels

### 3.2.1 Standard 2D Convolution Kernel

#### Description

These kernels perform the 2D convolution operation as  $z = \text{inp} (*) \text{kernel} + \text{bias}$ . A 3D input cube ( $\text{input\_height} \times \text{input\_width} \times \text{input\_channels}$ ), is convolved with a 3D kernel cube ( $\text{kernel\_height} \times \text{kernel\_width} \times \text{input\_channels}$ ) to produce a 2D convolution output plane ( $\text{out\_height} \times \text{out\_width}$ ). With `out_channels` number of such 3D kernels, output cube ( $\text{out\_height} \times \text{out\_width} \times \text{out\_channels}$ ) is produced. The bias having dimension (`out_channels`) is added after the convolution (one bias value is added to each output channel) to produce the final output.

**Note:** The depth or channels dimension (`input_channels`) of input and kernel must be identical for 2D convolution.

`bias_shift` and `acc_shift` arguments are provided in kernel API to adjust Q format of bias and output, respectively. Both `bias_shift` and `acc_shift` can be either positive or negative where positive value denotes a left shift and negative value denotes a right shift.

`bias_shift` is the shift in number of bits applied to the bias to make it in the same Q format as convolution - accumulation result. `acc_shift` is the shift in number of bits applied to the accumulator to obtain the output in desired Q format.

**Note:** `acc_shift` and `bias_shift` are not relevant in case of floating point kernels and quantized 8-bit kernels.

The `x_stride` and `y_stride` arguments in kernel API define the step size of the kernel when traversing the input in width and height dimensions respectively.

The `x_padding` argument defines padding to the left of the input in the width dimension and the `y_padding` argument defines padding to the top of the input in the height dimension.

The right padding is calculated based on `out_width` as  $\text{right\_padding} = \text{kernel\_width} + (\text{out\_width} - 1) * \text{x\_stride} - (\text{x\_padding} + \text{input\_width})$ .

The bottom padding is calculated based on `out_height` as  $\text{bottom\_padding} = \text{kernel\_height} + (\text{out\_height} - 1) * \text{y\_stride} - (\text{y\_padding} + \text{input\_height})$ .

For conversion from higher precision accumulator to lower precision output, symmetric rounding is used.

The kernel is expected to be padded in the depth or channels dimension if the number of `input_channels` is not a multiple of 4 in case of fixed point variants other than the 8x8, `asym8uxasym8u` and



per\_chan\_sym8sxasym8s variant, and 2 in case of floating point variant. No padding is needed for 8x8 and quantized 8-bit variants.

These kernels require temporary buffer for convolution computation. This temporary buffer is provided by p\_scratch argument of kernel API. The size of temporary buffer should be queried using xa\_nn\_conv2d\_std\_getsize() helper API.

These kernels expect input and kernel cubes in SHAPE\_CUBE\_DWH\_T shape type and can produce output cube in either SHAPE\_CUBE\_DWH\_T or SHAPE\_CUBE\_WHD\_T shape type. The out\_data\_format argument to kernel API controls the output cube shape type.

Function variants available are xa\_nn\_conv2d\_std\_[p]x[q], where:

- [p]: Kernel precision in bits
- [q]: Input precision in bits

## Precision

There are six variants available.

Type	Description
16x16	16-bit kernel, 16-bit input, 16-bit output
8x16	8-bit kernel, 16-bit input, 16-bit output
8x8	8-bit kernel, 8-bit input, 8-bit output
f32	float32 kernel, float32 input, float32 output
asym8u	asym8u kernel, asym8u input, asym8u output
per_chan_sym8sxasym8s	per channel quantized sym8s kernel, asym8s input, asym8s output

## Algorithm

$$Z_{h,w,d} = 2^{acc-shift} \left( \sum_{i=0}^{K_H-1} \sum_{j=0}^{K_W-1} \sum_{k=0}^{I_C-1} in_{pad}(h*y-stride+i),(w*x-stride+j),k \cdot ker_{pad,d,i,j,k} + 2^{bias-shift} b_d \right)$$

$$h = 0, \dots, \overline{out-height - 1}, w = 0, \dots, \overline{out-width - 1},$$

$$d = 0, \dots, \overline{out-channels - 1}$$

In case of floating-point kernels and quantized 8-bit kernels, acc\_shift=0 and bias\_shift=0.

Thus,  $2^{acc-shift} = 2^{bias-shift} = 1$

$in_{pad}, ker_{pad}$  denote the padded p\_inp and padded p\_ker shapes, respectively.

$K_H, K_W, I_C$  denote kernel\_height, kernel\_width, and input\_channels, respectively.

$b$  denotes the bias shape.

## Prototype

```
WORD32 xa_nn_conv2d_std_getsize
(WORD32 input_height,      WORD32 input_channels, WORD32 kernel_height,
 WORD32 kernel_width,     WORD32 y_stride,      WORD32 y_padding,
 WORD32 out_height,       WORD32 input_precision);

WORD32 xa_nn_conv2d_std_16x16
(WORD16 * p_out,          WORD16 * p_inp,        WORD16 * p_ker,
 WORD16 * p_bias,         WORD32 input_height,   WORD32 input_width,
 WORD32 input_channels,   WORD32 kernel_height, WORD32 kernel_width,
 WORD32 out_channels,     WORD32 x_stride,       WORD32 y_stride,
 WORD32 x_padding,        WORD32 y_padding,      WORD32 out_height,
 WORD32 out_width,        WORD32 bias_shift,     WORD32 acc_shift,
 WORD32 out_data_format,  VOID * p_scratch);

WORD32 xa_nn_conv2d_std_8x16
(WORD16 * p_out,          WORD16 * p_inp,        WORD8 * p_ker,
 WORD16 * p_bias,         WORD32 input_height,   WORD32 input_width,
 WORD32 input_channels,   WORD32 kernel_height, WORD32 kernel_width,
 WORD32 out_channels,     WORD32 x_stride,       WORD32 y_stride,
 WORD32 x_padding,        WORD32 y_padding,      WORD32 out_height,
 WORD32 out_width,        WORD32 bias_shift,     WORD32 acc_shift,
 WORD32 out_data_format,  VOID * p_scratch);

WORD32 xa_nn_conv2d_std_8x8
(WORD8 * p_out,           WORD8 * p_inp,         WORD8 * p_ker,
 WORD8 * p_bias,          WORD32 input_height,   WORD32 input_width,
 WORD32 input_channels,   WORD32 kernel_height, WORD32 kernel_width,
 WORD32 out_channels,     WORD32 x_stride,       WORD32 y_stride,
 WORD32 x_padding,        WORD32 y_padding,      WORD32 out_height,
 WORD32 out_width,        WORD32 bias_shift,     WORD32 acc_shift,
 WORD32 out_data_format,  VOID * p_scratch);

WORD32 xa_nn_conv2d_std_f32
(FLOAT32 * p_out,         const FLOAT32 * p_inp, const FLOAT32 * p_ker,
 const FLOAT32 * p_bias,  WORD32 input_height,   WORD32 input_width,
 WORD32 input_channels,   WORD32 kernel_height, WORD32 kernel_width,
 WORD32 out_channels,     WORD32 x_stride,       WORD32 y_stride,
 WORD32 x_padding,        WORD32 y_padding,      WORD32 out_height,
 WORD32 out_width,        WORD32 out_data_format, VOID * p_scratch);

WORD32 xa_nn_conv2d_std_asym8uxasym8u
(UWORD8 * p_out,          const UWORD8 * p_inp,  const UWORD8 * p_ker,
 const WORD32 * p_bias,   WORD32 input_height,   WORD32 input_width,
 WORD32 input_channels,   WORD32 kernel_height, WORD32 kernel_width,
 WORD32 out_channels,     WORD32 x_stride,       WORD32 y_stride,
 WORD32 x_padding,        WORD32 y_padding,      WORD32 out_height,
 WORD32 out_width,        WORD32 input_zero_bias, WORD32 kernel_zero_bias,
 WORD32 out_multiplier,   WORD32 out_shift,      WORD32 out_zero_bias,
 WORD32 out_data_format,  VOID * p_scratch);

WORD32 xa_nn_conv2d_std_per_chan_sym8sxasym8s
(WORD8 * p_out,           const WORD8 * p_inp,   const WORD8 * p_ker,
 const WORD32 * p_bias,   WORD32 input_height,   WORD32 input_width,
 WORD32 input_channels,   WORD32 kernel_height, WORD32 kernel_width,
 WORD32 out_channels,     WORD32 x_stride,       WORD32 y_stride,
 WORD32 x_padding,        WORD32 y_padding,      WORD32 out_height,
 WORD32 out_width,        WORD32 input_zero_bias, WORD32 * p_out_multiplier,
 WORD32 * p_out_shift,    WORD32 out_zero_bias,  WORD32 out_data_format,
 VOID * p_scratch);
```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
WORD16 *, WORD8 *, const FLOAT32 *, const UWORD8 *, const WORD8 *	p_inp	input_height* input width* input_channels	Input cube, fixed, floating point, asym8u or asym8s, in SHAPE_CUBE_DWH_T
WORD16 *, WORD8 *, const FLOAT32 *, const UWORD8 *, const WORD8 *	p_ker	out_channels* (kernel_height* kernel width* input_channels)	Kernel cube, fixed, floating point, asym8u or sym8s, in SHAPE_CUBE_DWH_T
WORD16 *, WORD8 *, FLOAT32 *, const WORD32 *	p_bias	out_channels	Bias vector, fixed or floating point
WORD32	input_height		Input height
WORD32	input_width		Input width
WORD32	input_channels		Number of input channels
WORD32	kernel_height		Kernel height
WORD32	kernel_width		Kernel width
WORD32	out_channels		Number of output channels
WORD32	x_stride		Horizontal stride over input
WORD32	y_stride		Vertical stride over input
WORD32	x_padding		Left padding width on input
WORD32	y_padding		Top padding height on input
WORD32	out_height		Output height
WORD32	out_width		Output width
WORD32	bias_shift		Shift applied to bias
WORD32	acc_shift		Shift applied to accumulator
WORD32	input_zero_bias		Zero offset of input
WORD32	kernel_zero_bias		Zero offset of kernel
WORD32	out_multiplier		Multiplier value of output
WORD32	out_shift		Shift value of output
WORD32	out_zero_bias		Zero offset of output
WORD32	out_data_format		Output data format 0:SHAPE_CUBE_DWH_T 1:SHAPE_CUBE_WHD_T
VOID *	p_scratch	xa_nn_conv2d_s td_getsize()	Scratch memory pointer

Output			
WORD16 *, WORD8 *, FLOAT32 *, UWORD8 *	p_out	(out_height* out_width)* out_channels	Output cube, fixed, floating point, asym8u or asym8s, as per the out_data_format argument.

## Returns

- 0: no error
- -1: error, invalid parameters
- 

## Restrictions

Arguments	Restrictions
p_out, p_inp, p_ker, p_bias, p_scratch	Cannot be NULL Should not overlap Aligned on 16-byte boundary except for quantized 8-bit kernels where only p_scratch is required to be 16-byte aligned. For p_scratch - memory size >= size returned by xa_nn_conv2d_std_getsize()
input_height, input_width, input_channels	Greater than or equal to 1
kernel_height	{1, 2, ..., input_height}
kernel_width	{1, 2, ..., input_width}
out_channels	Greater than or equal to 1
x_stride	{1, 2, ..., kernel_width}
y_stride	Greater than or equal to 1
x_padding, y_padding	Greater than or equal to 0
out_height, out_width	Greater than or equal to 1
acc_shift, bias_shift, out_shift	{-31 .... 31} for fixed point and quantized 8-bit APIs
input_zero_bias	{-255, ..., 0} for asym8u input, {-127, ..., 128} for asym8s input
kernel_zero_bias	{-255, ..., 0} for asym8u kernel
out_zero_bias	{0, ..., 255} for asym8u output, {-128, ..., 127} for asym8s output
out_multiplier	Greater than 0
out_data_format	Can be 0: SHAPE_CUBE_DWH_T or 1: SHAPE_CUBE_WHD_T

## 3.2.2 Standard 1D Convolution Kernel

### Description

These kernels perform the 1D convolution operation as  $z = \text{inp} (*) \text{kernel} + \text{bias}$ . A 3D input cube (input\_height x input\_width x input\_channels) is convolved with a 3D kernel cube

(`kernel_height` x `input_width` x `input_channels`) to produce a 1D convolution output vector (`out_height`). With `out_channels` number of such 3D kernels, output matrix (`out_height` x `out_channels`) is produced. The bias having dimension (`out_channels`) is added after the convolution (one bias value is added to each output column) to produce the final output.

**Note:** The depth or channels dimension (`input_channels`) of input and kernel must be identical, and width dimension (`input_width`) of input and kernel also must be identical for 1D convolution.

`bias_shift` and `acc_shift` arguments are provided in kernel API to adjust Q format of bias and output, respectively. Both `bias_shift` and `acc_shift` can be either positive or negative, where positive value denotes a left shift and negative value denotes a right shift.

`bias_shift` is the shift in number of bits applied to the bias to make it in the same Q format as convolution - accumulation result. `acc_shift` is the shift in number of bits applied to the accumulator to obtain the output in desired Q format.

**Note:** `acc_shift` and `bias_shift` are not relevant in case of floating point kernels.

The `y_stride` argument to kernel API defines the step size of the kernel when traversing the input in height dimension.

The `y_padding` argument defines padding to the top of the input in the height dimension.

The bottom padding is calculated based on `out_height` as `bottom_paddding = kernel_height + (out_height - 1) * y_stride - (y_padding + input_height)`.

For conversion from higher precision accumulator to lower precision output, symmetric rounding is used.

The kernel is expected to be padded if the product `input_channels*input_width` is not a multiple of 4 in case of fixed point variants, and 2 in case of floating point variant.

These kernels require temporary buffer for convolution computation. This temporary buffer is provided by `p_scratch` argument of kernel API. The size of temporary buffer should be queried using `xa_nn_conv1d_std_getsize()` helper API.

These kernels expect input and kernel cubes in `SHAPE_CUBE_DWH_T` shape type and can produce output matrix with either (`out_height` x `out_channels`) or (`out_channels` x `out_height`) dimensions. The `out_data_format` argument to kernel API controls the output matrix height and width order.

Function variants available are `xa_nn_conv1d_std_[p]`, where:

- `[p]`: precision in bits

## Precision

There are four variants available:

Type	Description
------	-------------

16x16	16-bit kernel, 16-bit input, 16-bit output
8x16	8-bit kernel, 16-bit input, 16-bit output
8x8	8-bit kernel, 8-bit input, 8-bit output
f32	float32 kernel, float32 input, float32 output

## Algorithm

$$z_{h,d} = 2^{acc-shift} \left( \sum_{i=0}^{K_H-1} \sum_{j=0}^{I_W-1} \sum_{k=0}^{I_C-1} in_{pad(h*y-stride+i),j,k} \cdot ker_{pad_{d,i,j,k}} + 2^{bias-shift} b_d \right)$$

$$h = 0, \dots, \overline{out-height} - 1, d = 0, \dots, \overline{out-channels} - 1$$

In case of floating point kernel, `acc_shift=0` and `bias_shift=0`.

Thus,  $2^{acc-shift} = 2^{bias-shift} = 1$

$n_{pad}, ker_{pad}$  denote the padded `p_inp` and padded `p_ker` shapes, respectively.

$K_H, I_W, I_C$  denote `kernel_height`, `input_width`, and `input_channels`, respectively.

$b$  denotes the `bias` shape.

## Prototype

```
WORD32 xa_nn_convld_std_getsize
(WORD32 kernel_height, WORD32 input_width, WORD32 input_channels,
 WORD32 input_precision);

WORD32 xa_nn_convld_std_16x16
(WORD16 * p_out, WORD16 * p_inp, WORD16 * p_ker,
 WORD16 * p_bias, WORD32 input_height, WORD32 input_width,
 WORD32 input_channels, WORD32 kernel_height, WORD32 out_channels,
 WORD32 y_stride, WORD32 y_padding, WORD32 out_height,
 WORD32 bias_shift, WORD32 acc_shift, WORD32 out_data_format,
 VOID * p_scratch);

WORD32 xa_nn_convld_std_8x16
(WORD16 * p_out, WORD16 * p_inp, WORD8 * p_ker,
 WORD16 * p_bias, WORD32 input_height, WORD32 input_width,
 WORD32 input_channels, WORD32 kernel_height, WORD32 out_channels,
 WORD32 y_stride, WORD32 y_padding, WORD32 out_height,
 WORD32 bias_shift, WORD32 acc_shift, WORD32 out_data_format,
 VOID * p_scratch);

WORD32 xa_nn_convld_std_8x8
(WORD8 * p_out, WORD8 * p_inp, WORD8 * p_ker,
 WORD8 * p_bias, WORD32 input_height, WORD32 input_width,
 WORD32 input_channels, WORD32 kernel_height, WORD32 out_channels,
 WORD32 y_stride, WORD32 y_padding, WORD32 out_height,
 WORD32 bias_shift, WORD32 acc_shift, WORD32 out_data_format,
 VOID * p_scratch);

WORD32 xa_nn_convld_std_f32
```

```

(FLOAT32 * p_out,          FLOAT32 * p_inp,          FLOAT32 * p_ker,
 FLOAT32 * p_bias,        WORD32 input_height,      WORD32 input_width,
 WORD32 input_channels,   WORD32 kernel_height, WORD32 out_channels,
 WORD32 y_stride,         WORD32 y_padding,       WORD32 out_height,
 WORD32 out_data_format,  VOID * p_scratch);

```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
WORD16 *, WORD8 *, FLOAT32 *,	p_inp	input_height* input_width* input_channels	Input cube, fixed or floating point, in SHAPE_CUBE_DWHT
WORD16 *, WORD8 *, FLOAT32 *,	p_ker	out_channels* (kernel_height* input_width* input_channels)	Kernel cube, fixed or floating point, in SHAPE_CUBE_DWHT
WORD16 *, WORD8 *, FLOAT32 *,	p_bias	out_channels	Bias vector, fixed or floating point
WORD32	input_height		Input height
WORD32	input_width		Input width
WORD32	input_channels		Number of input channels
WORD32	kernel_height		Kernel height
WORD32	out_channels		Number of output channels
WORD32	y_stride		Vertical stride over input
WORD32	y_padding		Top padding height on input
WORD32	out_height		Output height
WORD32	bias_shift		Shift applied to bias
WORD32	acc_shift		Shift applied to accumulator
WORD32	out_data_format		Output matrix order 0: out_height x out_channels 1: out_channels x out_height
VOID *	p_scratch	xa_nn_conv1d_std_getsize()	Scratch memory pointer
<b>Output</b>			
WORD16 *, WORD8 *, FLOAT32 *,	p_out	out_height* out_channels	Output matrix, fixed or floating point, as per the out_data_format argument.

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions

Arguments	Restrictions
p_out, p_inp, p_ker, p_bias, p_scratch	Cannot be NULL Should not overlap Aligned on 16-byte boundary For p_scratch - memory size >= size returned by xa_nn_conv1d_std_getsize()
input_height, input_width, input_channels	Greater than or equal to 1
kernel_height	{1,2,...,input_height}
out_channels	Greater than or equal to 1
y_stride	{1,2,...,kernel_height}
y_padding	Greater than or equal to 0
out_height	Greater than or equal to 1
acc_shift, bias_shift	{-31 .... 31} for fixed point APIs
out_data_format	Can be 0: out_height x out_channels or 1: out_channels x out_height

### 3.2.3 Depthwise Separable 2D Convolution Kernel

Depthwise Separable 2D Convolution is computed in two steps using following two low level kernels:

- First step: xa\_nn\_conv2d\_depthwise\_xx() low level kernel  
These kernels convolve each input 2D plane (input\_height x input\_width) from input cube (input\_height x input\_width x input\_channels) with channels\_multiplier number of 2D kernels (kernel\_height x kernel\_width) to produce channels\_multiplier number of 2D output planes (out\_height x out\_width). Thus, with kernel cube of dimension (kernel\_height x kernel\_width x (channels\_multiplier \* input\_channels)), output cube of dimension (out\_height x out\_width x (channels\_multiplier \* input\_channels)) is produced. Bias is added to the convolution output. There is one bias value for each output 2D plane; that is, bias is a vector of dimension (channels\_multiplier \* input\_channels).
- Second step: xa\_nn\_conv2d\_pointwise\_xx() low level kernel  
These kernels take output cube (out\_height x out\_width x (channels\_multiplier \* input\_channels)) of first step as input and perform pointwise multiplication with kernel vector (channels\_multiplier \* input\_channels) in depth dimension to produce output 2D plane (out\_height x out\_width). Thus, with out\_channels kernel vectors, output cube of dimension (out\_height x out\_width x out\_channels) is produced. Bias is added to the pointwise multiplication output. There is one bias value for each output 2D plane; that is, bias is a vector of dimension out\_channels.

Following are the descriptions for these two low level kernels.



### 3.2.3.1 Depthwise 2D Convolution Kernel

#### Description

These kernels perform the 2D depthwise convolution operation as  $z = \text{inp} (*) \text{kernel} + \text{bias}$ . These kernels convolve each input 2D plane ( $\text{input\_height} \times \text{input\_width}$ ) from input cube ( $\text{input\_height} \times \text{input\_width} \times \text{input\_channels}$ ) with  $\text{channels\_multiplier}$  number of 2D kernels ( $\text{kernel\_height} \times \text{kernel\_width}$ ) to produce  $\text{channels\_multiplier}$  number of 2D output planes ( $\text{out\_height} \times \text{out\_width}$ ). Thus, with kernel cube of dimension ( $\text{kernel\_height} \times \text{kernel\_width} \times (\text{channels\_multiplier} * \text{input\_channels})$ ), output cube of dimension ( $\text{out\_height} \times \text{out\_width} \times (\text{channels\_multiplier} * \text{input\_channels})$ ) is produced. Bias is added to the convolution output. There is one bias value for each output 2D plane; that is, bias is a vector of dimension ( $\text{channels\_multiplier} * \text{input\_channels}$ ).

`bias_shift` and `acc_shift` arguments are provided in kernel API to adjust Q format of bias and output respectively. Both `bias_shift` and `acc_shift` can be either positive or negative where positive value denotes a left shift and negative value denotes a right shift.

`bias_shift` is the shift in number of bits applied to the bias to make it in the same Q format as convolution - accumulation result. `acc_shift` is the shift in number of bits applied to the accumulator to obtain the output in desired Q format.

**Note:** `acc_shift` and `bias_shift` are not relevant in case of floating point kernels and quantized 8-bit kernels.

The `x_stride` and `y_stride` arguments in kernel API define the step size of the kernel when traversing the input in width and height dimensions, respectively.

The `x_padding` argument defines padding to the left of the input in the width dimension, and `y_padding` argument defines padding to the top of the input in the height dimension.

The right padding is calculated based on `out_width` as `right_padding = kernel_width + (out_width - 1) * x_stride - (x_padding + input_width)`.

The bottom padding is calculated based on `out_height` as `bottom_padding = kernel_height + (out_height - 1) * y_stride - (y_padding + input_height)`.

For conversion from higher precision accumulator to lower precision output, symmetric rounding is used.

These kernels require a temporary buffer for convolution computation. This temporary buffer is provided by the `p_scratch` argument of kernel API. The size of temporary buffer should be queried using `xa_nn_conv2d_depthwise_getsize()` helper API.

The arguments `input_zero_bias`, `kernel_zero_bias` are provided to convert the quantized 8-bit inputs into their real values and perform Depthwise 2D Convolution operation. The `out_zero_bias`, `out_multiplier` and `out_shift` values are used to quantize real values of output back to 8-bit.

The depthwise kernels expect input cube in SHAPE\_CUBE\_DWH\_T and SHAPE\_CUBE\_WHD\_T shape type and produce output cube in SHAPE\_CUBE\_DWH\_T shape type respectively. The `inp_data_format` argument to the kernel API can be 0 or 1 to indicate input cube shape respectively.

The `out_data_format` argument to the kernel API must be 0 for all the kernels to indicate output cube shape.

Function variants available are `xa_nn_conv2d_depthwise_[p]`, where:

- `[p]`: precision in bits

## Precision

There are six variants available:

Type	Description
16x16	16-bit kernel, 16-bit input, 16-bit output
8x16	8-bit kernel, 16-bit input, 16-bit output
8x8	8-bit kernel, 8-bit input, 8-bit output
f32	float32 kernel, float32 input, float32 output
asym8uxasym8u	asym8u kernel, asym8u input, asym8u output
per_chan_sym8sxasym8s	per channel quantized sym8s kernel, asym8s input, asym8s output

## Algorithm

$$z_{h,w,d \cdot C_M + m} = 2^{acc\_shift} \left( \sum_{i=0}^{K_H-1} \sum_{j=0}^{K_W-1} in_{pad_{(h \cdot y\_stride + i)(w \cdot x\_stride + j),d}} \cdot ker_{pad_{i,j,(d \cdot C_M + m)}} + 2^{bias\_shift} b_{0,0,d \cdot C_M + m} \right)$$

$$h = 0, \dots, \overline{out\_height - 1}, w = 0, \dots, \overline{out\_width - 1}, \\ d = 0, \dots, \overline{input\_channels - 1}, \\ m = 0, \dots, \overline{channels\_multiplier - 1}$$

In case of floating-point kernel and quantized 8-bit kernels, `acc_shift=0` and `bias_shift=0`.

Thus,  $2^{acc\_shift} = 2^{bias\_shift} = 1$

$in_{pad}, ker_{pad}$  denote the padded `p_inp` and padded `p_ker` shapes, respectively.

$K_H, K_W, C_M$  denote `kernel_height`, `kernel_width`, and `channels_multiplier`, respectively.

$b$  denotes the bias shape.

## Prototype

```
WORD32 xa_nn_conv2d_depthwise_getsize
(WORD32 input_height,      WORD32 input_width,      WORD32 input_channels,
 WORD32 kernel_height,    WORD32 kernel_width,    WORD32 channels_multiplier,
```

```

WORD32 x_stride,          WORD32 y_stride,          WORD32 x_padding,
WORD32 y_padding,        WORD32 output_height,      WORD32 output_width,
WORD32 circ_buf_precision, WORD32 inp_data_format);

WORD32 xa_nn_conv2d_depthwise_16x16
(WORD16 * p_out,          WORD16 * p_ker,          WORD16 * p_inp,
WORD16 * p_bias,          WORD32 input_height,      WORD32 input_width,
WORD32 input_channels,    WORD32 kernel_height,    WORD32 kernel_width,
WORD32 channels_multiplier, WORD32 x_stride,        WORD32 y_stride,
WORD32 x_padding,         WORD32 y_padding,        WORD32 out_height,
WORD32 out_width,         WORD32 acc_shift,         WORD32 bias_shift,
WORD32 inp_data_format,   WORD32 out_data_format, VOID * p_scratch);

WORD32 xa_nn_conv2d_depthwise_8x16
(WORD16 * p_out,          WORD8 * p_ker,          WORD16 * p_inp,
WORD16 * p_bias,          WORD32 input_height,      WORD32 input_width,
WORD32 input_channels,    WORD32 kernel_height,    WORD32 kernel_width,
WORD32 channels_multiplier, WORD32 x_stride,        WORD32 y_stride,
WORD32 x_padding,         WORD32 y_padding,        WORD32 out_height,
WORD32 out_width,         WORD32 acc_shift,         WORD32 bias_shift,
WORD32 inp_data_format,   WORD32 out_data_format, VOID * p_scratch);

WORD32 xa_nn_conv2d_depthwise_8x8
(WORD8 * p_out,           WORD8 * p_ker,           WORD8 * p_inp,
WORD8 * p_bias,           WORD32 input_height,      WORD32 input_width,
WORD32 input_channels,    WORD32 kernel_height,    WORD32 kernel_width,
WORD32 channels_multiplier, WORD32 x_stride,        WORD32 y_stride,
WORD32 x_padding,         WORD32 y_padding,        WORD32 out_height,
WORD32 out_width,         WORD32 acc_shift,         WORD32 bias_shift,
WORD32 inp_data_format,   WORD32 out_data_format, VOID * p_scratch);

WORD32 xa_nn_conv2d_depthwise_f32
(FLOAT32 * p_out,         const FLOAT32 * p_ker, const FLOAT32 * p_inp,
const FLOAT32 * p_bias,   WORD32 input_height,      WORD32 input_width,
WORD32 input_channels,    WORD32 kernel_height,    WORD32 kernel_width,
WORD32 channels_multiplier, WORD32 x_stride,        WORD32 y_stride,
WORD32 x_padding,         WORD32 y_padding,        WORD32 out_height,
WORD32 out_width,         WORD32 inp_data_format, WORD32 out_data_format,
VOID * p_scratch);

WORD32 xa_nn_conv2d_depthwise_asym8uxasym8u
(pUWORD8 p_out,           const UWORD8 * p_kernel, const UWORD8 * p_inp,
const WORD32 * p_bias,    WORD32 input_height,      WORD32 input_width,
WORD32 input_channels,    WORD32 kernel_height,    WORD32 kernel_width,
WORD32 channels_multiplier, WORD32 x_stride,        WORD32 y_stride,
WORD32 x_padding,         WORD32 y_padding,        WORD32 out_height,
WORD32 out_width,         WORD32 input_zero_bias, WORD32 kernel_zero_bias,
WORD32 out_multiplier,    WORD32 out_shift,         WORD32 out_zero_bias,
WORD32 inp_data_format,   WORD32 out_data_format, pVOID p_scratch);

WORD32 xa_nn_conv2d_depthwise_per_chan_sym8sxsasym8s
(pWORD8 p_out,            const WORD8 * p_kernel, const WORD8 * p_inp,
const WORD32 * p_bias,    WORD32 input_height,      WORD32 input_width,
WORD32 input_channels,    WORD32 kernel_height,    WORD32 kernel_width,
WORD32 channels_multiplier, WORD32 x_stride,        WORD32 y_stride,
WORD32 x_padding,         WORD32 y_padding,        WORD32 out_height,
WORD32 out_width,         WORD32 input_zero_bias, const WORD32 * p_out_multiplier,
const WORD32 * p_out_shift, WORD32 out_zero_bias, WORD32 inp_data_format,
WORD32 out_data_format,   pVOID p_scratch);

```

## Arguments

Type	Name	Size	Description
Input			

Type	Name	Size	Description
WORD16 *, WORD8 *, FLOAT32 *, const UWORD8 *, const WORD8 *	p_ker	kernel_height* kernel_width* input_channels* channels_multiplier	Kernel cube, fixed, floating point, asym8u or sym8s, in SHAPE_CUBE_D WH or SHAPE_CUBE_W HD_T
WORD16 *, WORD8 *, FLOAT32 *, const UWORD8 *, const WORD8 *	p_inp	input_height* input_width* input_channels	Input cube, fixed, floating point, asym8u or asym8s in SHAPE_CUBE_D WH or SHAPE_CUBE_W HD_T
WORD16 *, WORD8 *, FLOAT32 *, const WORD32 *	p_bias	input_channels*chann els_multiplier	Bias vector, fixed or floating point
WORD32	input_height		Input height
WORD32	input_width		Input width
WORD32	input_channels		Number of input channels
WORD32	kernel_height		Kernel height
WORD32	kernel_width		Kernel width
WORD32	channels_multipl ier		Multiplier value for each input channel
WORD32	x_stride		Horizontal stride over input
WORD32	y_stride		Vertical stride over input
WORD32	x_padding		Left padding width on input
WORD32	y_padding		Right padding height on input
WORD32	out_height		Output height
WORD32	out_width		Output width
WORD32	acc_shift		Shift applied to accumulator
WORD32	bias_shift		Shift applied to bias
WORD32	input_zero_bias		Zero offset of input
WORD32	kernel_zero_bias		Zero offset of kernel
WORD32	out_multiplier		Multiplier value of output
WORD32	out_shift		Shift value of output
WORD32 *	p_out_multiplier	input_channels*chann els_multiplier	Array of multiplier values of output

Type	Name	Size	Description
WORD32 *	p_out_shift	input_channels*channels_multiplier	Array of shift values of output
WORD32	out_zero_bias		Zero offset of output
WORD32	inp_data_format		Input and Kernel data format 0:SHAPE_CUBE_DWH_T 1:SHAPE_CUBE_WHD_T
WORD32	out_data_format		Output data format 0:SHAPE_CUBE_DWH_T
VOID *	p_scratch	xa_nn_conv2d_depthwise_getsize()	Scratch memory pointer
<b>Output</b>			
WORD16 *, WORD8 *, UWORD8 *, FLOAT32 *	p_out	out_height* out width* input_channels* channels_multiplier	Output cube, fixed, floating point, asym8u or asym8s, in SHAPE_CUBE_DWH_T

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions

Arguments	Restrictions
p_out, p_ker, p_inp, p_bias,	Cannot be NULL
	Should not overlap
	Aligned on <size of one element> boundary
p_scratch	Cannot be NULL
	Should not overlap with other buffers
	Aligned on 16-byte boundary
	For p_scratch - memory size >= size returned by xa_nn_conv2d_depthwise_getsize()
p_out_multiplier	Cannot be NULL
	Should not overlap
	Aligned on 4-byte boundry
p_out_shift	Cannot be NULL
	Should not overlap
	Aligned on 4-byte boundry
	Each 32-bit value should be in range [-31 ... 31]

input_height, input_width, input_channels	Greater than or equal to 1
kernel_height	{1,2,...,input_height}
kernel_width	{1,2,...,input_width}
channels_multiplier	Greater than or equal to 1
x_stride	{1,2,...,kernel_width}
y_stride	{1,2,...,kernel_height}
x_padding, y_padding	Greater than or equal to 0
out_height, out_width	Greater than or equal to 1
acc_shift, bias_shift, out_shift	{-31 .... 31} for fixed point and quantized 8-bit APIs
input_zero_bias	{-255,....., 0} for asym8u input, {-127....., 128} for asym8s input
kernel_zero_bias	{-255....., 0} for asym8u kernel
out_zero_bias	{0,.....,255} for asym8u output, {-128....., 127} for asym8s output
out_multiplier	Greater than 0
inp_data_format	can be 0: SHAPE_CUBE_DWH_T or 1: SHAPE_CUBE_WHD_T
out_data_format	must be 0: SHAPE_CUBE_DWH_T

### 3.2.3.2 Pointwise 2D Convolution Kernel

#### Description

These kernels perform pointwise multiplication of input cube (input\_height x input\_width x input\_channels) with kernel vector (input\_channels) in depth dimension to produce output 2D plane (input\_height x input\_width). Thus, with out\_channels kernel vectors, output cube of dimension (input\_height x input\_width x out\_channels) is produced. Bias is added to the pointwise multiplication output. There is one bias value for each output 2D plane; that is, bias is a vector of dimension out\_channels.

The bias\_shift and acc\_shift arguments are provided in kernel API to adjust Q format of bias and output respectively. Both bias\_shift and acc\_shift can be either positive or negative, where positive value denotes a left shift and negative value denotes a right shift.

bias\_shift is the shift in number of bits applied to the bias to make it in the same Q format as convolution - accumulation result. acc\_shift is the shift in number of bits applied to the accumulator to obtain the output in desired Q format.

**Note:** acc\_shift and bias\_shift are not relevant in case of floating point kernels and quantized 8-bit kernels.

For conversion from higher precision accumulator to lower precision output, symmetric rounding is used.

These kernels expect input cube in SHAPE\_CUBE\_DWH\_T shape type, kernel as matrix, bias as vector and produce output cube in SHAPE\_CUBE\_WHD\_T or SHAPE\_CUBE\_DWH\_T (only for 8x8, asym8uxasym8u and per\_chan\_sym8sxasym8s kernels) shape type. The out\_data\_format argument

to kernel API must be always 1 except for 8x8 and quantized 8-bit kernels for which it can be 0 or 1 indicating SHAPE\_CUBE\_DWH\_T and SHAPE\_CUBE\_WHD\_T respectively.

Function variants available are `xa_nn_conv2d_pointwise_[p]`, where:

- `[p]`: precision in bits

## Precision

There are six variants available:

Type	Description
16x16	16-bit kernel, 16-bit input, 16-bit output
8x16	8-bit kernel, 16-bit input, 16-bit output
8x8	8-bit kernel, 8-bit input, 8-bit output
f32	float32 kernel, float32 input, float32 output
asym8u/asym8s	asym8u kernel, asym8u input, asym8u output
per_chan_sym8s/asym8s	per channel quantized sym8s kernel, asym8s input, asym8s output

## Algorithm

$$z_{h,w,d} = 2^{acc-shift} \left( \sum_{k=0}^{I_C-1} in_{h,w,k} \cdot ker_{d,0,0,k} + 2^{bias-shift} b_{0,0,d} \right)$$

$$h = 0, \dots, \overline{input-height} - 1, w = 0, \dots, \overline{input-width} - 1,$$

$$d = 0, \dots, \overline{out\_channels} - 1$$

In case of floating-point kernel and quantized 8-bit kernels, `acc_shift=0` and `bias_shift=0`. Thus,  $2^{acc-shift} = 2^{bias-shift} = 1$

`in`, `ker` denote the `p_inp`, and `p_ker` shapes respectively.

$I_C$  denotes `input_channels`

`b` denotes the bias shape

## Prototype

```
WORD32 xa_nn_conv2d_pointwise_16x16
(WORD16 * p_out,          WORD16 * p_ker,          WORD16 * p_inp,
 WORD16 * p_bias,         WORD32 input_height,     WORD32 input_width,
 WORD32 input_channels,   WORD32 out_channels,     WORD32 acc_shift,
 WORD32 bias_shift,      WORD32 out_data_format);

WORD32 xa_nn_conv2d_pointwise_8x16
(WORD16 * p_out,          WORD8 * p_ker,           WORD16 * p_inp,
 WORD16 * p_bias,         WORD32 input_height,     WORD32 input_width,
 WORD32 input_channels,   WORD32 out_channels,     WORD32 acc_shift,
 WORD32 bias_shift,      WORD32 out_data_format);

WORD32 xa_nn_conv2d_pointwise_8x8
(WORD8 * p_out,           WORD8 * p_ker,           WORD8 * p_inp,
 WORD8 * p_bias,         WORD32 input_height,     WORD32 input_width,
```

```

WORD32 input_channels,      WORD32 out_channels,      WORD32 acc_shift,
WORD32 bias_shift,         WORD32 out_data_format);
WORD32 xa_nn_conv2d_pointwise_f32
(FLOAT32 * p_out,          const FLOAT32 * p_ker, const FLOAT32 * p_inp,
 const FLOAT32 * p_bias,      WORD32 input_height,   WORD32 input_width,
WORD32 input_channels,      WORD32 out_channels,
WORD32 out_data_format);
WORD32 xa_nn_conv2d_pointwise_asym8uxasym8u
(UWORD8 * p_out,          const UWORD8 * p_ker, const UWORD8 * p_inp,
WORD32 * p_bias,          WORD32 input_height,   WORD32 input_width,
WORD32 input_channels,      WORD32 out_channels,   WORD32 input_zero_bias,
WORD32 kernel_zero_bias,    WORD32 out_multiplier, WORD32 out_shift,
WORD32 out_zero_bias,      WORD32 out_data_format);
WORD32 xa_nn_conv2d_pointwise_asym8uxasym8u
(UWORD8 * p_out,          const UWORD8 * p_ker, const UWORD8 * p_inp,
 const WORD32 * p_bias,      WORD32 input_height,   WORD32 input_width,
WORD32 input_channels,      WORD32 out_channels,   WORD32 input_zero_bias,
WORD32 kernel_zero_bias,    WORD32 out_multiplier, WORD32 out_shift,
WORD32 out_zero_bias,      WORD32 out_data_format);
WORD32 xa_nn_conv2d_pointwise_per_chan_sym8sxasym8s
(WORD8 * p_out,          const WORD8 * p_ker, const WORD8 * p_inp,
 const WORD32 * p_bias,      WORD32 input_height,   WORD32 input_width,
WORD32 input_channels,      WORD32 out_channels,   WORD32 input_zero_bias,
WORD32 * p_out_multiplier, WORD32 * p_out_shift,   WORD32 out_zero_bias,
WORD32 out_data_format);

```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
WORD16 *, WORD8 *, const FLOAT32 *, const UWORD8 *, const WORD8 *	p_ker	out_channels * input_channels	Kernel matrix, fixed, floating point, asym8u or asym8s, (out_channels x input_channels)
WORD16 *, WORD8 *, const FLOAT32 *, const UWORD8 *, const WORD8 *	p_inp	input_height* input width* input_channels	Input cube, fixed or floating point, asym8u or sym8s, in SHAPE_CUBE_DWHLT
WORD16 *, WORD8 *, FLOAT32 *, const WORD32 *	p_bias	out_channels	Bias vector, fixed or floating point
WORD32	input_height		Input height
WORD32	input_width		Input width
WORD32	input_channels		Number of input channels
WORD32	out_channels		Number of output channels
WORD32	acc_shift		Shift applied to accumulator
WORD32	bias_shift		Shift applied to bias
WORD32	input_zero_bias		Zero offset of input



WORD32	kernel_zero_bias		Zero offset of kernel
WORD32	out_multiplier		Multiplier value of output
WORD32	out_shift		Shift value of output
WORD32	out_zero_bias		Zero offset of output
WORD32	out_data_format		Output data format 0: SHAPE_CUBE_DWH_T 1: SHAPE_CUBE_WHD_T
<b>Output</b>			
WORD16 *, WORD8 *, FLOAT32 *, UWORD8 *	p_out	(out_height* out_width)* out_channels	Output cube, fixed, floating point, asym8u or asym8s, as per the out_data_format argument.

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions

Arguments	Restrictions
p_out, p_ker, p_inp, p_bias	Cannot be NULL Should not overlap Aligned on 16-byte boundary except for 8x8 and quantized 8-bit kernels
input_height, input_width	Greater than or equal to 1
input_channels	Greater than or equal to 4, multiple of 4 except for 8x8 and asym8u kernels
out_channels	Greater than or equal to 1
acc_shift, bias_shift, out_shift	{-31 ..... 31} for fixed point and quantized 8-bit APIs
input_zero_bias	{-255,....., 0} for asym8u input, {-127....., 128} for asym8s input
kernel_zero_bias	{-255....., 0} for asym8u kernel
out_zero_bias	{0,.....,255} for asym8u output, {-128....., 127} for asym8s output
out_multiplier	Greater than 0
out_data_format	Can be 0: SHAPE_CUBE_DWH_T or 1: SHAPE_CUBE_WHD_T for 8x8 and quantized 8-bit kernels. Must be 1 for other kernels.

## 3.3 Activation Kernels

### 3.3.1 Sigmoid

#### Description

These kernels perform the sigmoid operation on input vector  $x$  and give output vector as  $y = \text{sigmoid}(x)$ . Both the input and output vectors have size `vec_length`.

The 32-bit input fixed-point kernels accept 32-bit input in Q6.25 format and give output in Q16.15 (32-bit), Q15 (16-bit), or Q7 (8-bit) format. The 16-bit input/output fixed-point kernel accepts the input in Q3.12 and give output in Q15 (16-bit) format.

For the `asym8u` and `asym8s` kernels both the input and output are of `asym8u` and `asym8s` datatype respectively.

The 16-bit fixed point variant and the quantized 8-bit variants of sigmoid are based on Tensorflow implementations.

**Note:** The `input_range_radius` argument for quantized 8-bit variants is derived from other input parameters in Tensorflow. The kernel does not perform dependency check on the `input_range_radius` and the user will have to ensure that correct value is passed.

Function variants available are `xa_nn_vec_sigmoid_[p]_[q]`, where:

- `[p]`: Input precision in bits
- `[q]`: Output precision in bits

#### Precision

There are seven variants available.

Type	Description
<code>32_32</code>	32-bit input, 32-bit output
<code>32_16</code>	32-bit input, 16-bit output
<code>32_8</code>	32-bit input, 8-bit output
<code>16_16</code>	16-bit input, 16-bit output
<code>f32_f32</code>	float32 input, float32 output
<code>asym8u_x_asym8u</code>	<code>asym8u</code> input, <code>asym8u</code> output
<code>asym8s_x_asym8s</code>	<code>asym8s</code> input, <code>asym8s</code> output

#### Algorithm

$$y_n = \frac{1}{1 + \exp(-x_n)}, \quad n = 0, \dots, \text{vec\_length} - 1$$

## Prototype

```

WORD32 xa_nn_vec_sigmoid_32_32
(WORD32 * p_out,          const WORD32 * p_vec,      WORD32 vec_length);
WORD32 xa_nn_vec_sigmoid_32_16
(WORD16 * p_out,          const WORD32 * p_vec,      WORD32 vec_length);
WORD32 xa_nn_vec_sigmoid_32_8
(WORD8 * p_out,           const WORD32 * p_vec,      WORD32 vec_length);
WORD32 xa_nn_vec_sigmoid_f32_f32
(FLOAT32 * p_out,         const FLOAT32 * p_vec,     WORD32 vec_length);
WORD32 xa_nn_vec_sigmoid_asym8u_asym8u
(UWORD8 * p_out,          const UWORD8 * p_vec,      WORD32 zero_point,
WORD32 input_range_radius, WORD32 input_multiplier, WORD32 input_left_shift,
WORD32 vec_length);
WORD32 xa_nn_vec_sigmoid_asym8s_asym8s
(WORD8 * p_out,           const WORD8 * p_vec,       WORD32 zero_point,
WORD32 input_range_radius, WORD32 input_multiplier, WORD32 input_left_shift,
WORD32 vec_length);
WORD32 xa_nn_vec_sigmoid_16_16
(WORD16 * p_out,          const WORD16 * p_vec,      WORD32 vec_length);

```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
const WORD32 *, const WORD16 *, const UWORD8 *, const FLOAT32 *, const WORD8 *	p_vec	vec_length	Input vector, Q6.25, Q3.12, floating point, asym8u or asym8s
WORD32	zero_point		bias value
WORD32	input_range_radius		Range radius: For asym8u output = ((x <sub>i</sub> - zero_point) < radius)? sigmoid() : 255 output = ((x <sub>i</sub> - zero_point) > (-radius))? sigmoid() : 0 For asym8s output = ((x <sub>i</sub> - zero_point) < radius)? sigmoid() : 127 output = ((x <sub>i</sub> - zero_point) > (-radius))? sigmoid() : -128
WORD32	input_multiplier		Multiplier value of input
WORD32	input_left_shift		Left Shift value of input
WORD32	vec_length		Length of input vector
<b>Output</b>			
WORD32 *, WORD16 *, WORD8 *, UWORD8 *, FLOAT32 *	p_out	vec_length	Output vector, fixed (Q16.15, Q15, Q7), floating point, asym8u or asym8s

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions

Arguments	Restrictions
p_vec, p_out	Should not overlap Cannot be NULL
zero_point	[0, 255] for asym8u [-128, 127] for asym8s
input_range_radius	[0, 255]
input_left_shift	[-31, 31]
input_multiplier	Shouldn't be less than 0.
vec_length	Greater than 0

## 3.3.2 Tanh

### Description

These kernels perform the hyperbolic tangent operation on input vector  $x$  and give output vector as  $y = \tanh(x)$ . Both the input and output vectors have size `vec_length`.

The 32-bit input fixed-point kernels accept 32-bit input in Q6.25 format and give output in Q16.15 (32-bit), Q15 (16-bit), or Q7 (8-bit) format. The 16-bit fixed-point kernel has input argument `integer_bits` to specify the number of integer bits in input so input Q format is  $Q(\text{integer\_bits}).(15 - \text{integer\_bits})$ , output is given in Q15 (16-bit) format.

For the asym8s kernels both the input and output are of asym8s datatype.

The 16-bit fixed point variant and the quantized 8-bit variants of tanh are based on Tensorflow implementations.

**Note:** The `input_range_radius` argument for quantized 8-bit variant is derived from other input parameters in Tensorflow. The kernel does not perform dependency check on the `input_range_radius` and the user will have to ensure that correct value is passed.

Function variants available are `xa_nn_vec_tanh_[p]_[q]`, where:

- [p]: Input precision in bits
- [q]: Output precision in bits

### Precision

There are six variants available:

Type	Description
------	-------------

32_32	32-bit input, 32-bit output
32_16	32-bit input, 16-bit output
32_8	32-bit input, 8-bit output
16_16	16-bit input, 16-bit output
f32_f32	float32 input, float32 output
asym8sxasym8s	asym8s input, asym8s output

## Algorithm

$$y_n = \tanh(x_n), \quad n = 0, \dots, \overline{vec\_length} - 1$$

## Prototype

```
WORD32 xa_nn_vec_tanh_32_32
(WORD32 * p_out,          const WORD32 * p_vec,      WORD32 vec_length);
WORD32 xa_nn_vec_tanh_32_16
(WORD16 * p_out,          const WORD32 * p_vec,      WORD32 vec_length);
WORD32 xa_nn_vec_tanh_32_8
(WORD8 * p_out,           const WORD32 * p_vec,      WORD32 vec_length);
WORD32 xa_nn_vec_tanh_f32_f32
(FLOAT32 * p_out,         const FLOAT32 * p_vec,     WORD32 vec_length);
WORD32 xa_nn_vec_tanh_asym8s_asym8s
(WORD8 * p_out,           const WORD8 * p_vec,       WORD32 zero_point,
WORD32 input_range_radius, WORD32 input_multiplier, WORD32 input_left_shift,
WORD32 vec_length);
WORD32 xa_nn_vec_tanh_16_16
(WORD16 * p_out,          const WORD16 * p_vec,      WORD32 integer_bits,
WORD32 vec_length);
```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
const WORD32 *, const WORD16 *, const FLOAT32 *, const WORD8 *	p_vec	vec_length	Input vector, Q6.25, Q(integer_bits).(15- integer_bits), floating point or asym8s
WORD32	zero_point		Bias value
WORD32	input_range_radius		Range radius: output = ((x <sub>i</sub> - zero_point) < radius) ? tanh() : 127 output = ((x <sub>i</sub> - zero_point) > (-radius)) ? tanh() : -128
WORD32	input_multiplier		Multiplier value of input
WORD32	input_left_shift		Left shift value of input
WORD32	vec_length		Length of input vector
WORD32	integer_bits		Number of integer bits in the 16-bit input
<b>Output</b>			

WORD32 *, WORD16 *, WORD8 *, FLOAT32 *	p_out	vec_length	Output vector, fixed (Q16.15, Q15, Q7), floating point or asym8s
---	-------	------------	--

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions

Arguments	Restrictions
p_vec, p_out	Should not overlap Cannot be NULL
zero_point	[-128, 127]
input_range_radius	[0, 255]
input_multiplier	Shouldn't be less than 0
vec_length	Greater than 0
integer_bits	[0, 6]

### 3.3.3 Rectifier Linear Unit (ReLU)

#### Description

These kernels compute the rectifier linear unit function of input vector  $x$  and give output vector as  $y = \text{relu}(x)$ . Both the input and output vectors have size `vec_length`.

The fixed-point routines accept 32-bit input in Q6.25 format and gives 32-bit output in Q16.15 format.

The `threshold` argument to `relu` kernel API allows to set upper threshold for proper compression of output signal and is expected in Q16.15 format. In `relu1` and `relu6` kernels, the thresholds are set to 1 and 6, respectively.

For the `asym8u` and `asym8s` kernels, the quantized input is requantized and applied the standard ReLU function to give the output. The `threshold` argument is not applicable for quantized ReLU kernels.

The standard ReLU kernels `relu_std` can be used when the `threshold` is not required.

Function variants available are `xa_nn_vec_relu_[p]_[q]`, `xa_nn_vec_relu1_[p]_[q]`, and `xa_nn_vec_relu6_[p]_[q]`, where:

- `[p]`: Input precision in bits
- `[q]`: Output precision in bits

#### Precision

There are six variants available:

Type	Description
<code>32_32</code>	32-bit input, 32-bit output
<code>f32_f32</code>	float32 input, float32 output
<code>16_16</code>	16-bit input, 16-bit output
<code>8_8</code>	8-bit input, 8-bit output
<code>asym8s_asym8s</code>	asym8s input, asym8s output

#### Algorithm

$$y_n = \max(0, \min(x_n, K)), \quad n = 0, \dots, \text{vec\_length} - 1$$

$K$  represents threshold

#### Prototype

```
WORD32 xa_nn_vec_relu_32_32
(WORD32 * p_out,      const WORD32 * p_vec,      WORD32 threshold,
 WORD32 vec_length);
WORD32 xa_nn_vec_relu_f32_f32
(FLOAT32 * p_out,     const FLOAT32 * p_vec,     FLOAT32 threshold,
 WORD32 vec_length);
WORD32 xa_nn_vec_relu_16_16
```

```

(WORD16 * p_out,      const WORD16 * p_vec,   WORD16 threshold,
 WORD32 vec_length);
WORD32 xa_nn_vec_relu_8_8
(WORD8 * p_out,      const WORD8 * p_vec,   WORD8 threshold,
 WORD32 vec_length);
WORD32 xa_nn_vec_relu_asym8s_asym8s
(WORD8 * p_out,      const WORD8 * p_vec,   WORD32 inp_zero_bias,
 WORD32 out_multiplier, WORD32 out_shift,   WORD32 out_zero_bias,
 WORD32 quantized_activation_min, WORD32 quantized_activation_max,
 WORD32 vec_length);
WORD32 xa_nn_vec_relu1_32_32
(WORD32 * p_out,      const WORD32 * p_vec,   WORD32 vec_length);
WORD32 xa_nn_vec_relu1_f32_f32
(FLOAT32 * p_out,      const FLOAT32 * p_vec,   WORD32 vec_length);
WORD32 xa_nn_vec_relu6_32_32
(WORD32 * p_out,      const WORD32 * p_vec,   WORD32 vec_length);
WORD32 xa_nn_vec_relu6_f32_f32
(FLOAT32 * p_out,      const FLOAT32 * p_vec,   WORD32 vec_length);
WORD32 xa_nn_vec_relu_std_32_32
(WORD32 * p_out,      const WORD32 * p_vec,   WORD32 vec_length);
WORD32 xa_nn_vec_relu_std_f32_f32
(FLOAT32 * p_out,      const FLOAT32 * p_vec,   WORD32 vec_length);
WORD32 xa_nn_vec_relu_std_16_16
(WORD16 * p_out,      const WORD16 * p_vec,   WORD32 vec_length);
WORD32 xa_nn_vec_relu_std_8_8
(WORD8 * p_out,      const WORD8 * p_vec,   WORD32 vec_length);

```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
const WORD32 *, const FLOAT32 *, const WORD16 *, const WORD8 *, const UWORD8 *	p_vec	vec_length	Input vector, fixed-point, floating point, asym8u or asym8s
WORD32	inp_zero_bias		Zero bias value for input vector
WORD32	out_multiplier		Fixed-point multiplier value for output
WORD32	out_shift		Shift value for output
WORD32	vec_length		length of input vector
WORD32	out_zero_bias		Zero bias value for output vector
WORD32	quantized_act ivation_min		Lower threshold value, quantized.
WORD32, FLOAT32	quantized_act ivation_max		Upper threshold value, quantized
WORD32 FLOAT32 WORD16 WORD8	threshold		threshold, fixed or floating point
<b>Output</b>			
WORD32 *, FLOAT32 *, WORD16 *, WORD8 *, UWORD8 *	p_out	vec_length	Output vector, fixed-point, floating point, asym8u or asym8s



## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions

Arguments	Restrictions
p_vec, p_out	Should not overlap Cannot be NULL
inp_zero_bias, out_zero_bias	{0,.....,255} for asym8u, {-128....., 127} for asym8s input
out_multiplier	Shouldn't be less than 0.
out_shift	{-31, ..., 31}
quantized_activation_min quantized_activation_max	{0,.....,255} for asym8u output, {-128....., 127} for asym8s output quantized_activation_min < quantized_activation_max

### 3.3.4 Softmax

#### Description

These kernels compute the softmax (normalized exponential function) of input vector  $x$  and give output vector as  $y = \text{softmax}(x)$ . Both the input and output vectors have size `vec_length`.

The fixed-point kernels accept 32-bit input in Q6.25 format and give 32-bit output in Q16.15 format.

For the `asym8u` kernels, both the input and output are of the same precision and for `asym8s` kernels, the input is `asym8s` and the output precision can be `asym8s` or 16-bit.

Function variants available are `xa_nn_vec_softmax_[p]_[q]`, where:

- `[p]`: Input precision in bits
- `[q]`: Output precision in bits

#### Precision

There are five variants available:

Type	Description
<code>32_32</code>	32-bit input, 32-bit output
<code>f32_f32</code>	float32 input, float32 output
<code>asym8u_asym8u</code>	asym8u input, asym8u output
<code>asym8s_asym8s</code>	asym8s input, asym8s output
<code>asym8s_16</code>	asym8s input, 16-bit output

#### Algorithm

$$y_n = \frac{\exp(x_n)}{\sum_k \exp(x_k)}, \quad n = 0, \dots, \text{vec\_length} - 1$$

#### Prototype

```
WORD32 xa_nn_vec_softmax_32_32
(WORD32 * p_out, const WORD32 * p_vec, WORD32 vec_length);
WORD32 xa_nn_vec_softmax_f32_f32
(FLOAT32 * p_out, const FLOAT32 * p_vec, WORD32 vec_length);
WORD32 xa_nn_vec_softmax_asym8u_asym8u
(UWORD8 * p_out, const UWORD8 * p_vec, WORD32 diffmin,
WORD32 input_left_shift, WORD32 input_multiplier,
WORD32 vec_length, pVOID p_scratch);
WORD32 xa_nn_vec_softmax_asym8s_asym8s
(WORD8 * p_out, const WORD8 * p_vec, WORD32 diffmin,
WORD32 input_left_shift, WORD32 input_multiplier,
WORD32 vec_length, pVOID p_scratch);
WORD32 xa_nn_vec_softmax_asym8s_16
(WORD16 * p_out, const WORD8 * p_vec, WORD32 diffmin,
WORD32 input_left_shift, WORD32 input_multiplier,
WORD32 vec_length, pVOID p_scratch);
```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
WORD32 *, FLOAT32 *, const UWORD8 *, const WORD8 *	p_vec	vec_length	Input vector, Q6.25, floating point, asym8u or asym8s
WORD32	diffmin		Diffmin value: output = $((x_i - \max) > \text{diffmin}) ? \text{softmax}() : 0$
WORD32	input_left_shift		left shift value of input
WORD32	input_multiplier		multiplier value of input
WORD32	vec_length		Length of input vector
<b>Output</b>			
WORD32 *, FLOAT32 *, UWORD8 *, WORD8 *, WORD16 *	p_out	vec_length	Output vector, Q16.15, floating point, asym8u, asym8s or 16-bit.
<b>Temporary</b>			
VOID *	p_scratch		Scratch (temporary) memory pointer

## Returns

- 0: no error
- 1: error, invalid parameters

## Restrictions

Arguments	Restrictions
p_vec, p_out	Should not overlap Cannot be NULL
input_left_shift	[-31, 31]
input_multiplier	Shouldn't be less than 0.
vec_length	Greater than 0

## 3.3.5 Activation Min Max

### Description

These kernels compute the activation minimum and maximum value of input vector  $x$  and give output vector as  $y = \text{activation\_min\_max}(x)$ . Both the input and output vectors have size `num_elm`.

For activation min max kernels, the input precision and the output precision are same.

The `activation_min` and `activation_max` arguments to the kernel API allow to set the threshold for proper compression of the output. The kernel is a generic implementation of the ReLU function.

Function variant available is `xa_nn_vec_activation_min_max_[p]_[q]`, where:

- [p]: Input precision in bits
- [q]: Output precision in bits

## Precision

There are four variants available:

Type	Description
f32_f32	float32 input, float32 output
asym8u_asym8u	asym8u input, asym8u output
16_16	16-bit input, 16-bit output
8_8	8-bit input, 8-bit output

## Algorithm

$$y_n = \max(\text{activation\_min}, \min(x_n, \text{activation\_max})), \quad n = 0, \dots, \text{vec\_length} - 1$$

`activation_min` represents lower threshold.

`activation_max` represents upper threshold.

## Prototype

```
WORD32 xa_nn_vec_activation_min_max_f32_f32
(FLOAT32 * p_out, const FLOAT32 * p_vec, FLOAT32 activation_min,
 FLOAT32 activation_max, WORD32 vec_length);
WORD32 xa_nn_vec_activation_min_max_asym8u_asym8u
(UWORD8 * p_out, const UWORD8 * p_vec, int activation_min,
 int activation_max, WORD32 vec_length);
WORD32 xa_nn_vec_activation_min_max_16_16
(WORD16 * p_out, const WORD16 * p_vec, int activation_min,
 int activation_max, WORD32 vec_length);
WORD32 xa_nn_vec_activation_min_max_8_8
(WORD8 * p_out, const WORD8 * p_vec, int activation_min,
 int activation_max, WORD32 vec_length);
```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
const UWORD8 *, const FLOAT32 *, const WORD16 *, const WORD8 *	p_vec	vec_length	Input vector, floating-point, asym8u or fixed point.
WORD32	vec_length		Length of input vector
WORD32, FLOAT32	activation_min		Lower threshold value, floating-point, asym8u or fixed point.
WORD32, FLOAT32	activation_max		Upper threshold value, floating-point, asym8u or fixed point

Output			
UWORD8 *, FLOAT32 *, WORD16 *, WORD8 *	p_out	vec_length	Output vector, floating-point, asym8u or fixed point

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions

Arguments	Restrictions
p_vec, p_out	Aligned on (size of one element)-byte boundary Cannot be NULL

## 3.3.6 Hard Swish

### Description

These kernels compute the hard-swish function of input vector  $x$  and give output vector as  $y = \text{hard\_swish}(x)$ . Both the input and output vectors have size `vec_length`.

The hard-swish activation function is a type of activation function based on swish but replaces the computationally expensive sigmoid function by ReLU6.

Function variants available are `xa_nn_vec_hard_swish_[p]_[q]`, where:

- [p]: Input precision in bits
- [q]: Output precision in bits

### Precision

There is one variant available:

Type	Description
asym8s_asym8s	asym8s input, asym8s output

### Algorithm

$$y_n = x_n * [\text{ReLU6}(x_n + 3)/6], \quad n = 0, \dots, \overline{vec\_length} - 1$$

### Prototype

```
WORD32 xa_nn_vec_hard_swish_asym8s_asym8s
(WORD8 * p_out, const WORD8 * p_vec, WORD32 inp_zero_bias,
WORD16 reluish_multiplier, WORD32 reluish_shift, WORD16 out_multiplier,
WORD32 out_shift, WORD32 out_zero_bias, WORD32 vec_length);
```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
const WORD8 *	p_vec	vec_length	Input vector, asym8s
WORD32	inp_zero_bias		Zero bias value for input vector
WORD16	reluish_multiplier		Fixed-point multiplier value for reluish scale
WORD32	reluish_shift		Shift value for reluish scale
WORD16	out_multiplier		Fixed-point multiplier value for output
WORD32	out_shift		Shift value for output
WORD32	out_zero_bias		Zero bias value for output vector
WORD32	vec_length		length of input vector
<b>Output</b>			
WORD8 *	p_out	vec_length	Output vector, asym8s

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions

Arguments	Restrictions
p_vec, p_out	Cannot be NULL Should not overlap (the two pointers could be same, inplace operation is possible)
inp_zero_bias, out_zero_bias	{-128....., 127} for asym8s datatype
out_multiplier, reluish_multiplier	Shouldn't be less than 0
out_shift, reluish_shift	{-31, ..., 31}

## 3.3.7 Parametric ReLU (PReLU)

### Description

These kernels compute the Parametric ReLU function of input vector  $x$  and give output vector as  $y = \text{prelu}(x)$ . Both the input and output vectors have size `vec_length`.

The PReLU activation function acts like a standard ReLU function for input values greater than or equal to 0. For input values less than 0, a learnable negative slope parameter  $\alpha(a)$  is multiplied with input to get the output. This slope value for all the input elements is determined based on the  $\alpha$  input vector.

Function variants available are `xa_nn_vec_prelu_[p]_[q]`, where:

- [p]: Input precision in bits
- [q]: Output precision in bits

## Precision

There is one variant available:

Type	Description
asym8s_asym8s	asym8s input, asym8s output

## Algorithm

$$y_n = x_n, \quad \text{when } x_n \geq 0 \quad n = 0, \dots, \text{vec\_length} - 1$$

$$y_n = ax_n, \quad \text{when } x_n < 0$$

where a is the learnable negative slope parameter: alpha.

## Prototype

```
WORD32 xa_nn_vec_prelu_asym8s_asym8s
(WORD8 * p_out, const WORD8 * p_vec, const WORD8 * p_vec_alpha,
WORD32 inp_zero_bias, WORD32 alpha_zero_bias, WORD32 alpha_multiplier,
WORD32 alpha_shift, WORD32 out_multiplier, WORD32 out_shift,
WORD32 out_zero_bias, WORD32 vec_length);
```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
const WORD8 *	p_vec	vec_length	Input vector, asym8s
const WORD8 *	p_vec_alpha	vec_length	alpha input vector, asym8s
WORD32	inp_zero_bias		Zero bias value for input vector
WORD32	alpha_zero_bias		Zero bias value for alpha input vector
WORD16	alpha_multiplier		Fixed-point multiplier value for alpha input.
WORD32	alpha_shift		Shift value for alpha input.
WORD16	out_multiplier		Fixed-point multiplier value for output
WORD32	out_shift		Shift value for output
WORD32	out_zero_bias		Zero bias value for output vector
WORD32	vec_length		length of input vector
<b>Output</b>			
WORD8 *	p_out	vec_length	Output vector, asym8s

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions

Arguments	Restrictions
p_vec, p_out, p_vec_alpha	Cannot be NULL
	Should not overlap (the two pointers could be same, inplace operation is possible)
inp_zero_bias, alpha_zero_bias	{-127....., 128} for asym8s datatype
out_zero_bias	{-128....., 127} for asym8s datatype
out_multiplier, alpha_multiplier	Shouldn't be less than 0
out_shift,alpha_shift	{-31, ..., 31}

## 3.3.8 Leaky ReLU

### Description

These kernels compute the Leaky ReLU function of input vector  $x$  and give output vector as  $y = \text{leaky\_relu}(x)$ . Both the input and output vectors have size `vec_length`.

The Leaky ReLU activation function acts like a standard ReLU function for input values greater than or equal to 0. For input values less than 0, a negative slope parameter  $\alpha$  is multiplied with input to get the output. The slope value is constant for all the input elements.

Function variants available are `xa_nn_vec_leaky_relu_[p]_[q]`, where:

- [p]: Input precision in bits
- [q]: Output precision in bits

### Precision

There is one variant available:

Type	Description
asym8s_asym8s	asym8s input, asym8s output

### Algorithm

$$y_n = x_n, \quad \text{when } x_n \geq 0 \quad n = 0, \dots, \overline{vec\_length} - 1$$

$$y_n = \alpha x_n, \quad \text{when } x_n < 0$$

where  $\alpha$  is the negative slope parameter:  $\alpha$ .

### Prototype

```
WORD32 xa_nn_vec_leaky_relu_asym8s_asym8s
(WORD8 * p_out, const WORD8 * p_vec, WORD32 inp_zero_bias,
```



```
WORD32 alpha_multiplier, WORD32 alpha_shift, WORD32 out_multiplier,
WORD32 out_shift, WORD32 out_zero_bias, WORD32 vec_length);
```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
const WORD8 *	p_vec	vec_length	Input vector, asym8s
WORD32	inp_zero_bias		Zero bias value for input vector
WORD16	alpha_multiplier		Fixed-point multiplier value for alpha input.
WORD32	alpha_shift		Shift value for alpha input.
WORD16	out_multiplier		Fixed-point multiplier value for output
WORD32	out_shift		Shift value for output
WORD32	out_zero_bias		Zero bias value for output vector
WORD32	vec_length		length of input vector
<b>Output</b>			
WORD8 *	p_out	vec_length	Output vector, asym8s

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions

Arguments	Restrictions
p_vec, p_out	Cannot be NULL
	Should not overlap (the two pointers could be same, inplace operation is possible)
inp_zero_bias	{-127....., 128} for asym8s datatype
out_zero_bias	{-128....., 127} for asym8s datatype
out_multiplier, alpha_multiplier	Shouldn't be less than 0
out_shift, alpha_shift	{-31, ..., 31}

## 3.4 Pooling Kernels

### 3.4.1 Average Pool Kernel

#### Description

These kernels compute 2D average pool on a set of input planes (matrices)  $x$  and give a set of planes  $y$  as output.

The pooling region is defined by `kernel_height` and `kernel_width`. It is shifted over the input plane in steps of `x_stride` horizontally and in steps of `y_stride` vertically to generate the specified output plane size. The input is extended by zero padding as specified by the padding region. The padding is determined by the parameters `x_padding`, `y_padding` for left and top side padding respectively, and `out_width`, `out_height` for right and bottom padding respectively. Around the edges of input planes, if only a part of pooling region is covering input plane then only the average of those elements is calculated, and the denominator is the number of elements from input in current pooling region.

The average pool kernels accept input as 8-bit, 16-bit integer, `asym8u` or single precision floating point format and give output in the same precision as input.

These kernels require temporary buffer for average pool computation. This temporary buffer is provided by the `p_scratch` argument of kernel API. The size of the temporary buffer should be queried using `xa_nn_avgpool_getsize()` helper API.

These kernels expect input cube in `SHAPE_CUBE_WHD_T` and `SHAPE_CUBE_DWH_T` shape type and produce output cube in `SHAPE_CUBE_WHD_T` and `SHAPE_CUBE_DWH_T` shape type respectively. The `inp_data_format` and `out_data_format` arguments to the kernel API can be 0 or 1 to indicate input and output cube shapes respectively.

The value of `inp_data_format` and `out_data_format` must be equal.

Note, the fixed-point 8-bit average pool kernel, `xa_nn_avgpool_8` can be used for the quantized `int8` datatype.

Function variants available are `xa_nn_avgpool_[p]`, where:

- `[p]`: Input and Output precision in bits

#### Precision

There are four variants available:

Type	Description
8	8-bit input, 8-bit output
16	16-bit input, 16-bit output
f32	float32 input, float32 output
asym8u	asym8u input, asym8u output

## Algorithm

$$z_{h,w,d} = \frac{1}{K_H K_W} \left( \sum_{i=0}^{K_H-1} \sum_{j=0}^{K_W-1} in_{(h*y-stride+i),(w*x-stride+j),d} \right)$$

$h = 0, \dots, \overline{out-height - 1}, w = 0, \dots, \overline{out-width - 1},$   
 $d = 0, \dots, \overline{out-channels - 1}$

$in$  denotes padded input cube,  $z$  denotes output

$K_H, K_W$  denote kernel\_height, kernel\_width respectively.

## Prototype

```
WORD32 xa_nn_avgpool_getsize
(WORD32 input_channels, WORD32 inp_precision, WORD32 out_precision,
 WORD32 input_height, WORD32 input_width, WORD32 kernel_height,
 WORD32 kernel_width, WORD32 x_stride, WORD32 y_stride,
 WORD32 x_padding, WORD32 y_padding, WORD32 out_height,
 WORD32 out_width, WORD32 inp_data_format, WORD32 out_data_format);

WORD32 xa_nn_avgpool_8
(WORD8 * p_out, const WORD8 * p_inp, WORD32 input_height,
 WORD32 input_width, WORD32 input_channels, WORD32 kernel_height,
 WORD32 kernel_width, WORD32 x_stride, WORD32 y_stride,
 WORD32 x_padding, WORD32 y_padding, WORD32 out_height,
 WORD32 out_width, WORD32 inp_data_format, WORD32 out_data_format,
 VOID * p_scratch);

WORD32 xa_nn_avgpool_16
(WORD16 * p_out, const WORD16 * p_inp, WORD32 input_height,
 WORD32 input_width, WORD32 input_channels, WORD32 kernel_height,
 WORD32 kernel_width, WORD32 x_stride, WORD32 y_stride,
 WORD32 x_padding, WORD32 y_padding, WORD32 out_height,
 WORD32 out_width, WORD32 inp_data_format, WORD32 out_data_format,
 VOID * p_scratch);

WORD32 xa_nn_avgpool_f32
(FLOAT32 * p_out, const FLOAT32 * p_inp, WORD32 input_height,
 WORD32 input_width, WORD32 input_channels, WORD32 kernel_height,
 WORD32 kernel_width, WORD32 x_stride, WORD32 y_stride,
 WORD32 x_padding, WORD32 y_padding, WORD32 out_height,
 WORD32 out_width, WORD32 inp_data_format, WORD32 out_data_format,
 VOID * p_scratch);

WORD32 xa_nn_avgpool_asym8u
(UWORD8 * p_out, const UWORD8 * p_inp, WORD32 input_height,
 WORD32 input_width, WORD32 input_channels, WORD32 kernel_height,
 WORD32 kernel_width, WORD32 x_stride, WORD32 y_stride,
 WORD32 x_padding, WORD32 y_padding, WORD32 out_height,
 WORD32 out_width, WORD32 inp_data_format, WORD32 out_data_format,
 VOID * p_scratch);
```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
const WORD8 *, const WORD16 *,	p_inp	input_height * input_width * input_channels	Input cube

const UWORD8 *, const FLOAT32 *			
WORD32	input_height		Input height
WORD32	input_width		Input width
WORD32	input_channels		Input number of channels
WORD32	kernel_height		Pooling window height
WORD32	kernel_width		Pooling window width
WORD32	x_stride		Horizontal stride over input
WORD32	y_stride		Vertical stride over input
WORD32	x_padding		Left padding width on input
WORD32	y_padding		Top padding height on input
WORD32	out_height		Output height
WORD32	out_width		Output width
WORD32	inp_data_format		Input data format 0:SHAPE_CUBE_DWH_T 1:SHAPE_CUBE_WHD_T
WORD32	out_data_format		Output data format: 0:SHAPE_CUBE_DWH_T 1:SHAPE_CUBE_WHD_T
<b>Output</b>			
WORD8 *, WORD16 *, UWORD8 *, FLOAT32 *	p_out	out_height * out_width * input_channels	Output cube
<b>Temporary</b>			
VOID *	p_scratch	xa_nn_avgpool_ getsize()	Temporary / scratch memory

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions

Arguments	Restrictions
<code>p_inp, p_out</code>	Cannot be NULL Should not overlap
<code>p_scratch</code>	Cannot be NULL Should not overlap Memory size $\geq$ size returned by <code>xa_nn_avgpool_getsize()</code>
<code>input_height, input_width</code>	Greater than or equal to 1
<code>input_channels</code>	Greater than or equal to 1
<code>kernel_height</code>	$\{1, 2, \dots, \min(\text{input\_height}, 256)\}$ (for 8-bit and 16-bit) $\{1, 2, \dots, \text{input\_height}\}$ (for float32)
<code>kernel_width</code>	$\{1, 2, \dots, \min(\text{input\_width}, 256)\}$ (for 8-bit and 16-bit) $\{1, 2, \dots, \text{input\_width}\}$ (for float32)
<code>x_stride, y_stride</code>	Greater than or equal to 1
<code>x_padding, y_padding</code>	Greater than or equal to 0
<code>out_height, out_width</code>	greater than or equal to 1
<code>inp_data_format</code>	Can be 0: SHAPE_CUBE_DWH_T or 1: SHAPE_CUBE_WHD_T
<code>out_data_format</code>	Must be equal to <code>inp_data_format</code>

### 3.4.2 Max Pool Kernel

#### Description

These kernels perform 2D max pooling operation over a set of input planes `x` and give as output, a set of planes `y`.

The pooling region is defined by `kernel_height` and `kernel_width`. It is shifted over the input plane horizontally in steps of `x_stride` and vertically in steps of `y_stride` to generate the specified output plane size.

The input plane, padded with the maximum negative values is considered while performing the max pooling operation. The padding region is determined by the parameters `x_padding`, `y_padding` for left and top side padding respectively, and `out_width`, `out_height` for right and bottom padding respectively.

The max pool kernels accept input as 8-bit, 16-bit integer, `asym8u` or single precision floating point format and give output in the same precision as input.

These kernels require temporary buffer for max pool computation. This temporary buffer is provided by the `p_scratch` argument of kernel API. The size of the temporary buffer should be queried using the `xa_nn_maxpool_getsize()` helper API.

These kernels expect input cube in SHAPE\_CUBE\_WHD\_T and SHAPE\_CUBE\_DWH\_T shape type and produce output cube in SHAPE\_CUBE\_WHD\_T and SHAPE\_CUBE\_DWH\_T shape type respectively. The

`inp_data_format` and `out_data_format` arguments to the kernel API can be 0 or 1 to indicate input and output cube shapes respectively.

The value of `inp_data_format` and `out_data_format` must be equal.

Note, the fixed-point 8-bit max pool kernel, `xa_nn_maxpool_8` can be used for the quantized int8 datatype.

Function variants available are `xa_nn_maxpool_[p]`, where:

- `[p]`: Input and Output precision in bits

## Precision

There are four variants available:

Type	Description
8	8-bit input, 8-bit output
16	16-bit input, 16-bit output
f32	float32 input, float32 output
asym8u	asym8u input, asym8u output

## Algorithm

$$z_{h,w,d} = \max(in_{(h \times y\text{-stride} + i), (w \times x\text{-stride} + j), d})$$

$$h = 0, \dots, out\text{-height} - 1, \quad w = 0, \dots, out\text{-width} - 1,$$

$$d = 0, \dots, out\text{-channels} - 1$$

$$i = 0, \dots, K_H - 1, \quad j = 0, \dots, K_W - 1$$

`in` denotes padded input cube, `z` denotes output.

$K_H, K_W$  denote `kernel_height`, `kernel_width` respectively.

## Prototype

```

WORD32 xa_nn_maxpool_getsize
(WORD32 input_channels, WORD32 inp_precision, WORD32 out_precision,
 WORD32 input_height, WORD32 input_width, WORD32 kernel_height,
 WORD32 kernel_width, WORD32 x_stride, WORD32 y_stride,
 WORD32 x_padding, WORD32 y_padding, WORD32 out_height,
 WORD32 out_width, WORD32 inp_data_format, WORD32 out_data_format);

WORD32 xa_nn_maxpool_8
(WORD8 * p_out, const WORD8 * p_inp, WORD32 input_height,
 WORD32 input_width, WORD32 input_channels, WORD32 kernel_height,
 WORD32 kernel_width, WORD32 x_stride, WORD32 y_stride,
 WORD32 x_padding, WORD32 y_padding, WORD32 out_height,
 WORD32 out_width, WORD32 out_data_format,
 VOID * p_scratch);
WORD32 xa_nn_maxpool_16
(WORD16 * p_out, const WORD16 * p_inp, WORD32 input_height,
 WORD32 input_width, WORD32 input_channels, WORD32 kernel_height,
 WORD32 kernel_width, WORD32 x_stride, WORD32 y_stride,
 WORD32 x_padding, WORD32 y_padding, WORD32 out_height,
 WORD32 out_width, WORD32 out_data_format,
 VOID * p_scratch);
WORD32 xa_nn_maxpool_f32
(FLOAT32 * p_out, const FLOAT32 * p_inp, WORD32 input_height,
 WORD32 input_width, WORD32 input_channels, WORD32 kernel_height,
 WORD32 kernel_width, WORD32 x_stride, WORD32 y_stride,
 WORD32 x_padding, WORD32 y_padding, WORD32 out_height,
 WORD32 out_width, WORD32 out_data_format,
 VOID * p_scratch);
WORD32 xa_nn_maxpool_asym8u
(UWORD8* p_out, const UWORD8* p_inp, WORD32 input_height,
 WORD32 input_width, WORD32 input_channels, WORD32 kernel_height,
 WORD32 kernel_width, WORD32 x_stride, WORD32 y_stride,
 WORD32 x_padding, WORD32 y_padding, WORD32 out_height,
 WORD32 out_width, WORD32 inp_data_format, WORD32 out_data_format,
 VOID *p_scratch);

```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
const WORD8 *, const WORD16 *, const UWORD8 *, const FLOAT32 *	p_inp	input_height * input_width * input_channels	Input cube
WORD32	input_height		Input height
WORD32	input_width		Input width
WORD32	input_channels		Input number of channels
WORD32	kernel_height		Pooling window height
WORD32	kernel_width		Pooling window width
WORD32	x_stride		Horizontal stride over input
WORD32	y_stride		Vertical stride over input
WORD32	x_padding		Left padding width on input
WORD32	y_padding		Top padding height on input

WORD32	out_height		Output height
WORD32	out_width		Output width
WORD32	inp_data_format		Input data format 0:SHAPE_CUBE_DWH_T 1:SHAPE_CUBE_WHD_T
WORD32	out_data_format		Input data format 0:SHAPE_CUBE_DWH_T 1:SHAPE_CUBE_WHD_T
<b>Output</b>			
WORD8 *, WORD16 *, UWORD8 *, FLOAT32 *	p_out	out_height * out_width * input_channels	Output cube
<b>Temporary</b>			
VOID *	p_scratch	xa_nn_maxpool_ getsize()	Temporary / scratch memory

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions

Arguments	Restrictions
p_inp, p_out	Cannot be NULL Should not overlap Should not overlap Memory size $\geq$ size returned by xa_nn_maxpool_getsize()
input_height, input_width	Greater than or equal to 1
input_channels	Greater than or equal to 1
kernel_height	{1, 2, ..., input_height}
kernel_width	{1, 2, ..., input_width}
x_stride, y_stride	Greater than or equal to 1
x_padding, y_padding	Greater than or equal to 0
out_height, out_width	Greater than or equal to 1
inp_data_format	Can be 0: SHAPE_CUBE_DWH_T or 1: SHAPE_CUBE_WHD_T
out_data_format	Must be equal to inp_data_format



## 3.5 Fully connected Layer

### 3.5.1 Fully Connected Kernel

#### Description

These kernels perform the operation of multiplication of weight matrix with input vectors in a fully connected neural network layer i.e.  $z = \text{weight} * \text{input} + \text{bias}$ . The column dimension of `weight` must match the row dimension of `input`. Bias and resulting output vector `z` have as many numbers of rows as `weight` matrix.

The `bias_shift` and `acc_shift` arguments are provided in kernel API to adjust Q format of bias and output, respectively. Both `bias_shift` and `acc_shift` can be either positive or negative, where positive value denotes a left shift and negative value denotes a right shift.

`bias_shift` is the shift in number of bits applied to the bias to make it in the same Q format as weight X input multiplication – accumulation result. `acc_shift` is the shift in number of bits applied to the accumulator to obtain the output in desired Q format.

**Note:** `acc_shift` and `bias_shift` are not relevant in the case of floating point and quantized 8-bit kernels.

For conversion from higher precision accumulator to lower precision output, symmetric rounding is used.

The precision of output is the same as precision of input vector.

The arguments `input_zero_bias`, `weight_zero_bias` are provided to convert the quantized 8-bit inputs into their real values and perform Fully Connected kernel operation. The `out_zero_bias`, `out_multiplier` and `out_shift` values are used to quantize real values of output back to 8-bit.

Function variants available (for fixed point) are `xa_nn_fully_connected_[p]x[q]_[r]`, where:

- `[p]`: Weight matrix precision in bits
- `[q]`: Input vector precision in bits
- `[r]`: Output vector precision in bits

#### Precision

There are six variants available:

Type	Description
16x16_16	16-bit weight matrix, 16-bit input vector, 16-bit output
8x16_16	8-bit weight matrix, 16-bit input vector, 16-bit output
8x8_8	8-bit weight matrix, 8-bit input vector, 8-bit output
f32	float32 weight matrix, float32 input vector, float32 output
asym8uxasym8u_asym8u	asym8u weight matrix, asym8u input vector, asym8u output
sym8sxsym8s_asym8s	sym8s weight matrix, asym8s input vector, asym8s output

## Algorithm

$$z_n = 2^{acc-shift} \left( \sum_{m=0}^{W_D-1} weight_{n,m} \cdot input_m + 2^{bias-shift} bias_n \right),$$

$n = 0, \dots, out\_depth - 1$

where  $W_D$  represents weight\_depth

For floating point and quantized 8-bit routines, acc\_shift=0 and bias\_shift=0

Thus,  $2^{acc-shift} = 2^{bias-shift} = 1$

## Prototype

```
WORD32 xa_nn_fully_connected_16x16_16
(WORD16 * p_out,          WORD16 * p_weight,          WORD16 * p_inp,
 WORD16 * p_bias,         WORD32 weight_depth,        WORD32 out_depth,
 WORD32 acc_shift,        WORD32 bias_shift);
WORD32 xa_nn_fully_connected_8x16_16
(WORD16 * p_out,          WORD8 * p_weight,           WORD16 * p_inp,
 WORD16 * p_bias,         WORD32 weight_depth,        WORD32 out_depth,
 WORD32 acc_shift,        WORD32 bias_shift);
WORD32 xa_nn_fully_connected_8x8_8
(WORD8 * p_out,           WORD8 * p_weight,           WORD8 * p_inp,
 WORD8 * p_bias,         WORD32 weight_depth,        WORD32 out_depth,
 WORD32 acc_shift,        WORD32 bias_shift);
WORD32 xa_nn_fully_connected_f32
(FLOAT32 * p_out,         FLOAT32 * p_weight,         FLOAT32 * p_inp,
 FLOAT32 * p_bias,        WORD32 weight_depth,        WORD32 out_depth);
WORD32 xa_nn_fully_connected_asym8uxasym8u_asym8u
(UWORD8 * p_out,          const UWORD8 * p_weight,    const UWORD8 * p_inp,
 const WORD32 * p_bias,    WORD32 weight_depth,        WORD32 out_depth,
 WORD32 input_zero_bias,  WORD32 weight_zero_bias,  WORD32 out_multiplier,
 WORD32 out_shift,        WORD32 out_zero_bias);
WORD32 xa_nn_fully_connected_sym8sxasym8s_asym8s
(WORD8 * p_out,           const WORD8 * p_weight,     const WORD8 * p_inp,
 const WORD32 * p_bias,    WORD32 weight_depth,        WORD32 out_depth,
 WORD32 input_zero_bias,  WORD32 out_multiplier,  WORD32 out_shift,
 WORD32 out_zero_bias);
```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
WORD16 *, WORD8 *, pFLOAT32, const UWORD8 *, const WORD8 *	p_weight	out_depth* weight_depth	Weight matrix, fixed, floating point, asym8u or sym8s
WORD16 *, WORD8 *, pFLOAT32, const UWORD8 *	p_inp	weight_depth* 1	Input vector, fixed, floating point, asym8u or asym8s

const WORD8 *			
WORD16 *, WORD8 *, pFLOAT32, WORD32 *	p_bias	out_depth*1	Bias vector, fixed or floating point, 32-bit for quantized kernels
WORD32	out_depth		Number of rows in weight matrix, bias and output vector
WORD32	weight_depth		Number of columns in weight matrix and rows in input vector
WORD32	acc_shift		Shift applied to accumulator
WORD32	bias_shift		Shift applied to bias
WORD32	input_zero_bias		Zero offset of input
WORD32	weight_zero_bias		Zero offset of weights
WORD32	out_multiplier		Multiplier value of output
WORD32	out_shift		Shift value of output
WORD32	out_zero_bias		Zero offset of output
<b>Output</b>			
WORD8 *, WORD16 *, pFLOAT32, WORD8 *, UWORD8*	p_out	out_depth*1	Output vector, fixed, floating point, asym8u or asym8s

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions

Arguments	Restrictions
weight_depth	Multiple of 4 for fixed point and floating point kernels. No restriction for quantized 8-bit kernels.
p_weight, p_inp, p_bias, p_out	Aligned on 16-byte boundary, should not overlap.. (size of one element)-byte boundary in case of floating point and quantized 8-bit kernels).
p_weight, p_inp, p_out	Cannot be NULL
out_depth	Greater than or equal to 1
acc_shift, bias_shift, out_shift	{-31, ..., 31}
input_zero_bias	{-255, ..., 0} for asym8u, {-127, ..., 128} for asym8s
weight_zero_bias	{-255, ..., 0} for asym8u
out_zero_bias	{0, ..., 255} for asym8u, {-128, ..., 127} for asym8s

## 3.6 Basic Operations and Miscellaneous Kernels

### 3.6.1 Interpolation Kernel

#### Description

This kernel performs interpolation between two input vectors  $h$  and  $y$  using interpolation factor from vector  $x$  to get output vector  $z$ .

The interpolation kernel accepts 16-bit inputs and 16-bit interpolation factor in Q15 format and produces 16-bit output in Q15 format.

#### Precision

Type	Description
16-bit	16-bit input, 16-bit interpolation factor, 16-bit output

#### Algorithm

$$z_n = x_n * y_n + (1 - x_n) * h_n, \quad n = 0 \dots, \text{num\_elements} - 1$$

$x_n$  represents interpolation factor.

$y_n$  represents first input,  $h_n$  represents second input.

$z_n$  represents output.

#### Prototype

```
WORD32 xa_nn_vec_interpolation_q15(
    (WORD16 * p_out,      WORD16 * p_ifact,      WORD16 * p_inp1, WORD16 * p_inp2,      WORD32
    num_elements);
```

#### Arguments

Type	Name	Size	Description
<b>Input</b>			
WORD16 *	p_ifact	num_elements	Interpolation factor vector
WORD16 *	p_inp1	num_elements	First input vector
WORD16 *	p_inp2	num_elements	Second input vector
WORD32	num_elements		Number of elements
<b>Output</b>			
WORD16 *	p_out	num_elements	Output vector

#### Returns

- 0: no error

- -1: error, invalid parameters

## Restrictions

Arguments	Restrictions
p_ifact, p_inp1, p_inp2, p_out	Aligned on 8-byte boundary
	Should not overlap
	Cannot be NULL
num_elements	Multiple of 4

## 3.6.2 Dot Product Kernels

### Description

These kernels perform the dot product operations between two sets of input vectors `p_inp1` and `p_inp2` to get output vector `p_out`. The supported precisions are: `f32xf32_f32` and `16x16_asym8s`.

Function variants available are `xa_nn_elm_quantize_[p]x[q]_[r]`, where:

- `[p], [q]`: Input precision
- `[r]`: Output precision

### Precision

There are two variants available:

Type	Description
<code>f32xf32_f32</code>	float32 input, float32 output
<code>16x16_asym8s</code>	16-bit input, asym8s output

### Algorithm

$$z_n = \left( \sum_{m=0}^{vec\_length-1} inp1_m \cdot inp2_m + bias_n \right)$$

$$n = 0, \dots, \overline{vec\_count} - 1$$

### Prototype

```
WORD32 xa_nn_dot_prod_f32xf32_f32
(FLOAT32 * p_out,          const FLOAT32 * p_inp1,          const FLOAT32 * p_inp2,
 WORD32 vec_length,        WORD32 num_vecs);
WORD32 xa_nn_dot_prod_16x16_asym8s
(WORD8 * p_out,            const WORD16 * p_inp1,            const WORD16 * p_inp2,
 const WORD32 * bias_ptr,   WORD32 vec_length,              WORD32 out_multiplier,
 WORD32 out_shift,          WORD32 out_zero_bias,            WORD32 vec_count);
```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
const FLOAT32 * const WORD16 *	p_inp1	vec_length	First input vector
const FLOAT32 * const WORD16 *	p_inp2	vec_length	Second input vector
const WORD32 *	bias_ptr	vec_count	Bias vector
WORD32	vec_length		Length of each vector
WORD32	out_multiplier		Multiplier value of output
WORD32	out_shift		Shift value of output
WORD32	out_zero_bias		Zero offset of output
WORD32	num_vecs, vec_count		Number of input vectors
<b>Output</b>			
FLOAT32 * WORD8 *	p_out	num_vecs	Output vector

## Returns

- 0: no error
- 1: error, invalid parameters

## Restrictions:

Arguments	Restrictions
p_inp1, p_inp2, p_out	Aligned on (size of one element)-byte boundary Cannot be NULL
vec_length, num_vecs	Greater than 0
out_shift	{-31, ..., 31}
out_multiplier	Greater than 0
out_zero_bias	{-128, ..., 127} for out type asym8s

## 3.6.3 Elementwise Requantize Kernels

### Description

These kernels perform the requantization operation of the `p_inp1` input vector elements to get the output vector `p_out`. The kernels are developed in reference to the Quantize operator implementation in TensorFlow Lite Micro.

Function variants available are `xa_nn_elm_requantize_[p]_[q]`, where:

- [p]: Input precision
- [q]: Output precision

## Algorithm

```
for itr = 0:(num_elm-1)
    p_out[itr] = ((2^out_shift) * (out_multiplier) * (p_inp[itr] - inp_zero_bias)) + out_zero_bias
```

## Precision

Type	Description
asym16s_asym8s	asym16s input, asym8s output

## Prototype

```
WORD32 xa_nn_elm_requantize_asym16s_asym8s
(WORD8 * __restrict__ p_out, const WORD16 * __restrict__ p_inp, WORD32 inp_zero_bias,
 WORD32 out_zero_bias, WORD32 out_shift, WORD32 out_multiplier,
 WORD32 num_elm);
```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
const WORD16 *, const WORD8 *	p_inp	num_elm	Input vector
WORD32	inp_zero_bias		Zero offset of input
WORD32	out_zero_bias		Zero offset of output
WORD32	out_shift		Shift value of output
WORD32	out_multiplier		Multiplier value of output
WORD32	num_elm		Number of input elements
<b>Output</b>			
WORD8 *, WORD32 *	p_out	num_elm	Output vector

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions:

Arguments	Restrictions
p_inp, p_out	Aligned on (size of one element)-byte boundary
	Cannot be NULL
	Should not overlap
num_elm	Greater than 0
out_shift	{-31, ..., 31}
out_multiplier	Greater than 0
inp_zero_bias	{-32768, ..., 32767} for inp type asym16s {-128, ..., 127} for inp type asym8s
out_zero_bias	{-128, ..., 127} for out type asym8s

	Signed 32-bit integer value for out type asym32s
--	--

## 3.6.4 Elementwise Dequantize Kernels

### Description

These kernels perform the dequantization operation of the `p_inp1` input vector elements to get the output vector `p_out`. The kernels are developed in reference to the Dequantize operator implementation in Tensorflow Lite Micro.

Function variants available are `xa_nn_elm_dequantize_[p]_[q]`, where:

- `[p]`: Input precision
- `[p]`: Output precision

### Precision

Type	Description
<code>asym8s_f32</code>	asym8s input, float output

### Algorithm

for `itr = 0:(num_elm-1)`

$$p\_out[itr] = (p\_inp[itr] - inp\_zero\_bias) * inp\_scale$$

### Prototype

```
WORD32 xa_nn_elm_dequantize_asym8s_f32
(FLOAT32 * __restrict__ p_out, const WORD8 * __restrict__ p_inp, WORD32 inp_zero_bias,
 FLOAT32 inp_scale, WORD32 num_elm);
```

### Arguments

Type	Name	Size	Description
<b>Input</b>			
<code>const WORD8 *</code>	<code>p_inp</code>	<code>num_elm</code>	Input vector
<code>WORD32</code>	<code>inp_zero_bias</code>		Zero offset of input
<code>FLOAT32</code>	<code>inp_scale</code>		Input scale
<code>WORD32</code>	<code>num_elm</code>		Number of input elements
<b>Output</b>			
<code>FLOAT32 *</code>	<code>p_out</code>	<code>num_elm</code>	Output vector

### Returns

- 0: no error
- -1: error, invalid parameters



**Restrictions:**

Arguments	Restrictions
p_inp, p_out	Aligned on (size of one element)-byte boundary
	Cannot be NULL
	Should not overlap
num_elm	Greater than 0
inp_zero_bias	{-128....,127} for inp type asym8s

### 3.6.5 Basic Kernels

**Description**

These kernels perform basic elementwise operations on one or two input vectors  $x$  and  $y$  to get output vector  $z$ . The supported operations are: add, subtract, multiply, floor, minimum, maximum, sine, cosine, log (natural), absolute, ceil, round (banker's), negative, square, square-root and inverse square-root. The supported precisions are: 8-bit, float32 and asym8s.

The 8-bit elementwise minimum and maximum kernels can be also used for asym8s datatype.

Function variants available are `xa_nn_[o]_[p]_[q]`, where:

- `[o]`: Operations: `elm_add`, `elm_sub`, `elm_mul`, `elm_floor`, `elm_min`, `elm_max`, `elm_sine`, `elm_cosine`, `elm_logn`, `elm_abs`, `elm_ceil`, `elm_round`, `elm_neg`, `elm_square`, `elm_sqrt`, `elm_rsqrt`
- `[p]`: Input Precision in bits- input1input2 or input1
- `[q]`: Output Precision in bits

**Precision**

Type	Description
f32xf32_f32	2 float32 inputs, float32 output
f32_f32	float32 input, float32 output
8x8_8	2 8-bit input, 8-bit output
asym8sxasym8s_asym8s	2 asym8s inputs, asym8s output

**Algorithm**

$$\begin{array}{lll}
 \text{elm\_floor:} & z_n = \lfloor x_n \rfloor, & n = 0 \dots, \text{num\_elm} - 1 \\
 \text{elm\_sine:} & z_n = \sin(x_n), & n = 0 \dots, \text{num\_elm} - 1 \\
 \text{elm\_cosine:} & z_n = \cos(x_n), & n = 0 \dots, \text{num\_elm} - 1 \\
 \text{elm\_logn:} & z_n = \log_e(x_n), & n = 0 \dots, \text{num\_elm} - 1 \\
 \text{elm\_abs:} & z_n = \text{abs}(x_n), & n = 0 \dots, \text{num\_elm} - 1
 \end{array}$$

elm_ceil:	$z_n = \lceil x_n \rceil,$	$n = 0 \dots, \overline{num\_elm - 1}$
elm_round <sup>6</sup> :	$z_n = \text{round}(x_n),$	$n = 0 \dots, \overline{num\_elm - 1}$
elm_neg:	$z_n = -x_n,$	$n = 0 \dots, \overline{num\_elm - 1}$
elm_square:	$z_n = x_n * x_n,$	$n = 0 \dots, \overline{num\_elm - 1}$
elm_sqrt:	$z_n = \sqrt{x_n},$	$n = 0 \dots, \overline{num\_elm - 1}$
elm_rsqrt:	$z_n = 1 \div \sqrt{x_n},$	$n = 0 \dots, \overline{num\_elm - 1}$

$x_n$  represents first input,  $y_n$  represents second input.

$z_n$  represents output.

## Prototype

```
WORD32 xa_nn_elm_floor_f32_f32
(FLOAT32 * p_out, const FLOAT32 * p_inp, WORD32 num_elm);
WORD32 xa_nn_elm_sine_f32_f32
(FLOAT32 * __restrict__ p_out, const FLOAT32 * __restrict__ p_inp, WORD32 num_elm);
WORD32 xa_nn_elm_cosine_f32_f32
(FLOAT32 * __restrict__ p_out, const FLOAT32 * __restrict__ p_inp, WORD32 num_elm);
WORD32 xa_nn_elm_logn_f32_f32
(FLOAT32 * __restrict__ p_out, const FLOAT32 * __restrict__ p_inp, WORD32 num_elm);
WORD32 xa_nn_elm_abs_f32_f32
(FLOAT32 * __restrict__ p_out, const FLOAT32 * __restrict__ p_inp, WORD32 num_elm);
WORD32 xa_nn_elm_ceil_f32_f32
(FLOAT32 * __restrict__ p_out, const FLOAT32 * __restrict__ p_inp, WORD32 num_elm);
WORD32 xa_nn_elm_round_f32_f32
(FLOAT32 * __restrict__ p_out, const FLOAT32 * __restrict__ p_inp, WORD32 num_elm);
WORD32 xa_nn_elm_neg_f32_f32
(FLOAT32 * __restrict__ p_out, const FLOAT32 * __restrict__ p_inp, WORD32 num_elm);
WORD32 xa_nn_elm_square_f32_f32
(FLOAT32 * __restrict__ p_out, const FLOAT32 * __restrict__ p_inp, WORD32 num_elm);
WORD32 xa_nn_elm_sqrt_f32_f32
(FLOAT32 * __restrict__ p_out, const FLOAT32 * __restrict__ p_inp, WORD32 num_elm);
WORD32 xa_nn_elm_add_f32xf32_f32
(FLOAT32 * __restrict__ p_out, const FLOAT32 * __restrict__ p_inpl,
const FLOAT32 * __restrict__ p_inp2, WORD32 num_elm);
WORD32 xa_nn_elm_rsqrt_f32_f32
(FLOAT32 * __restrict__ p_out, const FLOAT32 * __restrict__ p_inp, WORD32 num_elm);
```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
const WORD8 * FLOAT32 *	p_inpl, p_inp, p_inl	num_elm	First input vector
const WORD8 * FLOAT32 *	p_inp2, p_in2	num_elm	Second input vector
WORD32	num_elm/num_element		Number of elements
WORD32	out_zero_bias		Zero bias of output
WORD32	out_shift		Shift value of output
WORD32	out_multiplier		Multiplier value of output

<sup>6</sup> The round variant is banker's rounding. It is also called as "Round half to even". In this rounding method, if fractional part of input is 0.5, then output is the even integer nearest to input. Thus, for example, +23.5 becomes 24, as does 24.5; while -23.5 becomes -24, as does -24.5

WORD32	out_activation_min		Activation min of output
WORD32	out_activation_max		Activation max of output
WORD32	inp1_zero_bias		Zero bias of input 1
WORD32	inp1_shift		Shift value of input 1
WORD32	inp1_multiplier		Multiplier value of input 1
WORD32	inp2_zero_bias		Zero bias of input 2
WORD32	inp2_shift		Shift value of input 2
WORD32	inp2_multiplier		Multiplier value of input 2
WORD32	left_shift		Global left shift value for inputs.
<b>Output</b>			
WORD8 *	p_out	num_elm	Output vector
FLOAT32 *			

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions:

Arguments	Restrictions
p_inp1, p_inp2, p_inp, p_in1, p_in2 p_out	Aligned on (size of one element)-byte boundary Cannot be NULL
p_out	Should not overlap with the input pointers (could be same as one of the input pointers, inplace operation is possible)
num_elm, num_element	Greater than 0
inp1_zero_bias, inp2_zero_bias	{-128....., 128} for asym8s input
inp1_shift, inp2_shift, out_shift	{-31 .... 31} for fixed point and quantized 8-bit APIs
left_shift	{0 .... 31}
inp1_multiplier, inp2_multiplier out_multiplier	Shouldn't be less than 0.
out_zero_bias	{-128....., 127} for asym8s output
out_activation_min, out_activation_max	{-128....., 127} for asym8s output out_activation_min < out_activation_max

## 3.6.6 Broadcast Kernel

### Description

This kernel broadcasts an input shape into the specified output shape. The input and output shapes must be compatible for the broadcast operation to succeed.

Details of the broadcast operation can be found at [Tensorflow Broadcasting semantics](#) <sup>[4]</sup>.

The dimensions of input and output tensors are passed as `in_shape` and `out_shape` and the number of dimensions specified by `numDims` must be the same for both. In case, the number of input and output dimensions are unequal, the empty leading dimensions of the smaller shape must be filled with ones to equalize them. For example, if the input dimension is 2x1x3 and the output dimension is 4x2x5x3, then `in_shape` must be passed as 1x2x1x3.

A simple illustration for broadcasting a 1x4x1 tensor into 1x4x3 and 2x4x3 is shown below.




Figure 3-1 Broadcasting a 1x4x1 tensor to 1x4x3 and 2x4x3

### Precision

Type	Description
8_8	8-bit input, 8-bit output

### Prototype

```
WORD32 xa_nn_broadcast_8_8
(WORD8* __restrict__ p_out, const int* const out_shape,
 const WORD8* __restrict__ p_in, const int* const in_shape,
 int numDims);
```

### Arguments

Type	Name	Size	Description
Input			

const WORD8 *	p_in	$\prod_{i=0}^{i=num\_dims-1} in\_shape[i]$	Input tensor
const int * const	in_shape out_shape	num_dims	Input/output shapes
int	num_dims	-	Number of dimensions
<b>Output</b>			
WORD8 *	p_out	$\prod_{i=0}^{i=num\_dims-1} out\_shape[i]$	Output tensor

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions:

Arguments	Restrictions
p_in, p_out	Aligned on (size of one element)-byte boundary Cannot be NULL
inp_shape, out_shape	Aligned on 4-byte boundary Cannot be NULL All elements should be greater than zero inp_shape[i] should be either equal to out_shape[i] or 1 for i = [0, numDims-1]
num_dims	In the range [1, 8]

## 3.6.7 Memory Operation Kernels

### Description

These kernels perform basic memory related operations. The supported operations are: memmove and memset. The supported precisions are: 8-bit, float32.

Memmove kernel does byte-level transfer and takes generic pointers, num\_elm should be set to number of 1-byte elements or simply number of bytes to be transferred for data types with sizes bigger than 1-byte.

Function variants available are xa\_nn\_[o]\_[p]\_[q], where:

- [o]: Operations: memmove, memset
- [p]: Input Precision in bits
- [q]: Output Precision in bits

## Precision

Type	Description
f32_f32	float32 input, float32 output
8_8	8-bit input, 8-bit output

## Algorithm

memmove:  $z_n = x_n$ ,  $n = 0 \dots, \overline{num\_elm} - 1$   
 memset:  $z_n = x_0$ ,  $n = 0 \dots, \overline{num\_elm} - 1$ ;  $x_0 < scalar >$

$x_n$  represents input

$z_n$  represents output.

## Prototype

```
WORD32 xa_nn_memset_f32_f32
(FLOAT32 * __restrict__ p_out, FLOAT32 val, WORD32 num_elm);
WORD32 xa_nn_memmove_8_8
(void * p_out, const void * p_inp, WORD32 num_elm);
```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
const FLOAT32 * void *	p_inp	num_elm	First input vector
FLOAT32	val		Memset value
WORD32	num_elm		Number of 1-byte elements or Number of bytes
<b>Output</b>			
FLOAT32 * void *	p_out	num_elm	Output vector

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions:

Arguments	Restrictions
p_inp, p_out	Aligned on (size of one element)-byte boundary Cannot be NULL
num_elm	Greater than 0

## 3.7 Normalization Kernels

### 3.7.1 L2 Normalization Kernel

#### Description

This kernel performs L2 normalization of an input vector  $x$  to get output vector  $z$ , which means every element of input vector  $x$  is divided by L2 norm of  $x$ , this gives an output vector  $z$  whose L2 norm is 1.

The L2 Normalization kernel accepts asym8s input vector and produces asym8s output vector.

#### Precision

Type	Description
asym8s	asym8s input, asym8s output

#### Algorithm

$$z_n = \frac{x_n}{\sqrt{\sum_{n=1}^N |x_n|^2}}, \quad n = 1 \dots, \overline{\text{num-elements}}$$

$x_n$  represents input vector.

$z_n$  represents output vector.

#### Prototype

```
WORD32 xa_nn_l2_norm_asym8s_asym8s
(WORD8 * p_out, const WORD8 * p_inp, WORD32 zero_point, WORD32 num_elm);
```

#### Arguments

Type	Name	Size	Description
<b>Input</b>			
const WORD8 *	p_inp	num_elm	Input vector
WORD32	zero_point		Input zero bias
WORD32	num_elm		Number of elements
<b>Output</b>			
WORD8 *	p_out	num_elm	Output vector

#### Returns

- 0: no error
- -1: error, invalid parameters

#### Restrictions

Arguments	Restrictions
-----------	--------------

p_inp, p_out	Aligned on (size of one element)-byte boundary
	Should not overlap
	Cannot be NULL
zero_point	{-128....., 127} for asym8s input
num_elm	Greater than 0

## 3.8 Reorg Kernels

### 3.8.1 Depth to Space Kernels

#### Description

These kernels convert the depth dimension of an input cube into the spatial dimensions of an output cube controlled by a block size parameter.

These kernels are based on DEPTH\_TO\_SPACE operator in TFLM<sup>[3]</sup>, which collects all elements from the input depth dimension and spreads it across the output spatial dimension using a `block_size` factor. The operation is illustrated below




Figure 3-2 Depth to space conversion for 4x4x8 input with block size of 2

Given an input cube of shape  $H \times W \times C$  and a `block_size` of  $K$ , this kernel will output cube of dimensions  $HK \times WK \times C/K^2$ . The specified output shape i.e `out_height/width/channels` must therefore equal  $HK$ ,  $WK$  and  $C/K^2$  respectively.



Since the elements collected from one dimension must be spread across two, the input depth dimension C (i.e. `input_channels`) must be divisible by  $K^2$  (i.e. `block_size^2`).

## Precision

Type	Description
8_8	8-bit input, 8-bit output

## Prototype

```
WORD32 xa_nn_depth_to_space_8_8
(pWORD8 __restrict__ p_out, const WORD8 * __restrict__ p_inp,
 WORD32 input_height, WORD32 input_width, WORD32 input_channels,
 WORD32 block_size,
 WORD32 out_height, WORD32 out_width, WORD32 out_channels,
 WORD32 inp_data_format, WORD32 out_data_format);
```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
const WORD8 *	p_inp	input_height* input_width* input_channels	Input cube data
WORD32	input_height		Input cube height
WORD32	input_width		Input cube width
WORD32	input_channels		Input cube channels
WORD32	block_size		Spatial dimension block size
WORD32	out_height		Output cube height
WORD32	out_width		Output cube width
WORD32	out_channels		Output cube channels
WORD32	inp_data_format		Input data format
WORD32	out_data_format		Output data format
<b>Output</b>			
WORD8 *	p_out	output_height* output_width* output_channels	Output cube data

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions

Arguments	Restrictions
p_inp, p_out	Aligned on (size of one element)-byte boundary
	Cannot be NULL
	Should not overlap

input_height	Must be greater than 0
input_width	Must be greater than 0
input_channels	Must be greater than 0 and divisible by <code>block_size<sup>2</sup></code>
block_size	Must be greater than 0
out_height	Must be <code>input_height*block_size</code>
out_width	Must be <code>input_width*block_size</code>
out_channels	Must be <code>input_channels/(block_size<sup>2</sup>)</code>
inp_data_format	Must be 0 (NHWC)
out_data_format	Must be 0 (NHWC)

## 3.8.2 Space to Depth Kernels

### Description

These kernels convert the spatial dimension of an input cube into the depth dimensions of an output cube controlled by a block size parameter.

These kernels perform the opposite operation of [depth to space kernels](#) which is illustrated in the figure below




Figure 3-3 Space to depth conversion for a 8x8x2 input with a block size of 2

Given an input of shape  $H \times W \times C$  with a `block_size` of  $K$ , this kernel will collect  $K \times K \times C$  elements from the input cube and serialize it into  $CK^2$  elements across the depth dimension of the output resulting in an output of shape  $(H/K) \times (W/K) \times (CK^2)$ .

The output shape specified i.e `out_height/width/channels` must equal  $H/K$ ,  $W/K$  and  $CK^2$  respectively.

Since the elements collected from in input 2D spatial dimension must be serialized into one output depth dimension, output\_channels specified must equal input\_channels\*block\_size<sup>2</sup>.

## Precision

Type	Description
8_8	8-bit input, 8-bit output

## Prototype

```
WORD32 xa_nn_space_to_depth_8_8
(pWORD8 __restrict__ p_out, const WORD8 * __restrict__ p_inp,
WORD32 input_height, WORD32 input_width, WORD32 input_channels,
WORD32 block_size,
WORD32 out_height, WORD32 out_width, WORD32 out_channels,
WORD32 inp_data_format, WORD32 out_data_format);
```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
const WORD8 *	p_inp	input_height* input_width* input_channels	Input cube data
WORD32	input_height		Input cube height
WORD32	input_width		Input cube width
WORD32	input_channels		Input cube channels
WORD32	block_size		Spatial dimension block size
WORD32	out_height		Output cube height
WORD32	out_width		Output cube width
WORD32	out_channels		Output cube channels
WORD32	inp_data_format		Input data format
WORD32	out_data_format		Output data format
<b>Output</b>			
WORD8 *	p_out	output_height* output_width* output_channels	Output cube data

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions

Arguments	Restrictions
p_inp, p_out	Aligned on (size of one element)-byte boundary
	Cannot be NULL
	Should not overlap

input_height	Must be greater than 0 and divisible by block_size
input_width	Must be greater than 0 and divisible by block_size
input_channels	Must be greater than 0
block_size	Must be greater than 0
out_height	Must be input_height/block_size
out_width	Must be input_width/block_size
out_channels	Must be input_channels*(block_size <sup>2</sup> )
inp_data_format	Must be 0 (NHWC)
out_data_format	Must be 0 (NHWC)

### 3.8.3 Pad Kernel

#### Description

This kernel pads an input with given `pad_value` according to the values specified in `p_pad_values`. `p_pad_values` is an integer array with size  $(2 * \text{input\_dimensions})$ , giving a pair of values for each input dimension. For each dimension of input, `p_pad_values` will contain a pair of values which will indicate how many values to add before the contents of input in that dimension and how many values to add after the contents of input in that dimension. This kernel is based on Pad and PadV2 operators in TFLM.

Input dimensions must be less than or equal to 4. 1/2/3-dimensional input will be scaled up to 4D. Output dimension must be equal to input dimension. Size of `p_pad_values` should be exactly  $(2 * \text{input\_dimensions})$ . The value to be padded can be given through `pad_value`.

Naming convention used for pad kernel is:

`xa_nn_pad_[p]`

Where `[p] = [input_precision]_[out_precision]`

#### Precision

Type	Description
8_8	Signed 8-bit input, signed 8-bit output

#### Algorithm

If

`ob = ib + p_pad_values[0]` ; `ib = [0, p_inp_shape[0]-1]`

`oh = ih + p_pad_values[2]`; `ih = [0, p_inp_shape[1]-1]`

`ow = iw + p_pad_values[4]`; `iw = [0, p_inp_shape[2]-1]`

`od = id + p_pad_values[6]`; `id = [0, p_inp_shape[3]-1]`

$$\text{Output}_{ob,oh,ow,od} = \text{Input}_{ib,ih,iw,id}$$

else

$$\text{Output}_{ob,oh,ow,od} = \text{pad\_value}$$

The shape of output after padding will be:

for  $D=0:(\text{num\_inp\_dims}-1)$

$$p\_out\_shape[D] = p\_pad\_values[2 * D] + p\_inp\_shape[D] + p\_pad\_values[2 * D + 1]$$

## Prototype

```
WORD32 xa_nn_pad_8_8
(WORD8 *__restrict__ p_out, const WORD32 *const p_out_shape,
const WORD8 *__restrict__ p_inp, const WORD32 *const p_inp_shape,
const WORD32 *__restrict__ p_pad_values, const WORD32 *const p_pad_shape,
WORD32 num_out_dims, WORD32 num_inp_dims, WORD32 num_pad_dims,
WORD32 pad_value);
```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
const WORD32 *const	p_out_shape	num_out_dims	Shape of output
const WORD8 *	p_inp	$\prod_{i=0}^{i=\text{num\_inp\_dims}-1} p\_inp\_shape[i]$	Input (set of cubes)
const WORD32 *const	p_inp_shape	num_inp_dims	Shape of input
const WORD32 *	p_pad_values	$\prod_{i=0}^{i=\text{num\_pad\_dims}-1} p\_pad\_shape[i]$	Pair of values (corresponds to before pad value and after pad value) for each input dimension
const WORD32 *const	p_pad_shape	num_pad_dims	Shape of pad_values
WORD32	num_out_dims		Number of output dimensions
WORD32	num_inp_dims		Number of input dimensions
WORD32	num_pad_dims		Number of pad dimensions
WORD32	pad_value		Value for padding
<b>Output</b>			
WORD8 *	p_out	$\prod_{i=0}^{i=\text{num\_out\_dims}-1} p\_out\_shape[i]$	Output (set of cubes)

## Returns

- 0: no error
- -1: error, invalid parameters

**Restrictions:**

Arguments	Restrictions
p_out, p_inp	Aligned on (size of one element)-byte boundary
	Cannot be NULL
	Should not overlap
p_out_shape, p_inp_shape, p_pad_shape	Aligned on 4-byte boundary
	Cannot be NULL
	Should not overlap
	All elements should be greater than zero
p_pad_values	Aligned on 4-byte boundary
	Cannot be NULL
	Should not overlap with other buffers
	All elements should be greater than or equal to zero
	Pair of values for each input dimension
num_out_dims	Must be in range [1, 4]
num_inp_dims	Must be in range [1, 4]
num_pad_dims	Must be in range [1, 4]
pad_value	Must be in range [-128, 127]

### 3.8.4 Batch to Space Kernels

**Description**

These kernels performs batch to space conversion on a set of input cube in (input\_batch x input\_height x input\_width x input\_depth) and outputs a set of output cubes out of dimension (out\_batch x out\_height x out\_width x out\_depth). These kernels are based on BATCH\_TO\_SPACE\_ND operator in TFLM<sup>[3]</sup>.

Input can be 4 dimensional (dimensions are in order – batch, height, width and depth) or 3 dimensional (for 3 dimensional input width is assumed to be 1), output is always 4 dimensional. The conversion is determined by parameters block\_sizes (num\_inp\_dims - 2) which determine conversion of a set of vectors in input (input\_batch x input\_depth) to a set of cubes (out\_batch x block\_size\_height x block\_size\_width x out\_depth) (out\_depth must be equal to input\_depth), this conversion is repeated over all (input\_height x input\_width) sets of vectors in input. Additionally, some parts of output in height and width dimensions can be cropped by using crop\_sizes.

For 4 dimensional input, number of block\_sizes are 2 (in\_order - block\_size\_height, block\_size\_width), for 3 dimensional input only block\_size\_height is used and block\_size\_width is ignored.

For 4 dimensional input, number of crop\_sizes are 4 (in order – crop\_top, crop\_bottom, crop\_left, crop\_right), crop\_top and crop\_left are used for 4 dimensional input, and only crop\_top is used for 3 dimensional input.

Naming convention used for batch\_to\_space\_nd kernels is:

xa\_nn\_batch\_to\_space\_nd\_[p]

Where [p] = [input\_precision]\_[out\_precision]

## Precision

Type	Description
8_8	Signed 8-bit input, signed 8-bit output

## Algorithm

$$out_{ob,oh,ow,d} = in_{ib,ih,iw,d}$$

$$ob = ib \% out\_batch$$

$$oh = ih * block\_size\_height - \left( \frac{ib}{out\_batch} \right) / block\_size\_width - crop\_left$$

$$ow = iw * block\_size\_width - \left( \frac{ib}{out\_batch} \right) \% block\_size\_width - crop\_top$$

% represents mod operator in C.

/ represents integer division in C.

Please refer to Figure 3-4 for visualization of batch to space conversion.

## Prototype

```
WORD32 xa_nn_batch_to_space_nd_8_8
(WORD8 *__restrict__ p_out, const WORD32 *const p_out_shape,
const WORD8 *__restrict__ p_inp, const WORD32 *const p_inp_shape,
const WORD32 *const p_block_sizes, const WORD32 *const p_crop_sizes,
WORD32 num_out_dims, WORD32 num_inp_dims);
```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
const WORD32 *const	p_out_shape	num_out_dims	Shape of output
const WORD8 *	p_inp	$\prod_{i=0}^{i=num\_inp\_dims-1} p\_inp\_shape[i]$	Input (set of cubes)
const WORD32 *const	p_inp_shape	num_inp_dims	Shape of input
const WORD32 *const	p_block_sizes	num_inp_dims - 2	Block sizes for spatial dimension.
const WORD32 *const	p_crop_sizes	2*(num_inp_dims - 2)	Crop sizes for cropping output
WORD32	num_out_dims		Number of output dimensions
WORD32	num_inp_dims		Number of input dimensions
<b>Output</b>			
WORD8 *	p_out	$\prod_{i=0}^{i=num\_out\_dims-1} p\_out\_shape[i]$	Output (set of cubes)

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions:

Arguments	Restrictions
p_out, p_inp	Aligned on (size of one element)-byte boundary
	Cannot be NULL
	Should not overlap
p_out_shape, p_inp_shape	Aligned on 4-byte boundary
	Cannot be NULL
	Should not overlap
	All elements should be greater than zero
	$p\_out\_shape[num\_out\_dims - 1] == p\_inp\_shape[num\_inp\_dims - 1]$ (depth for input and output should be equal).
p_block_sizes	Aligned on 4-byte boundary
	Cannot be NULL
	Should not overlap with other buffers
	All elements should be greater than zero
	$p\_inp\_shape[0] == p\_out\_shape[0] * p\_block\_sizes[0] * p\_block\_sizes[1]$ <sup>7</sup>
p_crop_sizes	Aligned on 4-byte boundary
	Cannot be NULL
	Should not overlap with other buffers
	All elements should be greater than or equal to zero
num_out_dims	Must be equal to 4
num_inp_dims	Must be in range {3, 4}

<sup>7</sup> This restriction is for num\_inp\_dims 4, if num\_inp\_dims is 3, it becomes  $p\_inp\_shape[0] == p\_out\_shape[0] * p\_block\_size[0]$






Figure 3-4 batch\_to\_space and space to batch conversion (for simplicity crop\_sizes and pad\_sizes are assumed to be 0)

### 3.8.5 Space to Batch Kernels

#### Description

These kernels performs space to batch conversion on a set of input cube  $in$  ( $input\_batch \times input\_height \times input\_width \times input\_depth$ ) and outputs a set of output cubes  $out$  of dimension ( $out\_batch \times out\_height \times out\_width \times out\_depth$ ). These kernels are based on SPACE\_TO\_BATCH\_ND operator in Tensorflow Lite Micro<sup>[3]</sup>.

Input can be 4 dimensional (dimensions are in order – batch, height, width and depth) or 3 dimensional (for 3 dimensional input width is assumed to be 1), output must have same number of dimensions as input. The conversion is determined by parameters  $block\_sizes$  ( $num\_inp\_dims - 2$ ) which determine conversion of a set of cubes in input ( $input\_batch \times block\_size\_height \times block\_size\_width \times input\_depth$ ) to a set of vectors ( $out\_batch \times out\_depth$ ) ( $out\_depth$  must be equal to  $input\_depth$ ), this conversion is repeated over all of input. Additionally, output can be padded in height and width dimensions according to  $pad\_sizes$ .

For 4 dimensional input, number of  $block\_sizes$  are 2 ( $in\_order - block\_size\_height, block\_size\_width$ ), for 3 dimensional input only  $block\_size\_height$  is used and  $block\_size\_width$  is ignored.

For 4 dimensional input, number of  $pad\_sizes$  are 4 ( $in\_order - pad\_top, pad\_bottom, pad\_left, pad\_right$ ),  $pad\_top$  and  $pad\_left$  are used for 4 dimensional input, and only  $pad\_top$  is used for 3 dimensional input.

The value to be filled in padding regions can be specified by  $pad\_value$ .

Naming convention used for  $space\_to\_batch\_nd$  kernels is:

$xa\_nn\_batch\_to\_space\_nd\_ [p]$

Where  $[p] = [input\_precision]_{[out\_precision]}$

#### Precision

Type	Description
8_8	Signed 8-bit input, signed 8-bit output

#### Algorithm

$$out_{ob,oh,ow,d} = in_{ib,ih,iw,d}$$

$$ib = ob \% out\_batch$$

$$ih = oh * block\_size\_height - \left( \frac{ob}{input\_batch} \right) / block\_size\_width - crop\_left$$

$$iw = ow * block\_size\_width - \left( \frac{ob}{input\_batch} \right) \% block\_size\_width - crop\_top$$

% represents mod operator in C.

/ represents integer division in C.

Please refer to Figure 3-4 for visualization of space to batch conversion.

## Prototype

```
WORD32 xa_nn_space_to_batch_nd_8_8
(WORD8 *__restrict__ p_out, const WORD32 *const p_out_shape,
 const WORD8 *__restrict__ p_inp, const WORD32 *const p_inp_shape,
 const WORD32 *const p_block_sizes, const WORD32 *const p_pad_sizes,
 WORD32 num_out_dims, WORD32 num_inp_dims
 WORD32 pad_value);
```

## Arguments

Type	Name	Size	Description
<b>Input</b>			
const WORD32 *const	p_out_shape	num_out_dims	Shape of output
const WORD8 *	p_inp	$\prod_{i=0}^{l=num\_inp\_dims-1} p\_inp\_shape[i]$	Input (set of cubes)
const WORD32 *const	p_inp_shape	num_inp_dims	Shape of input
const WORD32 *const	p_block_sizes	num_inp_dims - 2	Block sizes for spatial dimension.
const WORD32 *const	p_pad_sizes	2*(num_inp_dims - 2)	Crop sizes for cropping output
WORD32	num_out_dims		Number of output dimensions
WORD32	num_inp_dims		Number of input dimensions
WORD32	pad_value		Value for padding
<b>Output</b>			
WORD8 *	p_out	$\prod_{i=0}^{l=num\_out\_dims-1} p\_out\_shape[i]$	Output (set of cubes)

## Returns

- 0: no error
- -1: error, invalid parameters

## Restrictions:

Arguments	Restrictions
p_out, p_inp	Aligned on (size of one element)-byte boundary

	Cannot be NULL
	Should not overlap
p_out_shape, p_inp_shape	Aligned on 4-byte boundary
	Cannot be NULL
	Should not overlap
	All elements should be greater than zero
	$p\_out\_shape[num\_out\_dims - 1] == p\_inp\_shape[num\_inp\_dims - 1]$ (depth for input and output should be equal).
p_block_sizes	Aligned on 4-byte boundary
	Cannot be NULL
	Should not overlap with other buffers
	All elements should be greater than zero
	$p\_out\_shape[0] == p\_inp\_shape[0] * p\_block\_sizes[0] * p\_block\_sizes[1]^8$
p_pad_sizes	Aligned on 4-byte boundary
	Cannot be NULL
	Should not overlap with other buffers
	All elements should be greater than or equal to zero
num_out_dims	Must be in range {3, 4}
num_inp_dims	Must be in range {3, 4}
pad_value	Must be in range [-128, 127]

<sup>8</sup> This restriction is for num\_inp\_dims 4, if num\_inp\_dims is 3, it becomes  $p\_out\_shape[0] == p\_inp\_shape[0] * p\_block\_size[0]$

## 4. HiFi 1 NN Library – Layers

This section explains the APIs of each layer implementation in the NN library. All the layers conform to the “generic NN Layer API” and flow explained in Section 2.

The NN library is a single archive containing all layers and low-level kernels implementations. Each layer has its own header file that defines the APIs specific to the layer. The following sections explain each layer in detail.

---

**Note** This version of the library supports GRU, LSTM, and CNN layers.

---

### 4.1 GRU Layer

The GRU APIs are defined in `xa_nnlb_gru_api.h`. Refer to the overall signal flow diagram of GRU in [\[4\]](#).

#### 4.1.1 GRU Layer Specification

GRU layer implements the following input-output equations [\[4\]](#):

$$\begin{aligned} z_t &= \text{sigmoid}(W_z * x_t + U_z * \text{prev-h} + b_z) \\ r_t &= \text{sigmoid}(W_r * x_t + U_r * \text{prev-h} + b_r) \\ g &= \tanh(W_h * x_t + U_h * (r_t * \text{prev-h}) + b_h) \\ y_t &= h_t = z_t * g + (1 - z_t) * \text{prev-h} \\ \text{prev-h} &= h_t \end{aligned}$$

$x_t$  : input vector  
 $y_t, h_t$  : output vector  
 $W, U$  : weight matrices  
 $\text{prev-h}$ : previous output vector

$z_t$  : update gate vector  
 $r_t$  : reset gate vector  
 $b$  : bias vectors

## 4.1.2 Error Codes Specific to GRU

Other than common error codes explained in Section 2.3, the GRU layer may also report the following error codes, which may be generated during the initialization stage.

- `XA_NNLIB_GRU_CONFIG_FATAL_INVALID_IN_FEATS9`  
Number of input features is not supported
- `XA_NNLIB_GRU_CONFIG_FATAL_INVALID_OUT_FEATS`  
Number of output features is not supported
- `XA_NNLIB_GRU_CONFIG_FATAL_INVALID_PRECISION`  
I/O precision is not supported
- `XA_NNLIB_GRU_CONFIG_FATAL_INVALID_COEFF_QFORMAT`  
Number of fractional bits for coefficients is not supported.
- `XA_NNLIB_GRU_CONFIG_FATAL_INVALID_IO_QFORMAT`  
Number of fractional bits for input-output is not supported.
- `XA_NNLIB_GRU_CONFIG_FATAL_INVALID_MEMBANK_PADDING`  
Membank padding should be 0 or 1.
- `XA_NNLIB_GRU_CONFIG_FATAL_INVALID_PARAM_ID`  
Parameter identifier (`param_id`) is not valid

The following error codes may be generated during the execution stage.

- `XA_NNLIB_GRU_EXECUTE_FATAL_INSUFFICIENT_DATA`  
Input data passed in is insufficient
- `XA_NNLIB_GRU_EXECUTE_FATAL_INSUFFICIENT_OUTPUT_BUFFER_SPACE`  
Output Buffer Size is not sufficient

---

<sup>9</sup> FEATS := features

## 4.1.3 API Functions Specific to GRU

### 4.1.3.1 Query Functions

Table 4-1 GRU Get Persistent Size Function

<b>Function</b>	<code>xa_nnlb_gru_get_persistent_fast</code>
<b>Syntax</b>	<pre>Int32 xa_nnlb_gru_get_persistent_fast (     xa_nnlb_gru_init_config_t *config)</pre>
<b>Description</b>	Returns persistent memory size in bytes required by GRU layer.
<b>Parameters</b>	Input: <code>config</code> Initial configuration parameters (see Table 4-7).
<b>Errors</b>	<p>If return value is less than 0, then it is an error. Following are the possible error codes:</p> <ul style="list-style-type: none"> <li>■ <code>XA_NNLB_FATAL_MEM_ALLOC</code></li> <li>■ <code>XA_NNLB_GRU_CONFIG_FATAL_INVALID_IN_FEATS</code> Number of input features is not supported</li> <li>■ <code>XA_NNLB_GRU_CONFIG_FATAL_INVALID_OUT_FEATS</code> Number of output features is not supported</li> <li>■ <code>XA_NNLB_GRU_CONFIG_FATAL_INVALID_PRECISION</code> I/O precision is not supported</li> <li>■ <code>XA_NNLB_GRU_CONFIG_FATAL_INVALID_COEFF_QFORMAT</code> Number of fractional bits for coefficients is not supported.</li> <li>■ <code>XA_NNLB_GRU_CONFIG_FATAL_INVALID_IO_QFORMAT</code> Number of fractional bits for input-output is not supported.</li> </ul>

Table 4-2 GRU Get Scratch Size Function

<b>Function</b>	<code>xa_nnlb_gru_get_scratch_fast</code>
<b>Syntax</b>	<pre>Int32 xa_nnlb_gru_get_scratch_fast(     xa_nnlb_gru_init_config_t *config)</pre>
<b>Description</b>	Returns scratch memory size in bytes required by GRU layer.
<b>Parameters</b>	<p>Input: <code>config</code></p> <p>Initial configuration parameters (see Table 4-7).</p>
<b>Errors</b>	<p>If return value is less than 0, then it is an error. Following are the possible error codes:</p> <ul style="list-style-type: none"> <li>■ <code>XA_NNLB_FATAL_MEM_ALLOC</code></li> <li>■ <code>XA_NNLB_GRU_CONFIG_FATAL_INVALID_IN_FEATS</code> Number of input features is not supported</li> <li>■ <code>XA_NNLB_GRU_CONFIG_FATAL_INVALID_OUT_FEATS</code> Number of output features is not supported</li> <li>■ <code>XA_NNLB_GRU_CONFIG_FATAL_INVALID_PRECISION</code> I/O precision is not supported</li> <li>■ <code>XA_NNLB_GRU_CONFIG_FATAL_INVALID_COEFF_QFORMAT</code> Number of fractional bits for coefficients is not supported</li> <li>■ <code>XA_NNLB_GRU_CONFIG_FATAL_INVALID_IO_QFORMAT</code> Number of fractional bits for input-output is not supported</li> </ul>



### 4.1.3.2 Initialization Stage

Table 4-3 GRU InitFunction

<b>Function</b>	<code>xa_nnlb_gru_init</code>
<b>Syntax</b>	<pre>Int32 xa_nnlb_gru_init (     xa_nnlb_handle_t handle,     xa_nnlb_gru_init_config_t *config)</pre>
<b>Description</b>	Reset the GRU Layer API handle into its initial state. Set up the GRU Layer to the specified initial configuration parameters. This function sets <code>prev_h</code> vector to 0; the user can put the desired values in <code>prev_h</code> by using set config <code>XA_NNLB_GRU_RESTORE_CONTEXT</code> (refer to Table 4-11 for more information).
<b>Parameters</b>	<p>Input: <code>handle</code>            Pointer to the component persistent memory. This is the opaque handle.            Required size: see <code>xa_nnlb_gru_get_persistent_fast</code>.            Required alignment: 8 bytes.</p> <p>Input: <code>config</code>            Initial configuration parameters (see Table 4-7). Note that the initial configuration parameters <i>must</i> be identical to those passed to query functions.</p>
<b>Errors</b>	<p>If the return value is not <code>XA_NNLB_NO_ERROR</code>, it implies that the function has encountered one of the following errors:</p> <ul style="list-style-type: none"> <li>■ <code>XA_NNLB_FATAL_MEM_ALLOC</code> One of the pointers is invalid.</li> <li>■ <code>XA_NNLB_FATAL_MEM_ALIGN</code> One of the pointers is not properly aligned.</li> <li>■ <code>XA_NNLB_GRU_CONFIG_FATAL_INVALID_IN_FEATS</code> Number of input features is not supported</li> <li>■ <code>XA_NNLB_GRU_CONFIG_FATAL_INVALID_OUT_FEATS</code> Number of output features is not supported</li> <li>■ <code>XA_NNLB_GRU_CONFIG_FATAL_INVALID_PRECISION</code> I/O precision is not supported.</li> <li>■ <code>XA_NNLB_GRU_CONFIG_FATAL_INVALID_COEFF_QFORMAT</code> Number of fractional bits for coefficients is not supported.</li> <li>■ <code>XA_NNLB_GRU_CONFIG_FATAL_INVALID_IO_QFORMAT</code> Number of fractional bits for input-output is not supported.</li> </ul>

### 4.1.3.3 Execution Stage

Table 4-4 GRU Execution Function

<b>Function</b>	<code>xa_nnl-lib-gru-process</code>
<b>Syntax</b>	<pre> Int32 xa_nnl-lib-gru-process (     xa_nnl-lib-handle_t handle,     void *scratch,     void *input,     void *output,     xa_nnl-lib-shape_t *p_in_shape,     xa_nnl-lib-shape_t *p_out_shape) </pre>
<b>Description</b>	Processes one input shape to generate one output shape.
<b>Parameters</b>	<p><b>Input: handle</b> The opaque component handle. Required alignment: 8 bytes.</p> <p><b>Input: scratch</b> A pointer to the scratch buffer. Required alignment: 8 bytes.</p> <p><b>Input: input</b> A pointer to the input buffer. Input buffer contains input data. Required alignment: 8 bytes.</p> <p><b>Output: output</b> A pointer to the output buffer. Output is written to output buffer. Required alignment: 8 bytes.</p> <p><b>Input/Output: p_in_shape</b> Pointer to the shape containing input buffer dimensions. Contains the length of input data passed to GRU layer. Required alignment: 4 bytes.</p> <p><b>Input/Output: p_out_shape</b> Pointer to the shape for output buffer dimensions. On return, *p_out_shape is filled with the length of output generated by HiFi GRU Layer. Required alignment: 4 bytes.</p>
<b>Errors</b>	<p>If the return value is not <code>XA_NNL-LIB-NO-ERROR</code>, it implies that the function has encountered one of the following errors:</p> <ul style="list-style-type: none"> <li>■ <code>XA_NNL-LIB-FATAL-MEM-ALLOC</code> One of the pointers is NULL.</li> <li>■ <code>XA_NNL-LIB-FATAL-MEM-ALIGN</code> One of the pointers is not properly aligned.</li> </ul>

	<ul style="list-style-type: none"> <li>■ <b>XA_NNLIB_FATAL_INVALID_SHAPE</b> Either input or output shape is invalid.</li> <li>■ <b>XA_NNLIB_GRU_EXECUTE_FATAL_INSUFFICIENT_DATA</b> Input data passed in insufficient.</li> <li>■ <b>XA_NNLIB_GRU_EXECUTE_FATAL_INSUFFICIENT_OUTPUT_BUFFER_SPACE</b> Output buffer size is not sufficient.</li> </ul>
--	--

Table 4-5 GRU Set Parameter Function Details

<b>Function</b>	<code>xa_nnl-lib-gru-set-config</code>
<b>Syntax</b>	<pre>Int32 xa_nnl-lib-gru-set-config (     xa_nnl-lib-handle_t handle,     xa_nnl-lib-gru-param-id_t param_id,     void *params)</pre>
<b>Description</b>	Sets the parameter specified by <code>param_id</code> to the value passed in the buffer pointed to by <code>params</code> .
<b>Parameters</b>	<p><b>Input:</b> <code>handle</code> The opaque component handle. Required alignment: 8 bytes.</p> <p><b>Input:</b> <code>param_id</code> Identifies the parameter to be written. Refer to Table 4-11 for the list of supported parameters.</p> <p><b>Input:</b> <code>params</code> A pointer to a buffer that contains the parameter value. Required alignment: 4 bytes.</p>
<b>Errors</b>	<p>If the return value is not <code>XA_NNLIB_NO_ERROR</code>, it implies that function has encountered one of the following errors:</p> <ul style="list-style-type: none"> <li>■ <b>XA_NNLIB_FATAL_MEM_ALLOC</b> One of the pointers (<code>handle</code> or <code>params</code>) is NULL.</li> <li>■ <b>XA_NNLIB_FATAL_MEM_ALIGN</b> One of the pointers (<code>handle</code> or <code>params</code>) is not aligned correctly.</li> <li>■ <b>XA_NNLIB_GRU_CONFIG_FATAL_INVALID_PARAM_ID</b> Parameter identifier (<code>param_id</code>) is not valid.</li> </ul>

Table 4-6 GRU Get Parameter Function Details

<b>Function</b>	<code>xa_nnlib_gru_get_config</code>
<b>Syntax</b>	<pre>Int32 xa_nnlib_gru_get_config (     xa_nnlib_handle_t handle,     xa_nnlib_gru_param_id_t param_id,     void *params)</pre>
<b>Description</b>	Gets the value of the parameter specified by <code>param_id</code> in the buffer pointed to by <code>params</code> .
<b>Parameters</b>	<p><b>Input:</b> <code>handle</code> The opaque component handle. Required alignment: 8 bytes.</p> <p><b>Input:</b> <code>param_id</code> Identifies the parameter to be read. Refer to Table 4-11 for the list of supported parameters.</p> <p><b>Output:</b> <code>params</code> A pointer to a buffer that is filled with the parameter value when the function returns. Required alignment: 4 bytes.</p>
<b>Errors</b>	<p>If the return value is not <code>XA_NNLIB_NO_ERROR</code>, it implies that function has encountered one of the following errors:</p> <ul style="list-style-type: none"> <li>■ <code>XA_NNLIB_FATAL_MEM_ALLOC</code> One of the pointers (<code>handle</code> or <code>params</code>) is <code>NULL</code>.</li> <li>■ <code>XA_NNLIB_FATAL_MEM_ALIGN</code> One of the pointers (<code>handle</code> or <code>params</code>) is not aligned correctly.</li> <li>■ <code>XA_NNLIB_GRU_CONFIG_FATAL_INVALID_PARAM_ID</code> Parameter identifier (<code>param_id</code>) is not valid.</li> </ul>

## 4.1.4 Structures Specific to GRU

Table 4-7 GRU Config Structure `xa_nnlib_gru_init_config_t`

Element Type	Element Name	Range	Default	Description
Int32	in_feats	4-2048	256	Number of input features (must be multiple of 4)
Int32	out_feats	4-2048	256	Number of output features (must be multiple of 4)
Int32	pad	0, 1	1	Padding 16 bytes for HiFi 1
Int32	mat_prec	8, 16	16	Matrix input precision
Int32	vec_prec	16	16	Vector input precision
<code>xa_nnlib_gru_precision_t</code>	precision	XA_NNLIB_GRU_16bx16b, XA_NNLIB_GRU_8bx16b	XA_NNLIB_GRU_16bx16b	Coef and I/O precision. Note: Current library supports only 16bx16b and 8bx16b precision for GRU
Int16	coeff_Qformat	0-15	15	Number of fractional bits for weights and biases
Int16	io_Qformat	0-15	12	Number of fractional bits for input and output

Table 4-8 `xa_nnlib_gru_weights_t` Parameter Type

Element Type	Element Name	Range	Default	Description
<code>coeff_t *</code>	w_z	NA	NA	Pointer to coefficient matrix w_z.
<code>xa_nnlib_shape_t</code>	shape_w_z	NA	NA	Shape information about w_z.
<code>coeff_t *</code>	u_z	NA	NA	Pointer to coefficient matrix u_z.
<code>xa_nnlib_shape_t</code>	shape_u_z	NA	NA	Shape information about u_z.
<code>coeff_t *</code>	w_r	NA	NA	Pointer to coefficient matrix w_r.
<code>xa_nnlib_shape_t</code>	shape_w_r	NA	NA	Shape information about w_r.
<code>coeff_t *</code>	u_r	NA	NA	Pointer to coefficient matrix u_r.
<code>xa_nnlib_shape_t</code>	shape_u_r	NA	NA	Shape information about u_r.
<code>coeff_t *</code>	w_h	NA	NA	Pointer to coefficient matrix w_h.
<code>xa_nnlib_shape_t</code>	shape_w_h	NA	NA	Shape information about w_h.
<code>coeff_t *</code>	u_h	NA	NA	Pointer to coefficient matrix u_h.
<code>xa_nnlib_shape_t</code>	shape_u_h	NA	NA	Shape information about u_h.

Table 4-9 xa\_nnlib\_gru\_biases\_t Parameter Type

Element Type	Element Name	Range	Default	Description
coeff_t *	b_z	NA	NA	Pointer to coefficient matrix b_z.
xa_nnlib_shape_t	shape_b_z	NA	NA	Shape information about b_z.
coeff_t *	b_r	NA	NA	Pointer to coefficient matrix b_r.
xa_nnlib_shape_t	shape_b_r	NA	NA	Shape information about b_r.
coeff_t *	b_h	NA	NA	Pointer to coefficient matrix b_h.
xa_nnlib_shape_t	shape_b_h	NA	NA	Shape information about b_h.

**Note** GRU requires all weight matrices' and bias vectors' pointers to be 8 bytes aligned.

## 4.1.5 Enums Specific to GRU

Table 4-10 Enum xa\_nnlib\_gru\_precision\_t

Element	Description
XA_NNLIB_GRU_16bx16b	Coef: 16 bits, I/O: 16 bits Fixed Point
XA_NNLIB_GRU_8bx16b	Coef: 8 bits, I/O: 16 bits Fixed Point
XA_NNLIB_GRU_8bx8b	Not supported
XA_NNLIB_flt16xflt16	Not supported

**Note** Currently, GRU only supports XA\_NNLIB\_GRU\_16bx16b, XA\_NNLIB\_GRU\_8bx16b precision setting.

Table 4-11 describes parameter IDs for parameters supported by GRU. It contains the following columns:

- Parameter ID: Parameter identifier (`param_id`).
- Value type: A pointer (`params`) to a variable of this type is to be passed.
- RW: Indicates whether the parameter can be read (`get`) and/or written (`set`).
- Range: Indicates valid values of the parameter.
- Default: Default value of the parameter
- Description: Brief description of the parameter.

Table 4-11 GRU Specific Parameters

Parameter ID	Value Type	RW	Range	Default	Description
XA_NNLIB_GRU_RESTORE_CONTEXT	vect_t []	RW	NA	NA	Set previous output. This can be used to set prev_h to specific context (size should be equal to number of output features). Upon set config, the buffer passed is copied to persistent memory; upon get config, it returns the prev_h state in the given buffer.
XA_NNLIB_GRU_WEIGHT	xa_nnlib_gru_weights_t	RW	NA	NA	Weight matrices, pointers to weight matrices along with shape information must be passed via xa_nnlib_gru_weights_t structure for set config. Upon get config, it returns pointers to weight matrices along with their shape information in same structure.
XA_NNLIB_GRU_BIAS	xa_nnlib_gru_biases_t	RW	NA	NA	Bias vectors, pointers to bias vectors along with shape information must be passed via xa_nnlib_gru_biases_t structure for set config. Upon get config, it returns pointers to bias vectors along with their shape information in same structure.
XA_NNLIB_GRU_INPUT_SHAPE	xa_nnlib_shape_t	R	NA	NA	Input shape information, get information of the input shape expected by the layer.
XA_NNLIB_GRU_OUTPUT_SHAPE	xa_nnlib_shape_t	R	NA	NA	Output shape information, get information of the output shape expected by layer.

## 4.2 LSTM Layer

The LSTM APIs are defined in `xa_nnl-lib-lstm-api.h`.

### 4.2.1 LSTM Layer Specification

The LSTM layer implements the following forward path input-output equations:

$$\begin{aligned} f_f &= \text{sigmoid}(w_{xf} * \text{frame}_f + \text{prev-h} * w_{hf} + b_f) \\ i_f &= \text{sigmoid}(w_{xi} * \text{frame}_f + \text{prev-h} * w_{hi} + b_i) \\ c\text{-hat}_f &= \tanh(w_{xc} * \text{frame}_f + \text{prev-h} * w_{hc} + b_c) \\ c_f &= f_f \cdot \text{prev-c} + i_f * c\text{-hat}_f \\ o_f &= \text{sigmoid}(w_{xo} * \text{frame}_f + \text{prev-h} * w_{ho} + b_o) \\ h_f &= o_f * \tanh(c_f) \end{aligned}$$

$i_f$ : input gate

$h_t$ : output vector

$c\text{-hat}_f$ : intermediate cell state vector

$f_f$ : forget gate

$\text{frame}_f$ : Input vector

$w_x$ : weight matrices of input connections

$\text{prev-h}$ : previous output vector

$\text{prev-c}$ : previous cell output

$b$ : bias vectors

$o_f$ : output gate

$c_f$ : cell state vector

$w_h$ : weight matrices of recurrent connections

### 4.2.2 Error Codes Specific to LSTM

Other than common error codes explained in Section 2.3, the LSTM layer may also report the following error codes, which may be generated during the initialization stage:

- `XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_IN_FEATS`<sup>10</sup>  
Number of input features is not supported
- `XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_OUT_FEATS`  
Number of output features is not supported
- `XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_PRECISION`  
I/O precision is not supported
- `XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_COEFF_QFORMAT`  
Number of fractional bits for coefficients is not supported.
- `XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_CELL_QFORMAT`  
Number of fractional bits for cells is not supported

<sup>10</sup> FEATS: = features



- `XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_IO_QFORMAT`  
Number of fractional bits for input-output is not supported.
- `XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_MEMBANK_PADDING`  
Membank padding should be 0 or 1.
- `XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_PARAM_ID`  
Parameter identifier (param\_id) is not valid

The following error codes may be generated during the execution stage.

- `XA_NNLIB_LSTM_EXECUTE_FATAL_INSUFFICIENT_DATA`  
Input data passed in insufficient
- `XA_NNLIB_LSTM_EXECUTE_FATAL_INSUFFICIENT_OUTPUT_BUFFER_SPACE`  
Output Buffer Size is not sufficient

## 4.2.3 API Functions Specific to LSTM

### 4.2.3.1 Query Functions

Table 4-12 LSTM Get Persistent Size Function

<b>Function</b>	<code>xa_nnlib_lstm_get_persistent_fast</code>
<b>Syntax</b>	<pre>Int32 xa_nnlib_lstm_get_persistent_fast (     xa_nnlib_lstm_init_config_t *config)</pre>
<b>Description</b>	Returns persistent memory size in bytes required by LSTM layer.
<b>Parameters</b>	Input: <code>config</code> Initial configuration parameters (see Table 4-18).
<b>Errors</b>	<p>If return value is less than 0 then it is an error. Following are the possible error codes:</p> <ul style="list-style-type: none"> <li>■ <code>XA_NNLIB_FATAL_MEM_ALLOC</code></li> <li>■ <code>XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_IN_FEATS</code> Number of input features is not supported</li> <li>■ <code>XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_OUT_FEATS</code> Number of output features is not supported</li> <li>■ <code>XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_PRECISION</code> I/O precision is not supported</li> <li>■ <code>XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_COEFF_QFORMAT</code> Number of fractional bits for coefficients is not supported.</li> <li>■ <code>XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_CELL_QFORMAT</code> Number of fractional bits for cells is not supported</li> <li>■ <code>XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_IO_QFORMAT</code> Number of fractional bits for input-output is not supported.</li> <li>■ <code>XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_MEMBANK_PADDING</code> Membank padding should be 0 or 1.</li> </ul>

Table 4-13 LSTM Get Scratch Size Function

<b>Function</b>	<code>xa_nnlib_lstm_get_scratch_fast</code>
<b>Syntax</b>	<code>Int32 xa_nnlib_lstm_get_scratch_fast (</code> <code>xa_nnlib_lstm_init_config_t *config)</code>
<b>Description</b>	Returns scratch memory size in bytes required by LSTM layer.
<b>Parameters</b>	Input: <code>config</code> Initial configuration parameters (see Table 4-18).
<b>Errors</b>	<p>If return value is less than 0 then it is an error, the possible error codes are:</p> <ul style="list-style-type: none"> <li>■ <code>XA_NNLIB_FATAL_MEM_ALLOC</code></li> <li>■ <code>XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_IN_FEATS</code> Number of input features is not supported</li> <li>■ <code>XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_OUT_FEATS</code> Number of output features is not supported</li> <li>■ <code>XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_PRECISION</code> I/O precision is not supported</li> <li>■ <code>XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_COEFF_QFORMAT</code> Number of fractional bits for coefficients is not supported.</li> <li>■ <code>XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_CELL_QFORMAT</code> Number of fractional bits for cells is not supported</li> <li>■ <code>XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_IO_QFORMAT</code> Number of fractional bits for input-output is not supported.</li> <li>■ <code>XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_MEMBANK_PADDING</code> Membank padding should be 0 or 1.</li> </ul>

## 4.2.3.2 Initialization Stage

Table 4-14 LSTM InitFunction

<b>Function</b>	<code>xa_nnlb_lstm_init</code>
<b>Syntax</b>	<pre>Int32 xa_nnlb_lstm_init (     xa_nnlb_handle_t handle,     xa_nnlb_lstm_init_config_t *config)</pre>
<b>Description</b>	<p>Reset the LSTM layer API handle into its initial state. Set up the LSTM layer to the specified initial configuration parameters. This function sets <code>prev_h</code> vector and <code>prev_c</code> vector to 0; the user can put the desired values in <code>prev_h</code> and <code>prev_c</code> by using <code>set config</code> <code>XA_NNLB_LSTM_RESTORE_CONTEXT_OUTPUT</code> and <code>XA_NNLB_LSTM_RESTORE_CONTEXT_CELL</code> respectively (refer to Table 4-22 for more information).</p>
<b>Parameters</b>	<p><b>Input: <code>handle</code></b>            Pointer to the component persistent memory. This is the opaque handle.            Required size: see <code>xa_nnlb_lstm_get_persistent_fast</code>.            Required alignment: 8 bytes.</p> <p><b>Input: <code>config</code></b>            Initial configuration parameters (see Table 4-18). Note that the initial configuration parameters MUST be identical to those passed to query functions.</p>
<b>Errors</b>	<p>If the return value is not <code>XA_NNLB_NO_ERROR</code>, it implies that the function has encountered one of the following errors:</p> <ul style="list-style-type: none"> <li>■ <code>XA_NNLB_FATAL_MEM_ALLOC</code> One of the pointers is invalid.</li> <li>■ <code>XA_NNLB_FATAL_MEM_ALIGN</code> One of the pointers is not properly aligned.</li> <li>■ <code>XA_NNLB_LSTM_CONFIG_FATAL_INVALID_IN_FEATS</code> Number of input features is not supported</li> <li>■ <code>XA_NNLB_LSTM_CONFIG_FATAL_INVALID_OUT_FEATS</code> Number of output features is not supported</li> <li>■ <code>XA_NNLB_LSTM_CONFIG_FATAL_INVALID_PRECISION</code> I/O precision is not supported</li> <li>■ <code>XA_NNLB_LSTM_CONFIG_FATAL_INVALID_COEFF_QFORMAT</code> Number of fractional bits for coefficients is not supported.</li> </ul>

- `XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_CELL_QFORMAT`  
Number of fractional bits for cells is not supported
- `XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_IO_QFORMAT`  
Number of fractional bits for input-output is not supported
- `XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_MEMBANK_PADDING`  
Membank padding should be 0 or 1.

### 4.2.3.3 Execution Stage

Table 4-15 LSTM Execution Function

<b>Function</b>	<code>xa_nnl-lib_lstm_process</code>
<b>Syntax</b>	<pre> Int32 xa_nnl-lib_lstm_process (     xa_nnl-lib_handle_t handle,     void *scratch,     void *input,     void *output,     xa_nnl-lib_shape_t *p_in_shape,     xa_nnl-lib_shape_t *p_out_shape) </pre>
<b>Description</b>	Processes one input shape to generate one output shape.
<b>Parameters</b>	<p><b>Input: <code>handle</code></b> The opaque component handle. Required alignment: 8 bytes.</p> <p><b>Input: <code>scratch</code></b> A pointer to the scratch buffer. Required alignment: 8 bytes.</p> <p><b>Input: <code>input</code></b> A pointer to the input buffer. Input buffer contains input data. Required alignment: 8 bytes.</p> <p><b>Output: <code>output</code></b> A pointer to the output buffer. Output is written to the output buffer. Required alignment: 8 bytes.</p> <p><b>Input/Output: <code>p_in_shape</code></b> Pointer to the shape containing input buffer dimensions. Contains the length of input data passed to LSTM layer. Required alignment: 4 bytes.</p> <p><b>Input/Output: <code>p_out_shape</code></b></p>

	<p>Pointer to the shape for output buffer dimensions. On return, *p_out_shape is filled with the length of output generated by HiFi LSTM layer.</p> <p>Required alignment: 4 bytes.</p>
<b>Errors</b>	<p>If the return value is not XA_NNLIB_NO_ERROR, it implies that the function has encountered one of the following errors:</p> <ul style="list-style-type: none"><li>■ XA_NNLIB_FATAL_MEM_ALLOC One of the pointers is NULL.</li><li>■ XA_NNLIB_FATAL_MEM_ALIGN One of the pointers is not having proper alignment.</li><li>■ XA_NNLIB_FATAL_INVALID_SHAPE Either input or output shape is invalid.</li><li>■ XA_NNLIB_LSTM_EXECUTE_FATAL_INSUFFICIENT_DATA Input data passed in insufficient</li><li>■ XA_NNLIB_LSTM_EXECUTE_FATAL_INSUFFICIENT_OUTPUT_BUFFER_SPACE Output Buffer Size is not sufficient</li></ul>

Table 4-16 LSTM Set Parameter Function Details

<b>Function</b>	<code>xa_nnlib_lstm_set_config</code>
<b>Syntax</b>	<pre> Int32 xa_nnlib_lstm_set_config (     xa_nnlib_handle_t handle,     xa_nnlib_lstm_param_id_t param_id,     void *params) </pre>
<b>Description</b>	Sets the parameter specified by <code>param_id</code> to the value passed in the buffer pointed to by <code>params</code> .
<b>Parameters</b>	<p><b>Input:</b> <code>handle</code> The opaque component handle. Required alignment: 8 bytes.</p> <p><b>Input:</b> <code>param_id</code> Identifies the parameter to be written. Refer to Table 4-11 for the list of supported parameters.</p> <p><b>Input:</b> <code>params</code> A pointer to a buffer that contains the parameter value. Required alignment: 4 bytes.</p>
<b>Errors</b>	<p>If the return value is not <code>XA_NNLIB_NO_ERROR</code>, it implies that the function has encountered one of the following errors:</p> <ul style="list-style-type: none"> <li>■ <code>XA_NNLIB_FATAL_MEM_ALLOC</code> One of the pointers (<code>handle</code> or <code>params</code>) is NULL.</li> <li>■ <code>XA_NNLIB_FATAL_MEM_ALIGN</code> One of the pointers (<code>handle</code> or <code>params</code>) is not aligned correctly.</li> <li>■ <code>XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_PARAM_ID</code> Parameter identifier (<code>param_id</code>) is not valid.</li> </ul>

Table 4-17 LSTM Get Parameter Function Details

<b>Function</b>	<code>xa_nnlib_lstm_get_config</code>
<b>Syntax</b>	<pre>Int32 xa_nnlib_lstm_get_config (     xa_nnlib_handle_t handle,     xa_nnlib_lstm_param_id_t param_id,     void *params)</pre>
<b>Description</b>	Gets the value of the parameter specified by <code>param_id</code> in the buffer pointed to by <code>params</code> .
<b>Parameters</b>	<p><b>Input:</b> <code>handle</code> The opaque component handle. Required alignment: 8 bytes.</p> <p><b>Input:</b> <code>param_id</code> Identifies the parameter to be read. Refer to Table 4-11 for the list of supported parameters.</p> <p><b>Output:</b> <code>params</code> A pointer to a buffer that is filled with the parameter value when the function returns. Required alignment: 4 bytes.</p>
<b>Errors</b>	<p>If the return value is not <code>XA_NNLIB_NO_ERROR</code>, it implies that the function has encountered one of the following errors:</p> <ul style="list-style-type: none"> <li>■ <code>XA_NNLIB_FATAL_MEM_ALLOC</code> One of the pointers (<code>handle</code> or <code>params</code>) is <code>NULL</code>.</li> <li>■ <code>XA_NNLIB_FATAL_MEM_ALIGN</code> One of the pointers (<code>handle</code> or <code>params</code>) is not aligned correctly.</li> <li>■ <code>XA_NNLIB_LSTM_CONFIG_FATAL_INVALID_PARAM_ID</code> Parameter identifier (<code>param_id</code>) is not valid.</li> </ul>



## 4.2.4 Structures Specific to LSTM

Table 4-18 LSTM Config Structure `xa_nnlib_lstm_init_config_t`

Element Type	Element Name	Range	Default	Description
Int32	in_feats	4-2048	256	Number of input features (must be multiple of 4)
Int32	out_feats	4-2048	256	Number of output features (must be multiple of 4)
Int32	pad	0, 1	1	Padding 16 bytes for HiFi 1 DSP
Int32	mat_prec	8, 16	16	Matrix input precision
Int32	vec_prec	16	16	Vector input precision
<code>xa_nnlib_lstm_precision_t</code>	precision	XA_NNLIB_LSTM_16bx16b, XA_NNLIB_LSTM_8bx16b	XA_NNLIB_LSTM_16bx16b	Coef and I/O precision. Note: The current library supports only 16bx16b and 8bx16b precision for LSTM.
Int16	coeff_Qformat	0-15	15	Number of fractional bits for weights and biases
Int16	cell_Qformat	0-26		Number of fractional bits for cells.
Int16	io_Qformat	0-15	12	Number of fractional bits for input and output

Table 4-19 `xa_nnlib_lstm_weights_t` Parameter Type

Element Type	Element Name	Range	Default	Description
<code>coeff_t *</code>	w_xf	NA	NA	Pointer to coefficient matrix w_xf.
<code>xa_nnlib_shape_t</code>	shape_w_xf	NA	NA	Shape information about w_xf.
<code>coeff_t *</code>	w_xi	NA	NA	Pointer to coefficient matrix w_xi.
<code>xa_nnlib_shape_t</code>	shape_w_xi	NA	NA	Shape information about w_xi.
<code>coeff_t *</code>	w_xc	NA	NA	Pointer to coefficient matrix w_xc.
<code>xa_nnlib_shape_t</code>	shape_w_xc	NA	NA	Shape information about w_xc.
<code>coeff_t *</code>	w_xo	NA	NA	Pointer to coefficient matrix w_xo.
<code>xa_nnlib_shape_t</code>	shape_w_xo	NA	NA	Shape information about w_xo.
<code>coeff_t *</code>	w_hf	NA	NA	Pointer to coefficient matrix w_hf.
<code>xa_nnlib_shape_t</code>	shape_w_hf	NA	NA	Shape information about w_hf.
<code>coeff_t *</code>	w_hi	NA	NA	Pointer to coefficient matrix w_hi.

Element Type	Element Name	Range	Default	Description
xa_nnlib_shape_t	shape_w_hi	NA	NA	Shape information about w_hi.
coeff_t *	w_hc	NA	NA	Pointer to coefficient matrix w_hc.
xa_nnlib_shape_t	shape_w_hc	NA	NA	Shape information about w_hc.
coeff_t *	w_ho	NA	NA	Pointer to coefficient matrix w_ho.
xa_nnlib_shape_t	shape_w_ho	NA	NA	Shape information about w_ho.

Table 4-20 xa\_nnlib\_lstm\_biases\_t Parameter Type

Element Type	Element Name	Range	Default	Description
coeff_t *	b_f	NA	NA	Pointer to coefficient matrix b_f.
xa_nnlib_shape_t	shape_b_f	NA	NA	Shape information about b_f.
coeff_t *	b_i	NA	NA	Pointer to coefficient matrix b_i.
xa_nnlib_shape_t	shape_b_i	NA	NA	Shape information about b_i.
coeff_t *	b_c	NA	NA	Pointer to coefficient matrix b_c.
xa_nnlib_shape_t	shape_b_c	NA	NA	Shape information about b_c.
coeff_t *	b_o	NA	NA	Pointer to coefficient matrix b_o.
xa_nnlib_shape_t	shape_b_o	NA	NA	Shape information about b_o.

**Note** LSTM requires all weight matrices' and bias vectors' pointers to be 8 bytes aligned.

## 4.2.5 Enums Specific to LSTM

Table 4-21 Enum xa\_nnlib\_lstm\_precision\_t

Element	Description
XA_NNLIB_LSTM_16bx16b	Coef: 16 bits, I/O: 16 bits Fixed Point
XA_NNLIB_LSTM_8bx16b	Coef: 8 bits, I/O: 16 bits Fixed Point
XA_NNLIB_LSTM_8bx8b	Not supported
XA_NNLIB_flt16xflt16	Not supported

**Note** Currently, LSTM only supports the XA\_NNLIB\_LSTM\_16bx16b, XA\_NNLIB\_LSTM\_8bx16b precision setting.

Table 4-22 describes parameter IDs for parameters supported by LSTM. It contains the following columns:

- Parameter ID: Parameter identifier (`param_id`).
- Value type: A pointer (`params`) to a variable of this type is to be passed.
- RW: Indicates whether the parameter can be read (`get`) and/or written (`set`).
- Range: Indicates valid values of the parameter.
- Default: Default value of the parameter.
- Description: Brief description of the parameter.

Table 4-22 LSTM Specific Parameters

Parameter ID	Value Type	RW	Range	Default	Description
XA_NNLIB_LSTM_RESTORE_CONTEXT_OUTPUT	<code>vect_t []</code>	RW	NA	NA	Set previous output. This can be used to set <code>prev_h</code> to specific context (size should be equal to number of output features). Upon set config, the buffer passed is copied to persistent memory; upon get config, it returns the <code>prev_h</code> state in the given buffer.
XA_NNLIB_LSTM_RESTORE_CONTEXT_CELL	<code>vect_t []</code>	RW	NA	NA	Set previous cell state. This can be used to set <code>prev_c</code> to specific cell context (size should be equal to number of output features). Upon set config, the buffer passed is copied to persistent memory; upon get config, it returns the <code>prev_c</code> state in the given buffer.
XA_NNLIB_LSTM_WEIGHT	<code>xa_nnlib_lstm_weights_t</code>	RW	NA	NA	Weight matrices, pointers to weight matrices along with shape information needs to be passed via <code>xa_nnlib_lstm_weights_t</code> structure for set config. Upon get config, it returns pointers to weight matrices along with their shape information in same structure.
XA_NNLIB_LSTM_BIAS	<code>xa_nnlib_lstm_biases_t</code>	RW	NA	NA	Bias vectors, pointers to bias vectors along with shape information needs to be passed via <code>xa_nnlib_lstm_biases_t</code> structure for set config. Upon get config, it returns pointers to bias vectors along with their shape information in same structure.
XA_NNLIB_LSTM_INPUT_SHAPE	<code>xa_nnlib_shape_t</code>	R	NA	NA	Input shape information, get information of the input shape expected by the layer.
A_NNLIB_LSTM_OUTPUT_SHAPE	<code>xa_nnlib_shape_t</code>	R	NA	NA	Output shape information, get information of the output shape expected by layer.

## 4.3 CNN Layer

The CNN APIs are defined in `xa_nnlib_cnn_api.h`.

### 4.3.1 CNN Layer Specification

The CNN layer implements Standard 2D Convolution, Standard 1D Convolution, and Depthwise Separable 2D Convolution. Refer to the equations in Section 3.2.1 for Standard 2D Convolution, Section 3.2.2 for Standard 1D Convolution, and Section 3.2.3 for Depthwise Separable 2D Convolution.

### 4.3.2 Error Codes Specific to CNN

Other than common error codes explained in Section 2.3, the CNN layer may also report the following error codes, which may be generated during the initialization stage.

- `XA_NNLIB_CNN_CONFIG_FATAL_INVALID_ALGO`  
Algorithm is not supported
- `XA_NNLIB_CNN_CONFIG_FATAL_INVALID_PRECISION`  
I/O precision is not supported.
- `XA_NNLIB_CNN_CONFIG_FATAL_INVALID_BIAS_SHIFT`  
Value of Bias shift is not supported
- `XA_NNLIB_CNN_CONFIG_FATAL_INVALID_ACC_SHIFT`  
Value of Accumulator shift is not supported.
- `XA_NNLIB_CNN_CONFIG_FATAL_INVALID_STRIDE`  
Value of strides is not supported
- `XA_NNLIB_CNN_CONFIG_FATAL_INVALID_PADDING`  
Value of padding is not supported.
- `XA_NNLIB_CNN_CONFIG_FATAL_INVALID_INPUT_SHAPE`  
Input shape dimension is not supported.
- `XA_NNLIB_CNN_CONFIG_FATAL_INVALID_OUTPUT_SHAPE`  
Out shape dimension is not supported.
- `XA_NNLIB_CNN_CONFIG_FATAL_INVALID_KERNEL_SHAPE`  
Kernel shape dimension is not supported.
- `XA_NNLIB_CNN_CONFIG_FATAL_INVALID_BIAS_SHAPE`  
Bias shape dimension is not supported.
- `XA_NNLIB_CNN_CONFIG_FATAL_INVALID_PARAM_ID`  
Parameter identifier (`param_id`) is not valid

- `XA_NNLIB_CNN_CONFIG_FATAL_INVALID_PARAM_COMBINATION`

Parameter combination (param\_id) is not valid

The following error codes may be generated during the execution stage.

- `XA_NNLIB_CNN_CONFIG_FATAL_INVALID_INPUT_SHAPE`

Input shape passed during execution does not match with the input shape passed during initialization

## 4.3.3 API Functions Specific to CNN

### 4.3.3.1 Query Functions

Table 4-23 CNN Get Persistent Size Function

<b>Function</b>	<code>xa_nnlb_cnn_get_persistent_fast</code>
<b>Syntax</b>	<pre>Int32 xa_nnlb_cnn_get_persistent_fast (     xa_nnlb_cnn_init_config_t *config)</pre>
<b>Description</b>	Returns persistent memory size in bytes required by CNN layer.
<b>Parameters</b>	Input: <code>config</code> Initial configuration parameters (see Table 4-29).
<b>Errors</b>	If return value is less than 0, then it is an error. Following are the possible error codes: <ul style="list-style-type: none"> <li>■ <code>XA_NNLIB_FATAL_MEM_ALLOC</code></li> <li>■ <code>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_ALGO</code> Algorithm is not supported</li> <li>■ <code>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_PRECISION</code> I/O precision is not supported.</li> <li>■ <code>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_BIAS_SHIFT</code> Value of Bias shift is not supported</li> <li>■ <code>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_ACC_SHIFT</code> Value of Accumulator shift is not supported.</li> <li>■ <code>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_STRIDE</code> Value of strides is not supported</li> <li>■ <code>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_PADDING</code> Value of padding is not supported.</li> </ul>

	<ul style="list-style-type: none"> <li>■ <b>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_INPUT_SHAPE</b> Input shape dimension is not supported.</li> <li>■ <b>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_OUTPUT_SHAPE</b> Out shape dimension is not supported.</li> <li>■ <b>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_KERNEL_SHAPE</b> Kernel shape dimension is not supported.</li> <li>■ <b>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_BIAS_SHAPE</b> Bias shape dimension is not supported</li> <li>■ <b>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_PARAM_ID</b> Parameter identifier (param_id) is not valid</li> <li>■ <b>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_PARAM_COMBINATION</b> Parameter combination (param_id) is not valid</li> </ul>
--	--

Table 4-24 CNN Get Scratch Size Function

<b>Function</b>	<code>xa_nnl-lib_cnn_get_scratch_fast</code>
<b>Syntax</b>	<pre>Int32 xa_nnl-lib_cnn_get_scratch_fast (     xa_nnl-lib_cnn_init_config_t *config)</pre>
<b>Description</b>	Returns scratch memory size in bytes required by CNN layer.
<b>Parameters</b>	Input: <code>config</code> Initial configuration parameters (see Table 4-29).
<b>Errors</b>	If return value is less than 0, then it is an error. Following are the possible error codes: <ul style="list-style-type: none"> <li>■ <b>XA_NNLIB_FATAL_MEM_ALLOC</b></li> <li>■ <b>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_ALGO</b> Algorithm is not supported</li> <li>■ <b>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_PRECISION</b> I/O precision is not supported.</li> <li>■ <b>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_BIAS_SHIFT</b> Value of bias shift is not supported</li> <li>■ <b>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_ACC_SHIFT</b> Value of Accumulator shift is not supported.</li> <li>■ <b>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_STRIDE</b> Value of strides is not supported</li> </ul>

- `XA_NNLIB_CNN_CONFIG_FATAL_INVALID_PADDING`  
Value of padding is not supported.
- `XA_NNLIB_CNN_CONFIG_FATAL_INVALID_INPUT_SHAPE`  
Input shape dimension is not supported.
- `XA_NNLIB_CNN_CONFIG_FATAL_INVALID_OUTPUT_SHAPE`  
Out shape dimension is not supported.
- `XA_NNLIB_CNN_CONFIG_FATAL_INVALID_KERNEL_SHAPE`  
Kernel shape dimension is not supported.
- `XA_NNLIB_CNN_CONFIG_FATAL_INVALID_BIAS_SHAPE`  
Bias shape dimension is not supported.
- `XA_NNLIB_CNN_CONFIG_FATAL_INVALID_PARAM_ID`  
Parameter identifier (`param_id`) is not valid
- `XA_NNLIB_CNN_CONFIG_FATAL_INVALID_PARAM_COMBINATION`  
Parameter combination (`param_id`) is not valid

### 4.3.3.2 Initialization Stage

Table 4-25 CNN Init Function

<b>Function</b>	<code>xa_nnlb_cnn_init</code>
<b>Syntax</b>	<pre>int xa_nnlb_cnn_init (     xa_nnlb_handle_t handle,     xa_nnlb_cnn_init_config_t *config)</pre>
<b>Description</b>	Reset the CNN layer API handle into its initial state. Set up the CNN layer to the specified initial configuration parameters.
<b>Parameters</b>	<p>Input: <code>handle</code>            Pointer to the component persistent memory. This is the opaque handle. Required size: see <code>xa_nnlb_cnn_get_persistent_fast</code>. Required alignment: 8 bytes.</p> <p>Input: <code>config</code>            Initial configuration parameters (see Table 4-29). Note that the initial configuration parameters <i>must</i> be identical to those passed to query functions.</p>
<b>Errors</b>	<p>If the return value is not <code>XA_NNLB_NO_ERROR</code>, it implies that the function has encountered one of the following errors:</p> <ul style="list-style-type: none"> <li>■ <code>XA_NNLB_FATAL_MEM_ALLOC</code> One of the pointers is invalid.</li> <li>■ <code>XA_NNLB_FATAL_MEM_ALIGN</code> One of the pointers is not properly aligned.</li> <li>■ <code>XA_NNLB_CNN_CONFIG_FATAL_INVALID_ALGO</code> Algorithm is not supported.</li> <li>■ <code>XA_NNLB_CNN_CONFIG_FATAL_INVALID_PRECISION</code> I/O precision is not supported.</li> <li>■ <code>XA_NNLB_CNN_CONFIG_FATAL_INVALID_BIAS_SHIFT</code> Value of Bias shift is not supported.</li> <li>■ <code>XA_NNLB_CNN_CONFIG_FATAL_INVALID_ACC_SHIFT</code> Value of Accumulator shift is not supported.</li> <li>■ <code>XA_NNLB_CNN_CONFIG_FATAL_INVALID_STRIDE</code> Value of strides is not supported.</li> <li>■ <code>XA_NNLB_CNN_CONFIG_FATAL_INVALID_PADDING</code> Value of padding is not supported.</li> </ul>



- |  |  |
|--|--|
|  | <ul style="list-style-type: none"><li>■ <code>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_INPUT_SHAPE</code><br/>Input shape dimension is not supported.</li><li>■ <code>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_OUTPUT_SHAPE</code><br/>Out shape dimension is not supported.</li><li>■ <code>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_KERNEL_SHAPE</code><br/>Kernel shape dimension is not supported.</li><li>■ <code>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_BIAS_SHAPE</code><br/>Bias shape dimension is not supported.</li><li>■ <code>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_PARAM_ID</code><br/>Parameter identifier (param_id) is not valid.</li><li>■ <code>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_PARAM_COMBINATION</code><br/>Parameter combination (param_id) is not valid.</li></ul> |
|--|--|

### 4.3.3.3 Execution Stage

Table 4-26 CNN Execution Function

<b>Function</b>	<code>xa_nnlib_cnn_process</code>
<b>Syntax</b>	<pre>int xa_nnlib_cnn_process (     xa_nnlib_handle_t handle,     void *scratch,     void *input,     void *output,     xa_nnlib_shape_t *p_in_shape,     xa_nnlib_shape_t *p_out_shape)</pre>
<b>Description</b>	Processes one input shape to generate one output shape.
<b>Parameters</b>	<p><b>Input: <code>handle</code></b> The opaque component handle. Required alignment: 8 bytes.</p> <p><b>Input: <code>scratch</code></b> A pointer to the scratch buffer. Required alignment: 8 bytes.</p> <p><b>Input: <code>input</code></b> A pointer to the input buffer. Input buffer contains input data. Required alignment: 8 bytes.</p> <p><b>Output: <code>output</code></b> A pointer to the output buffer. Output is written to the output buffer. Required alignment: 8 bytes.</p> <p><b>Input/Output: <code>p_in_shape</code></b> Pointer to the shape containing input buffer dimensions. Contains the length of input data passed to the CNN layer. Required alignment: 4 bytes.</p> <p><b>Output: <code>p_out_shape</code></b> Pointer to the shape for output buffer dimensions. Upon return, <code>*p_out_shape</code> is filled with the length of output generated by the CNN layer. Required alignment: 4 bytes.</p>
<b>Errors</b>	<p>If the return value is not <code>XA&gt;NNLIB_NO_ERROR</code>, it implies that the function has encountered one of the following errors:</p> <ul style="list-style-type: none"> <li>■ <code>XA&gt;NNLIB_FATAL_MEM_ALLOC</code> One of the pointers is NULL</li> </ul>

	<ul style="list-style-type: none"> <li>■ <b>XA_NNLIB_FATAL_MEM_ALIGN</b> One of the pointers is not having required alignment</li> <li>■ <b>XA_NNLIB_FATAL_INVALID_SHAPE</b> Input shape passed during execution does not match with the input shape passed during initialization</li> </ul>
--	--

Table 4-27 CNN Set Parameter Function Details

<b>Function</b>	<code>xa_nnl-lib-cnn-set-config</code>
<b>Syntax</b>	<pre>int xa_nnl-lib-cnn-set-config (     xa_nnl-lib-handle_t handle,     xa_nnl-lib-cnn-param-id_t param_id,     void *params)</pre>
<b>Description</b>	Sets the parameter specified by <code>param_id</code> to the value passed in the buffer pointed to by <code>params</code> .
<b>Parameters</b>	<p><b>Input:</b> <code>handle</code> The opaque component handle. Required alignment: 8 bytes.</p> <p><b>Input:</b> <code>param_id</code> Identifies the parameter to be written. Refer to Table 4-32 for the list of supported parameters.</p> <p><b>Input:</b> <code>params</code> A pointer to a buffer that contains the parameter value. Required alignment: 4 bytes.</p>
<b>Errors</b>	<p>If the return value is not <code>XA_NNLIB_NO_ERROR</code>, it implies that the function has encountered one of the following errors:</p> <ul style="list-style-type: none"> <li>■ <b>XA_NNLIB_FATAL_MEM_ALLOC</b> One of the pointers (<code>handle</code> or <code>params</code>) is <code>NULL</code>.</li> <li>■ <b>XA_NNLIB_FATAL_MEM_ALIGN</b> One of the pointers (<code>handle</code> or <code>params</code>) is not aligned correctly.</li> <li>■ <b>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_PARAM_ID</b> Parameter identifier (<code>param_id</code>) is not valid.</li> </ul>

Table 4-28 CNN Get Parameter Function Details

<b>Function</b>	<code>xa_nnlib_cnn_get_config</code>
<b>Syntax</b>	<pre>int xa_nnlib_cnn_get_config(     xa_nnlib_handle_t handle,     xa_nnlib_cnn_param_id_t param_id,     void *params )</pre>
<b>Description</b>	Gets the value of the parameter specified by <code>param_id</code> in the buffer pointed to by <code>params</code> .
<b>Parameters</b>	<p><b>Input:</b> <code>handle</code> The opaque component handle. Required alignment: 8 bytes.</p> <p><b>Input:</b> <code>param_id</code> Identifies the parameter to be read. Refer to Table 4-32 for the list of supported parameters.</p> <p><b>Output:</b> <code>params</code> A pointer to a buffer that is filled with the parameter value when the function returns. Required alignment: 4 bytes.</p>
<b>Errors</b>	<p>If the return value is not <code>XA_NNLIB_NO_ERROR</code>, it implies that the function has encountered one of the following errors:</p> <ul style="list-style-type: none"> <li>■ <code>XA_NNLIB_FATAL_MEM_ALLOC</code> One of the pointers (<code>handle</code> or <code>params</code>) is <code>NULL</code>.</li> <li>■ <code>XA_NNLIB_FATAL_MEM_ALIGN</code> One of the pointers (<code>handle</code> or <code>params</code>) is not aligned correctly.</li> </ul> <p><code>XA_NNLIB_CNN_CONFIG_FATAL_INVALID_PARAM_ID</code> Parameter identifier (<code>param_id</code>) is not valid.</p>

### 4.3.4 Structures Specific to CNN

Table 4-29 CNN Config Structure xa\_nnlib\_cnn\_init\_config\_t

Element Type	Element Name	Range	Default	Description
xa_nnlib_shape_t	input_shape	NA	height = 16 width = 16 channels = 4	Input shape dimensions
Int32	output_height	NA	16	Output height
Int32	output_width	NA	16	Output width
Int32	output_channels	NA	4	Output depth or channels
Int32	output_format	0 or 1	0	Output data format 0: SHAPE_CUBE_DWH_T 1: SHAPE_CUBE_WHD_T
xa_nnlib_shape_t	kernel_std_shape	NA	height = 16 width = 16 channels = 4	Standard 1D/2D Convolution Kernel (Filter) shape dimensions output_channels indicate number of kernels
xa_nnlib_shape_t	kernel_ds_depth_shape	NA	NA	Depthwise Separable 2D Convolution - Depthwise Kernel (filter) Dimensions
xa_nnlib_shape_t	kernel_ds_point_shape	NA	NA	Depthwise Separable 2D Convolution - Pointwise Kernel (filter) Dimensions
xa_nnlib_shape_t	bias_std_shape	NA	channels = 4	Standard 1D/2D Convolution Bias dimensions
xa_nnlib_shape_t	bias_ds_depth_shape	NA	NA	Depthwise Separable 2D Convolution - Depthwise Bias) Dimensions
xa_nnlib_shape_t	bias_ds_point_shape	NA	NA	Depthwise Separable 2D Convolution – Pointwise Bias Dimensions
xa_nnlib_cnn_precision_t	precision	XA_NNLIB_CNN_16b16b, XA_NNLIB_CNN_8b16b, XA_NNLIB_CNN_8b8b, XA_NNLIB_CNN_f32xf32	XA_NNLIB_CNN_8b16b	Kernel (filter), input, output precision setting
Int32	bias_shift	-31 to 31	7	Q-format adjustment for bias before addition into accumulator, +/- value - left/right shift

Element Type	Element Name	Range	Default	Description
Int32	acc_shift	-31 to 31	-7	Q-format adjustment for accumulator before rounding to result, +/- value - left/right shift
Int32	channels_multiplier	NA	NA	Depthwise Separable 2D Convolution - channel multiplier. (channels_multiplier * input_channels) must be multiple of 4
Int32	x_padding	NA	2	Left side padding to be added to input
Int32	y_padding	NA	2	Top padding to be added to input
Int32	x_stride	NA	2	Strides over padded input in width dimension
Int32	y_stride	NA	2	Strides over padded input in height dimension
xa_nnlb_cnn_algo_t	algo	NA	XA_NNLB_CNN_CONV2D_STD	Convolution algorithm

### 4.3.5 Enums Specific to CNN

Table 4-30 Enum xa\_nnlb\_cnn\_precision\_t

Element	Description
XA_NNLB_CNN_16b_x16b	Coef: 16 bits, I/O: 16 bits fixed point
XA_NNLB_CNN_8b_x16b	Coef: 8 bits, I/O: 16 bits fixed point
XA_NNLB_CNN_8b_x8b	Coef: 8 bits, I/O: 8 bits fixed point
XA_NNLB_CNN_f32_xf32	Coef: single precision float, I/O: single precision float

Table 4-31 Enum xa\_nnlb\_cnn\_algo\_t

Element	Description
XA_NNLB_CNN_CONV1D_ST	Standard 1D Convolution
XA_NNLB_CNN_CONV2D_STD	Standard 2D Convolution
XA_NNLB_CNN_CONV2D_DS	Depthwise Separable 2D Convolution

Table 4-32 describes parameter IDs for parameters supported by CNN. It contains the following columns:

- Parameter ID: Parameter identifier (`param_id`).
- Value type: A pointer (`params`) to a variable of this type is to be passed.
- RW: Indicates whether the parameter can be read (`get`) and/or written (`set`).
- Range: Indicates valid values of the parameter.
- Default: Default value of the parameter
- Description: Brief description of the parameter.

Table 4-32 CNN Specific Parameters

Parameter ID	Value Type	RW	Range	Default	Description
XA_NNLIB_CNN_KERNEL	<code>vect_t</code> []	<i>RW</i>	NA	NA	Kernel shape information, get or set information of the kernel shape expected by the layer
XA_NNLIB_CNN_BIAS	<code>vect_t</code> []	<i>RW</i>	NA	NA	Bias shape information, get or set information of the bias shape expected by the layer
XA_NNLIB_CNN_INPUT_SHAPE	<code>xa_nnlib_shape_t</code>	<i>R</i>	NA	NA	Input shape information, get information of the input shape expected by the layer.
XA_NNLIB_CNN_OUTPUT_SHAPE	<code>xa_nnlib_shape_t</code>	<i>R</i>	NA	NA	Output shape information, get information of the output shape produced by layer.

## 5. Introduction to the Example Testbench

---

The HiFi1 NN library is released as .tgz file for linux/makefile based usage and .xws file for Xtensa Xplorer based usage.

### 5.1 Making the Library

If you have source code distribution, you must build the NN library before you can build the testbench. To do so, follow these steps:

1. Go to `build`.
2. From the command prompt, enter:  

```
xt-make -f makefile clean all install
```

The NN library `xa_nnlib.a` will be built and copied to the `lib` directory.

### 5.2 Making the Executable

To build the testbenches, follow these steps:

1. Go to `test/build`.
2. From the command-line prompt, enter:  

```
xt-make -f makefile_testbench_sample clean all
```

This will build the example testbenches for all the kernels and layers.

The following header files are common and used by all testbenches.

- Testbench header files (`test/include`)
  - `xt_profiler.h`
  - `cmdline_parser.h`
  - `file_io.h`
  - `xt_manage_buffers.h`

To build and execute the example testbenches from .xws based release package, please refer to the `readme.html` files available in the imported example testbench projects.

The following sections describe each low-level kernel and layer testbench.



## 5.3 Sample Testbench for Matrix X Vector Multiplication Kernels

The NN library Matrix X Vector Multiplication Kernels are provided with a sample testbench application. The supplied testbench consists of the following files:

- Testbench source files (test/src)
  - xa\_nn\_matXvec\_testbench.c

### 5.3.1 Usage

The NN library Matrix X Vector Multiplication Kernels executable can be run with command-line options as follows.

```
$ xt-run [--mem_model] [--turbo] xa_nn_matXvec_test [options]
```

Following are available options:

Option	Description	Additional Information
-rows	Rows of mat1.	
-cols1	Columns of mat1 and rows of mat2 (Default=32)	Columns of mat1 must be multiple of 4
-cols2	Columns of mat2 (Default=32)	Columns of mat2 must be multiple of 4
-row_stride1	Row stride for mat1 (Default=32)	
-row_stride2	Row stride for mat2 (Default=32)	
-vec_count	Vec count for Time batching (Default=1)	
-acc_shift	Accumulator shift (Default=-7)	
-bias_shift	Bias shift (Default=7)	
-mat_precision	8, 16, -1 (single precision float), -3 (asym8u) or -5 (sym8s); (Default=16)	
-inp_precision	8, 16, -1 (single precision float), -3 (asym8u) or -4 (asym8s); (Default=16)	
-out_precision	8, 16, 32, 64 -1 (single precision float), -3 (asym8u), -4 (asym8s) or -7 (asym16s); (Default=16)	
-bias_precision	8, 16, 64 -1 (single precision float), 32 (for quantized 8-bit kernels); (Default=16)	
-membank_padding	0 or 1 (Default=1)	
-frames	Positive number; (Default=2)	
-activation	Sigmoid, tanh, relu or softmax (Default= bypass i.e. no activation for output)	

Option	Description	Additional Information
-write_file	Set to 1 to write input and output vectors to file; (Default=0)	
-read_inp_file_name	Full filename for reading inputs (order - mat1, vec1, mat2, vec2, bias)	
-read_ref_file_name	Full filename for reading reference output	
-write_inp_file_name	Full filename for writing inputs (order - mat1, vec1, mat2, vec2, bias)	
-write_out_file_name	Full filename for writing output	
-verify	Verify output against provided reference	0: Disable, 1: Bit exact match (Default=1)
-batch	Flag to execute time batching kernels	0: Disable, 1: Enable (Default=0)
-fc	Flag to execute fully connected kernels	0: Disable, 1: Enable (Default=0)
-help	Prints help	

If no command line arguments are given, the Matrix X Vector Multiplication Kernels sample testbench runs with default values.

## 5.4 Sample Testbench for Convolution Kernels

The NN library convolutional kernels are provided with a sample testbench application. The supplied testbench consists of the following files:

- Testbench source files (test/src)
  - xa\_nn\_conv\_testbench.c

### 5.4.1 Usage

The NN Library convolutional kernels executable can be run with command-line options as follows.

```
$ xt-run [--mem_model] [--turbo] xa_nn_conv_test [options]
```

Following are available options:

Option	Description
-input_height	Input height (Default=16)
-input_width	Input width (Default=16)
-input_channels	Input channels (Default=4)
-kernel_height	Kernel height (Default=3)

Option	Description
-kernel_width	Kernel width (Default=3)
-out_channels	Out channels (Default=4)
-channels_multiplier	Channel Multiplier (Default=1)
-x_stride	Stride in width dimension (Default=2)
-y_stride	Stride in height dimension (Default=2)
-x_padding	Left padding in width dimension (Default=2)
-y_padding	Top padding in height dimension (Default=2)
-dilation_height	Kernel Dilation factor in height dimension (Default=1)
-dilation_width	Kernel Dilation factor in width dimension (Default=1)
-out_height	Output height (Default=16)
-out_width	Output width (Default=16)
-bias_shift	Bias shift (Default=7)
-acc_shift	Accumulator shift (Default=-7)
-inp_data_format	0 (SHAPE_CUBE_DWH_T), 1 (SHAPE_CUBE_WHD_T) (Default=1)
-out_data_format	0 (SHAPE_CUBE_DWH_T), 1 (SHAPE_CUBE_WHD_T) (Default=0)
-inp_precision	8, 16, -1(single precision float), -3 (asym8u) or -4 (asym8s); (Default=16)
-kernel_precision	8, 16, -1(single precision float), -3 (asym8u) or -5 (sym8s); (Default=8)
-out_precision	8, 16, -1(single precision float), -3 (asym8u) or -4 (asym8s); (Default=16)
-bias_precision	8, 16, -1(single precision float), 32(for quantized 8-bit kernels) (Default=16)
-input_zero_bias	Input zero bias for quantized 8-bit, -255 to 0 for asym8u, -127 to 128 for asym8s; Default=-127
-kernel_zero_bias	Kernel zero_bias for quantized 8-bit, -255 to 0 for asym8u, ignored for sym8s; Default=-127
-out_multiplier	Output multiplier in Q31 format for asym8, 0x0 to 0x7fffff; Default=0x4000000
-out_shift	Output shift for quantized 8-bit(asym8u and asym8s), 31 to -31; Default=-8

Option	Description
-out_zero_bias	Output zero bias for quantized 8-bit, 0 to 255 for asym8u, -128 to 127 for asym8s; Default=128
-frames	Positive number (Default=2)
-kernel_name	conv2d_std, conv2d_depth, conv1d_std, dilated_conv2d_std; (Default= conv2d_std)
-pointwise_profile_only	Applicable only when kernel_name is conv2d_depth, 0 (print conv2d depthwise and pointwise profile info), 1 (print only conv2d pointwise profile info); Default=0
-write_file	Set to 1 to write input and output vectors to file; (Default=0)
-read_inp_file_name	Full filename for reading inputs (order - input, kernel, bias, (pointwise kernel, pointwise bias for depth separable))
-read_ref_file_name	Full filename for reading reference output
-write_inp_file_name	Full filename for writing inputs (order - input, kernel, bias, (pointwise kernel, pointwise bias for depth separable))
-write_out_file_name	Full filename for writing output
-verify	Verify output against provided reference; 0: Disable, 1: Bit exact match (Default=1)
-help	Prints help

If no command line arguments are given, the Convolutional Kernels sample testbench runs with default values.

## 5.5 Sample Testbench for Activation Kernels

The NN library activation kernels are provided with a sample testbench application. The supplied testbench consists of the following files:

- Testbench source files (test/src)
  - xa\_nn\_activations\_testbench.c

### 5.5.1 Usage

The NN library activation kernels executable can be run with command-line options as follows.

```
$ xt-run [--mem_model] [--turbo] xa_nn_activation_test [options]
```

Following are available options:

Option	Description
-num_elements	Number of elements (Default=32)
-relu_threshold	Threshold for relu in Q16.15 (Default= 32768 i.e. =1 in Q16.15)
-inp_precision	8, 16, 32, -1 (single precision float), -3 (asym8u) or -4 (asym8s); (Default=32)
-out_precision	8, 16, 32, -1 (single precision float), -3 (asym8u) or -4 (asym8s); (Default=32)
-frames	Positive number (Default=2)
-activation	Sigmoid, tanh, relu, relu_std, relu1, relu6, activation_min_max, softmax, hard_swish, prelu or leaky_relu (Default= sigmoid)
-write_file	Set to 1 to write input and output vectors to file; (Default=0)
-read_inp_file_name	Full filename for reading input
-read_ref_file_name	Full filename for reading reference output
-write_inp_file_name	Full filename for writing input
-write_out_file_name	Full filename for writing output
-verify	Verify output against provided reference; 0: Disable, 1: Bit exact match (Default=1)
Quantized 8-bit specific parameters	
-diffmin	Diffmin; Default=-15
-input_left_shift	Input_left_shift; Default=27
-input_multiplier	Input_multiplier; Default=2060158080

Option	Description
-activation_max	asym8u input data activation max; Default=0
-activation_min	asym8u input data activation min; Default=0
-input_range_radius	Sigmoid_asym8u input parameter; Default=128
-zero_point	Sigmoid_asym8u input parameter; Default=0
-inp_zero_bias	Zero bias value for input Default=0
-alpha_zero_bias	Prelu parameter - Zero bias value for alpha Default=0
-alpha_multiplier	Prelu parameter - Multiplier value for alpha Default=0x40000000
-alpha_shift	Prelu parameter - Shift value for alpha Default=0
-reluish_multiplier	Hard Swish parameter - Multiplier value for relu scale Default=0x40000000
-reluish_shift	Hard Swish parameter - Shift value for relu scale Default=0
-out_multiplier	Multiplier value for output Default=0x40000000
-out_shift	Shift value for output Default=0
-out_zero_bias	Zero bias value for output Default=0
-help	Prints help

If no command line arguments are given, the Activation Kernels sample testbench runs with default values.

## 5.6 Sample Testbench for Pooling Kernels

The NN library pooling kernels are provided with a sample testbench application. The supplied testbench consists of the following files:

- Testbench source files (test/src)
  - xa\_nn\_pool\_testbench.c

### 5.6.1 Usage

The NN library pooling kernels executable can be run with command-line options as follows.

```
$ xt-run [--mem_model] [--turbo] xa_nn_pool_test [options]
```

Following are available options:

Option	Description
-input_height	Input height (Default=16)
-input_width	Input width (Default=16)
-input_channels	Input channels (Default=4)
-kernel_height	Kernel height (Default=3)
-kernel_width	Kernel width (Default=3)
-x_stride	Stride in width dimension (Default=2)
-y_stride	Stride in height dimension (Default=2)
-x_padding	Left padding in width dimension (Default=2)
-y_padding	Top padding in height dimension (Default=2)
-out_height	Output height (Default=16)
-out_width	Output width (Default=16)
-acc_shift	Accumulator shift (Default=-7)
-inp_data_format	0 (SHAPE_CUBE_DWH_T), 1 (SHAPE_CUBE_WHD_T) (Default=1)
-out_data_format	0 (SHAPE_CUBE_DWH_T), 1 (SHAPE_CUBE_WHD_T) (Default=1)
-inp_precision	8, 16, -1(single precision float), -3(asym8u) (Default=16)
-out_precision	8, 16, -1(single precision float), -3(asym8u) (Default=16)
-frames	Positive number (Default=2)
-kernel_name	avgpool, maxpool (Default=avgpool)
-write_file	set to 1 to write input and output vectors to file; (Default=0)
-read_inp_file_name	Full filename for reading inputs (order - inp)
-read_ref_file_name	Full filename for reading reference output
-write_inp_file_name	Full filename for writing inputs (order - inp)
-write_out_file_name	Full filename for writing output
-verify	Verify output against provided reference; 0: Disable, 1: Bit exact match (Default=1)
-help	Prints help

If no command line arguments are given, the Pooling Kernels sample testbench runs with default values.

## 5.7 Sample Testbench for Basic Operations Kernels

The NN library basic kernels are provided with a sample testbench application. The supplied testbench consists of the following files:

- Testbench source files (test/src)
  - xa\_nn\_basic\_testbench.c

### 5.7.1 Usage

The NN library basic kernels executable can be run with command-line options as follows.

```
$ xt-run [--mem_model] [--turbo] xa_nn_basic_test [options]
```

Following are available options:

Option	Description
-io_length	Input/output vector length; Default=1024
-inp_precision	16, -3 (asym8u), -1 (single prec float), -4(asym8s), 1(bool); Default=-1
-out_precision	-3 (asym8u), -1 (single prec float), -4(asym8s), 1(bool), -10(asym32s); Default=-1
-vec_count	Number of input vectors; Default=1
-frames	Positive number; Default=2
-kernel_name	elm_sub, elm_mul, elm_floor, dot_prod, elm_sine, elm_cosine, elm_logn, elm_abs, elm_ceil, elm_round, elm_neg, elm_square, elm_sqrt, elm_rsqrt, elm_requantize, elm_dequantize, memmove, memset; Default=elm_add
-write_file	Set to 1 to write input and output vectors to file; Default=0
-read_inp1_file_name	Full filename for reading inputs (order - inp)
-read_inp2_file_name	Full filename for reading inputs (order - inp)
-read_ref_file_name	Full filename for reading reference output
-write_inp1_file_name	Full filename for writing inputs (order - inp)
-write_inp2_file_name	Full filename for writing inputs (order - inp)
-write out file name	Full filename for writing output



Option	Description
-verify	Verify output against provided reference; 0: Disable, 1: Bit exact match; Default=1
-read_inp_shape_str	Takes the input shape dimensions(space '' separated) as a string
-read_out_shape_str	Takes the output shape dimensions(space '' separated) as a string
-read_axis_data_str	Takes the axis data (space '' separated) as a string
Broadcast specific parameters	
-input1_numElements	Number of elements in input
-input2_numElements	Number of elements in input
-input1_strides	Input strides
-input2_strides	Input strides
Quantized data types specific parameters	
-output_zero_bias	Output zero bias; Default=127
-output_left_shift	Output_left_shift; Default=1
-output_multiplier	Output_multiplier; Default=0x7fff
-output_activation_min	Output_activation_min; Default=0
-output_activation_max	Output_activation_max; Default = 225
-input1_zero_bias	Input1 zero bias; Default=-127
-input1_left_shift	Input1 left shift; Default=0
-input1_multiplier	Input1 multiplier; Default=0x7fff
-input2_zero_bias	Input2 zero bias; Default=-127
-input2_left_shift	Input2 left shift; Default=0
-input2_multiplier	Input2 multiplier; Default=0x7fff
-left_shift	Global left shift; Default=0
-input1_scale	Input scale; Default=0.5
-h	Prints help

If no command line arguments are given, the Basic Kernels sample testbench runs with default values.

## 5.8 Sample Testbench for Normalization Kernels

The NN library Normalization Kernels are provided with a sample testbench application. The supplied testbench consists of the following files:

- Testbench source files (test/src)
  - xa\_nn\_norm\_testbench.c

### 5.8.1 Usage

The NN library Normalization Kernels executable can be run with command-line options as follows.

```
$ xt-run [--mem_model] [--turbo] xa_nn_norm_test [options]
```

Following are available options:

Option	Description
-num_elms	Number of elements; Default=256
-inp_precision	-4(asym8s); Default=16
-out_precision	-4(asym8s); Default=16
-frames	Positive number; Default=2
-kernel_name	L2_norm; Default=L2_norm
-zero_point	Input Zero point; Default = 0
-write_file	Set to 1 to write input and output vectors to file; Default=0
-read_inp_file_name	Full filename for reading inputs (order - inp)
-read_ref_file_name	Full filename for reading reference output
-write_inp_file_name	Full filename for writing inputs (order - inp)
-write_out_file_name	Full filename for writing output
-verify	Verify output against provided reference; 0: Disable, 1: Bit exact match; Default=1
-h	Prints help

If no command line arguments are given, the Normalization Kernels sample testbench runs with default values.

## 5.9 Sample Testbench for Reorg Kernels

The NN library reorg kernels are provided with a sample testbench application. The supplied testbench consists of the following files:

- Testbench source files (test/src)

■ `xa_nn_reorg_testbench.c`

## 5.9.1 Usage

The NN library basic kernels executable can be run with command-line options as follows.

```
$ xt-run [--mem_model] [--turbo] xa_nn_reorg_test [options]
```

Following are available options:

Option	Description
<code>-inp_data_format</code>	Data format of input and output, 0 for nhwc; Default=0
<code>-num_inp_dims</code>	Number of input dimensions; Default=4
<code>-num_pad_dims</code>	Number of pad dimensions; Default=2
<code>-num_out_dims</code>	Number of output dimensions; Default=4
<code>-pad_value</code>	Input to be padded with this pad value; Default=0
<code>-input_height</code>	Input height; Default=16
<code>-input_width</code>	Input width; Default=16
<code>-input_channels</code>	Input channels; Default=16
<code>-block_size</code>	Block size; Default=2
<code>-out_height</code>	Output height; Default=16
<code>-out_width</code>	Output width; Default=16
<code>-out_channels</code>	Output channels; Default=4
<code>-inp_precision</code>	8; Default=8
<code>-out_precision</code>	8; Default=8
<code>-frames</code>	Positive number; Default=2
<code>-kernel_name</code>	depth_to_space, space_to_depth, pad, batch_to_space_nd, space_to_batch_nd; Default=depth_to_space
<code>-write_file</code>	Set to 1 to write input and output vectors to file; Default=0
<code>-read_inp_file_name</code>	Full filename for reading inputs (order - inp)
<code>-read_ref_file_name</code>	Full filename for reading reference output
<code>-write_inp_file_name</code>	Full filename for writing inputs (order - inp)
<code>-write_out_file_name</code>	Full filename for writing output
<code>-verify</code>	Verify output against provided reference; 0
<code>-inp_shape</code>	Takes the input shape dimensions (num_inp_dims values space '' separated)

Option	Description
-pad_shape	Takes the pad shape dimensions (num_pad_dims values space '' separated)
-out_shape	Takes the output shape dimensions (num_out_dims values space '' separated)
-pad_values	Takes the pad values (prod(pad_shape) values space '' separated)
-block_sizes	Takes the block sizes ((num_inp_dims-2) values space '' separated) for batch_to_space_nd and space_to_batch_nd kernels
-crop_or_pad_sizes	Takes the crop sizes for batch_to_space_nd or pad sizes for space_to_batch_nd (2*(num_inp_dims-2) values space '' separated)
-h	Prints help.

If no command line arguments are given, the Reorg Kernels sample testbench runs with default values.

## 5.10 Sample Testbench for GRU Layer

The NN library GRU layer is provided with a sample testbench application. The supplied testbench consists of the following files:

- Testbench source files (test/src)
  - xa\_nn\_gru\_testbench.c

### 5.10.1 Usage

The NN library GRU executable can be run with command-line options as follows.

```
$ xt-run [--mem_model] [--turbo] xa_nn_gru_test [options]
```

Following are available options:

Option	Description	Additional Information
--in_feats	Input length (Default=256)	Range: 4-2048 <b>Note:</b> Input length must be multiple of 4
--out_feats	Output length (Default=256)	Range: 4-2048 <b>Note:</b> Output length must be multiple of 4
--membank_padding	Memory bank padding (Default=1)	Must be 0 or 1

Option	Description	Additional Information
--mat_prec	Coefficient precision (Default=16)	Must be 8 or 16
--vec_prec	Input precision (Default=16)	Must be 16
--verify	Verify output against ref output (Default=1)	Supported values: 0: Disable, 1: Enable
--input_file	Input file name	
--filter_path	Path where file containing filter are stored	
--output_file	File to which output will be written	
--prev_h_file	File containing context data	
--ref_file	File which has ref output	
-help	Prints help	

If no command line arguments are given, the GRU sample testbench runs with default values.

## 5.11 Sample Testbench for LSTM Layer

The NN library LSTM layer is provided with a sample testbench application. The supplied testbench consists of the following files:

- Testbench source files (test/src)
  - xa\_nn\_lstm\_testbench.c

### 5.11.1 Usage

The NN library LSTM executable can be run with command-line options as follows.

```
$ xt-run [--mem_model] [--turbo] xa_nn_lstm_test [options]
```

Following are available options:

Option	Description	Additional Information
--in_feats	Input length (Default=256)	Range: 4-2048 <b>Note:</b> Input length must be multiple of 4
--out_feats	Output length (Default=256)	Range: 4-2048 <b>Note:</b> Output length must be multiple of 4
--membank_padding	Memory bank padding (Default=1)	Must be 0 or 1
--mat_prec	Coefficient precision (Default=16)	Must be 8 or 16
--vec_prec	Input precision (Default=16)	Must be 16
--verify	Verify output against ref output (Default=1)	Supported values: 0: Disable, 1: Enable
--input_file	File containing input shape	

Option	Description	Additional Information
--filter_path	Path where file containing filter are stored	
--output_file	File to which output will be written	
--output_cell_file	File to which cell output will be written	
--prev_h_file	File containing context (previous output) data	
--prev_c_file	File containing context (previous cell state) data	
--ref_file	File which has ref output	
--ref_cell_file	File which has ref cell output	
-help	Prints help	

If no command line arguments are given, the LSTM sample testbench runs with default values.

## 5.12 Sample Testbench for CNN Layer

The NN library CNN layer is provided with a sample testbench application. The supplied testbench consists of the following files:

- Testbench source files (test/src)
  - xa\_nn\_cnn\_testbench.c

### 5.12.1 Usage

The NN Library CNN executable can be run with command-line options as follows.

```
$ xt-run [--mem_model] [--turbo] xa_nn_cnn_test [options]
```

Following are available options:

Option	Description
-input_height	Input height (Default=16)
-input_width	Input width (Default=16)
-input_channels	Input channels (Default=4)
-kernel_height	Kernel height (Default=3)
-kernel_width	Kernel width (Default=3)
-out_channels	Out channels (Default=4)
-channels_multiplier	Channel Multiplier (Default=1)
-x_stride	Stride in width dimension (Default=2)
-y_stride	Stride in height dimension (Default=2)

Option	Description
-x_padding	Left padding in width dimension (Default=2)
-y_padding	Top padding in height dimension (Default=2)
-out_height	Output height (Default=16)
-out_width	Output width (Default=16)
-bias_shift	Bias shift (Default=7)
-acc_shift	Accumulator shift (Default=-7)
-out_data_format	Output data format, 0 (SHAPE_CUBE_DWH_T), 1 (SHAPE_CUBE_WHD_T); (Default=0)
-inp_precision	8, 16, -1(single precision float); (Default=16)
-kernel_precision	8, 16, -1(single precision float); (Default=8)
-out_precision	8, 16, -1(single precision float); (Default=16)
-bias_precision	8, 16, -1(single precision float); (Default=16)
-frames	Positive number; (Default=2)
-kernel_name	conv2d_std, conv2d_depth, conv1d_std; (Default=conv2d_std)
-write_file	Set to 1 to write input and output vectors to file; (Default=0)
-read_inp_file_name	Full filename for reading inputs (order - input, kernel, bias, (pointwise kernel, pointwise bias for depth separable))
-read_ref_file_name	Full filename for reading reference output
-write_inp_file_name	Full filename for writing inputs (order - input, kernel, bias, (pointwise kernel, pointwise bias for depth separable))
-write_out_file_name	Full filename for writing output
-verify	Verify output against provided reference; 0: Disable, 1: Bit exact match; Default=1
-help	Prints help

If no command line arguments are given, the CNN sample testbench runs with default values.

## 6. References

---

- [1] Reference Wiki page for GRU. [https://en.wikipedia.org/wiki/Gated\\_recurrent\\_unit](https://en.wikipedia.org/wiki/Gated_recurrent_unit)
- [2] TF Micro Lite speech recognition example:  
[https://github.com/tensorflow/tensorflow/tree/r2.3/tensorflow/lite/micro/examples/micro\\_speech](https://github.com/tensorflow/tensorflow/tree/r2.3/tensorflow/lite/micro/examples/micro_speech)
- [3] [TensorFlow Lite for Microcontrollers](#)
- [4] TensorFlow XLA Documentation: <https://www.tensorflow.org/xla/broadcasting>  
NumPy Theory: <https://numpy.org/devdocs/user/theory.broadcasting.html>  
General Broadcasting syntax: <https://www.tensorflow.org/guide/tensor#broadcasting>
- [5] 'strides' as defined in the structure 'NDArrayDesc' at  
<https://github.com/tensorflow/tensorflow/blob/master/tensorflow/lite/kernels/internal/common.h>